


Review

- **Agents**
- **agent function vs agent program**
- A **perfectly rational agent**
- **PEAS**



- **Environments**

- are categorized along several dimensions:

observable? deterministic? episodic? static? discrete? single-agent?

- Real world vs Virtual world

- Several basic agent architectures exist:

reflex, reflex with state, goal-based, utility-based

- **Questions:**

1. Why design of a rational intelligent agent or agents is usually difficult?

Where the difficulties come from?

2. Why design of a human-like intelligent agent or agents is more challenging?

Lecture 4: Problem-solving in AI

LI Xiaoan (Dustin) 李孝安, Dr. & Associate Prof.
MARS-lab
School of Computer Science and Engineering
Northwestern Polytechnical University

E-mail: dustinli@nwpu.edu.cn
WeChat group: AI-2021-Li
Cell phone: 18629662731

Outline

- Problem-solving agents
- Problem types
- Problem formulation
- Example problems
- Basic search algorithms
- Summary

Problem-solving agents

function SIMPLE-PROBLEM-SOLVING-AGENT(*percept*) **returns** an action

static: *seq*, an action sequence, initially empty

state, some description of the current world state

goal, a goal, initially null

problem, a problem formulation

state \leftarrow UPDATE-STATE(*state*, *percept*)

if *seq* is empty **then do**

goal \leftarrow FORMULATE-GOAL(*state*)

problem \leftarrow FORMULATE-PROBLEM(*state*, *goal*)

seq \leftarrow SEARCH(*problem*)

action \leftarrow FIRST(*seq*)

seq \leftarrow REST(*seq*)

return *action*


Goal formulation

Problem formulation

Problem-solving process
based on algorithms

Example 1: Road planning in Romania

- Agent: touring in Romania on holiday
 - currently in Arad.
 - Flight leaves tomorrow from Bucharest
- **Formulate goal:**
 - be in Bucharest
- **Formulate problem:**
 - **states**: various cities
 - **actions**: drive between cities
- **Find solution:**
Search(problem)?
 - sequence of cities, e.g., Arad, Sibiu, Fagaras, Bucharest




Challenge questions:

- Q1: If the map is not labelled, what will happen for this problem-solving?
- Q2: How to label the map?
- Q3: Is the path found the best solution in all practical cases?

Problem types

- Deterministic, fully observable → **single-state problem**
 - Agent knows exactly which state it will be in; solution is a sequence
- Non-observable → **sensorless problem (conformant problem)**
 - Agent may have no idea where it is; solution is a sequence
- Nondeterministic and/or partially observable → **contingency problem**
 - percepts provide **new** information about current state
 - often **interleave** search, execution
- Unknown state space → **exploration** problem e.g. ADRobotAir
- Known/labelled state space → **exploitation** problem


Example 1: ADRobot



Example 2: vacuum world


- P1: Single-state, start in #5.
Solution?

[Right, Suck]



Example: vacuum world

- **P2: Sensorless**, start in {1,2,3,4,5,6,7,8} e.g.,
Right goes to {2,4,6,8}
Solution?
[Right,Suck,Left,Suck]




Example: vacuum world

- P3: Contingency

- Nondeterministic: *Suck* may dirty a clean carpet
- Partially observable: location, dirt at current location.
- Percept: $[L, \text{Clean}]$, i.e., start in #5 or #7

Solution?

[Right, if dirt then Suck]



Single-state problem formulation


A **problem** is defined by four items:

1. **initial state** e.g., "at Arad"
 2. **actions** or **successor function** $S(x)$ = set of action–state pairs
 - e.g., $S(\text{Arad}) = \{ \langle \text{Arad} \rightarrow \text{Zerind}, \text{Zerind} \rangle, \dots \}$
 3. **goal test**, can be
 - **explicit**, e.g., $x = \text{"at Bucharest"}$
 - **implicit**, e.g., $\text{Checkmate}(x)$
 4. **path cost** (additive)
 - e.g., sum of distances, number of actions executed, etc.
 - $c(x, a, y)$ is the **step cost**, assumed to be ≥ 0
- A **solution** is a sequence of actions leading from the initial state to a goal state

Selecting a state space

- Real world is absurdly complex
 - state space must be **abstracted** for problem solving
- (Abstract) state = set of real states
- (Abstract) action = complex combination of real actions
 - e.g., "Arad → Zerind" represents a complex set of possible routes, detours, rest stops, etc.
- For guaranteed realizability, **any** real state "in Arad" must get to **some** real state "in Zerind"
- (Abstract) solution =
 - set of real paths that are solutions in the real world
- Each abstract action should be "easier" than the original problem

State space for the Vacuum world



- states? integer dirt and robot location
- actions? *Left, Right, Suck*
- goal test? no dirt at all locations
- path cost? 1 per action

Example 3: The 8-puzzle

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- states?
- actions?
- goal test?
- path cost?

Example 3: The 8-puzzle

7	2	4
5		6
8	3	1

Start State


	1	2
3	4	5
6	7	8

Goal State

- states? locations of tiles
- actions? move blank left, right, up, down
- goal test? = goal state (given)
- path cost? 1 per move

[Note: optimal solution of n -Puzzle family is **NP-hard**]

Example 4: robotic assembly



- states?: real-valued coordinates of robot joint angles, parts of the object to be assembled
- actions?: continuous motions of robot joints
- goal test?: complete assembly
- path cost?: time to execute

Search algorithms


- Strategies
- Algorithms

Tree search algorithms


- Basic idea:
 - offline, simulated exploration of state space by generating successors of already-explored states (a.k.a. ~expanding states)

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
```

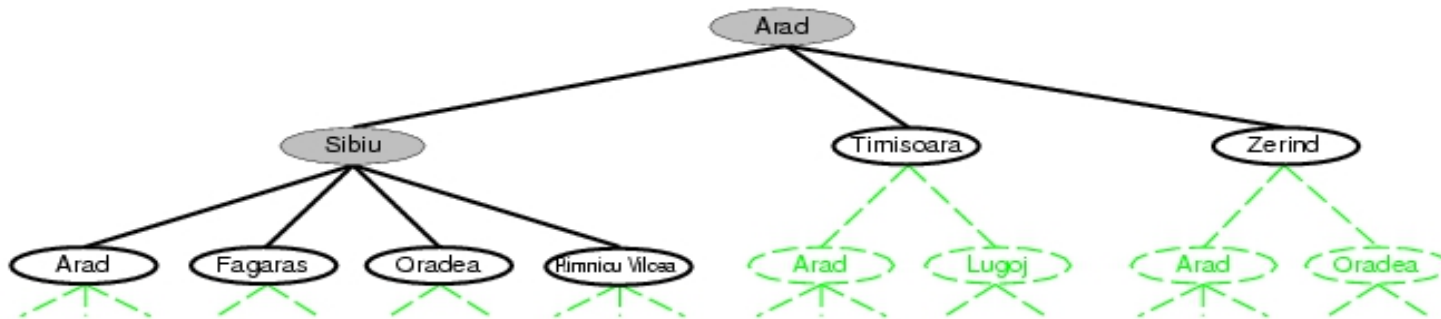
Tree search example



Tree search example



Tree search example




Implementation: general tree search

```
function TREE-SEARCH( problem, fringe) returns a solution, or failure  
  fringe  $\leftarrow$  INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)  
  loop do  
    if fringe is empty then return failure  
    node  $\leftarrow$  REMOVE-FRONT(fringe)  
    if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)  
    fringe  $\leftarrow$  INSERTALL(EXPAND(node, problem), fringe)
```

```
function EXPAND( node, problem) returns a set of nodes  
  successors  $\leftarrow$  the empty set  
  for each action, result in SUCCESSOR-FN[problem](STATE[node]) do  
    s  $\leftarrow$  a new NODE  
    PARENT-NODE[s]  $\leftarrow$  node; ACTION[s]  $\leftarrow$  action; STATE[s]  $\leftarrow$  result  
    PATH-COST[s]  $\leftarrow$  PATH-COST[node] + STEP-COST(node, action, s)  
    DEPTH[s]  $\leftarrow$  DEPTH[node] + 1  
    add s to successors  
  return successors
```

Implementation: states vs. nodes

- A **state** is a (representation of) a physical configuration
- A **node** is a data structure constituting part of a search tree includes **state**, **parent node**, **action**, **path cost $g(x)$** , **depth**



- The `Expand` function creates new nodes, filling in the various fields and using the `SuccessorFn` of the problem to create the corresponding states.

Search strategies *

- A search strategy is defined by picking the **order of node expansion**
- Strategies are evaluated along the following dimensions:
 - **completeness**: does it always find a solution if one exists?
 - **time complexity**: number of nodes generated
 - **space complexity**: maximum number of nodes in memory
 - **optimality**: does it always find a least-cost solution?
- Time and space complexity are measured in terms of
 - b : maximum branching factor of the search tree
 - d : depth of the least-cost solution
 - m : maximum depth of the state space (may be ∞)

Uninformed search strategies


- **Uninformed** search strategies use only the information available in the problem definition

Strategies:

- Breadth-first search
Uniform-cost search
Depth-first search
Depth-limited search
Iterative deepening search


Breadth-first search

- Expand shallowest unexpanded node
- **Implementation:**
 - *fringe* is a FIFO queue, i.e., new successors go at end



Breadth-first search

- Expand shallowest unexpanded node
- **Implementation:**
 - *fringe* is a FIFO queue, i.e., new successors go at end




Breadth-first search

- Expand shallowest unexpanded node


Implementation:

- *fringe* is a FIFO queue, i.e., new successors go at end



Breadth-first search

- Expand shallowest unexpanded node
- **Implementation:**
 - *fringe* is a FIFO queue, i.e., new successors go at end



Properties of breadth-first search

- Complete? Yes (if b is finite)
Time? $1+b+b^2+b^3+\dots +b^d + b(b^d-1) = O(b^{d+1})$
Space? $O(b^{d+1})$ (keeps every node in memory)
Optimal? Yes (if cost = 1 per step)
- **Space** is the bigger problem (more than time)

Time and space complexity are measured in terms of

b : maximum branching factor of the search tree

d : depth of the least-cost solution

m : maximum depth of the state space (may be ∞)

Uniform-cost search

- Expand least-cost unexpanded node
- **Implementation:**
 - *fringe* = queue ordered by path cost
- Equivalent to breadth-first if step costs all equal
- Complete? Yes, if step cost $\geq \epsilon$
- Time? # of nodes with $g \leq$ cost of optimal solution, $O(b^{\text{ceiling}(C^*/\epsilon)})$ where C^* is the cost of the optimal solution
- Space? # of nodes with $g \leq$ cost of optimal solution, $O(b^{\text{ceiling}(C^*/\epsilon)})$
- Optimal? Yes – nodes expanded in increasing order of $g(n)$

Depth-first search

- Expand deepest unexpanded node

Implementation:

- *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node

Implementation:

- *fringe* = LIFO queue, i.e., put successors at front



Depth-first search

- Expand deepest unexpanded node

Implementation:

- *fringe* = LIFO queue, i.e., put successors at front



Depth-first search

- Expand deepest unexpanded node

Implementation:

- *fringe* = LIFO queue, i.e., put successors at front



Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue, i.e., put successors at front



Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue, i.e., put successors at front



Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue, i.e., put successors at front



Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue, i.e., put successors at front



Properties of depth-first search

- Complete? No: fails in infinite-depth spaces, spaces with loops
 - Modify to avoid repeated states along path
 - complete in finite spaces
- Time? $O(b^m)$: terrible if m is much larger than d
 - but if solutions are dense, may be much faster than breadth-first
- Space? $O(bm)$, i.e., linear space!
- Optimal? No

Depth-limited search

= depth-first search with depth limit l ,
i.e., nodes at depth l have no successors

- ```
function DEPTH-LIMITED-SEARCH(problem, limit) returns soln/fail/cutoff
 RECURSIVE-DLS(MAKE-NODE(INITIAL-STATE[problem]), problem, limit)

function RECURSIVE-DLS(node, problem, limit) returns soln/fail/cutoff
 cutoff-occurred? \leftarrow false
 if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)
 else if DEPTH[node] = limit then return cutoff
 else for each successor in EXPAND(node, problem) do
 result \leftarrow RECURSIVE-DLS(successor, problem, limit)
 if result = cutoff then cutoff-occurred? \leftarrow true
 else if result \neq failure then return result
 if cutoff-occurred? then return cutoff else return failure
```

# Iterative deepening search

```
function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution, or fail-
ure
 inputs: problem, a problem
 for depth \leftarrow 0 to ∞ do
 result \leftarrow DEPTH-LIMITED-SEARCH(problem, depth)
 if result \neq cutoff then return result
```

# Iterative deepening search / =0

Limit = 0






# Iterative deepening search / =1

Limit = 1




# Iterative deepening search $l = 2$

Limit = 2



# Iterative deepening search / =3

Limit = 3



# Iterative deepening search

- Number of nodes generated in a ***depth-limited search*** to depth  $d$  with branching factor  $b$ :

$$N_{DLS} = b^0 + b^1 + b^2 + \dots + b^{d-2} + b^{d-1} + b^d$$

- Number of nodes generated in an ***iterative deepening search*** to depth  $d$  with branching factor  $b$ :

$$N_{IDS} = (d+1)b^0 + d b^1 + (d-1)b^2 + \dots + 3b^{d-2} + 2b^{d-1} + 1b^d$$

- For  $b = 10, d = 5$ ,
  - $N_{DLS} = 1 + 10 + 100 + 1,000 + 10,000 + 100,000 = 111,111$
  - $N_{IDS} = 6 + 50 + 400 + 3,000 + 20,000 + 100,000 = 123,456$
- Overhead =  $(123,456 - 111,111)/111,111 = 11\%$

# Properties of iterative deepening search



- Complete? Yes
- Time?  $(d+1)b^0 + d b^1 + (d-1)b^2 + \dots + b^d = O(b^d)$
- Space?  $O(bd)$
- Optimal? Yes, if step cost = 1

# Summary of algorithms\*

| Criterion | Breadth-First | Uniform-Cost                        | Depth-First | Depth-Limited | Iterative Deepening |
|-----------|---------------|-------------------------------------|-------------|---------------|---------------------|
| Complete? | Yes           | Yes                                 | No          | No            | Yes                 |
| Time      | $O(b^{d+1})$  | $O(b^{\lceil C^*/\epsilon \rceil})$ | $O(b^m)$    | $O(b^l)$      | $O(b^d)$            |
| Space     | $O(b^{d+1})$  | $O(b^{\lceil C^*/\epsilon \rceil})$ | $O(bm)$     | $O(bl)$       | $O(bd)$             |
| Optimal?  | Yes           | Yes                                 | No          | No            | Yes                 |

# Repeated states

- Failure to detect repeated states can turn a linear problem into an exponential one!



# Graph search

```
function GRAPH-SEARCH(problem, fringe) returns a solution, or failure
 closed ← an empty set
 fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
 loop do
 if fringe is empty then return failure
 node ← REMOVE-FRONT(fringe)
 if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)
 if STATE[node] is not in closed then
 add STATE[node] to closed
 fringe ← INSERTALL(EXPAND(node, problem), fringe)
```



# Summary

- **Problem formulation** usually requires abstracting away real-world details to define a **state space** that can feasibly be **explored**
- Variety of **uninformed search** strategies
- Iterative deepening search uses only linear space and not much more time than other uninformed algorithms
- **Challenging problem** for **uninformed search** strategies
  - Complexity in space and in time

# Assignment

– Chap 3: exercise 3.1, 3.6

\*Handed in next Tuesday