

Search algorithms

- Uninformed search strategies
 - **blind** search: none extra information or knowledge beyond the definition of the problem itself
 - **Challenging problem** for Uninformed search
- Informed search strategies
 - **heuristic** search: some problem-specific knowledge used
- Questions and Answers

Summary of algorithms*

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening
Complete?	Yes	Yes	No	No	Yes
Time	$O(b^{d+1})$	$O(b^{\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^l)$	$O(b^d)$
Space	$O(b^{d+1})$	$O(b^{\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(bl)$	$O(bd)$
Optimal?	Yes	Yes	No	No	Yes

Challenging from Complexity in space and in time

Lecture 5: Informed search algorithms

Chapter 4

Material

- Chapter 4 Section 1 - 3
- Exclude memory-bounded heuristic search

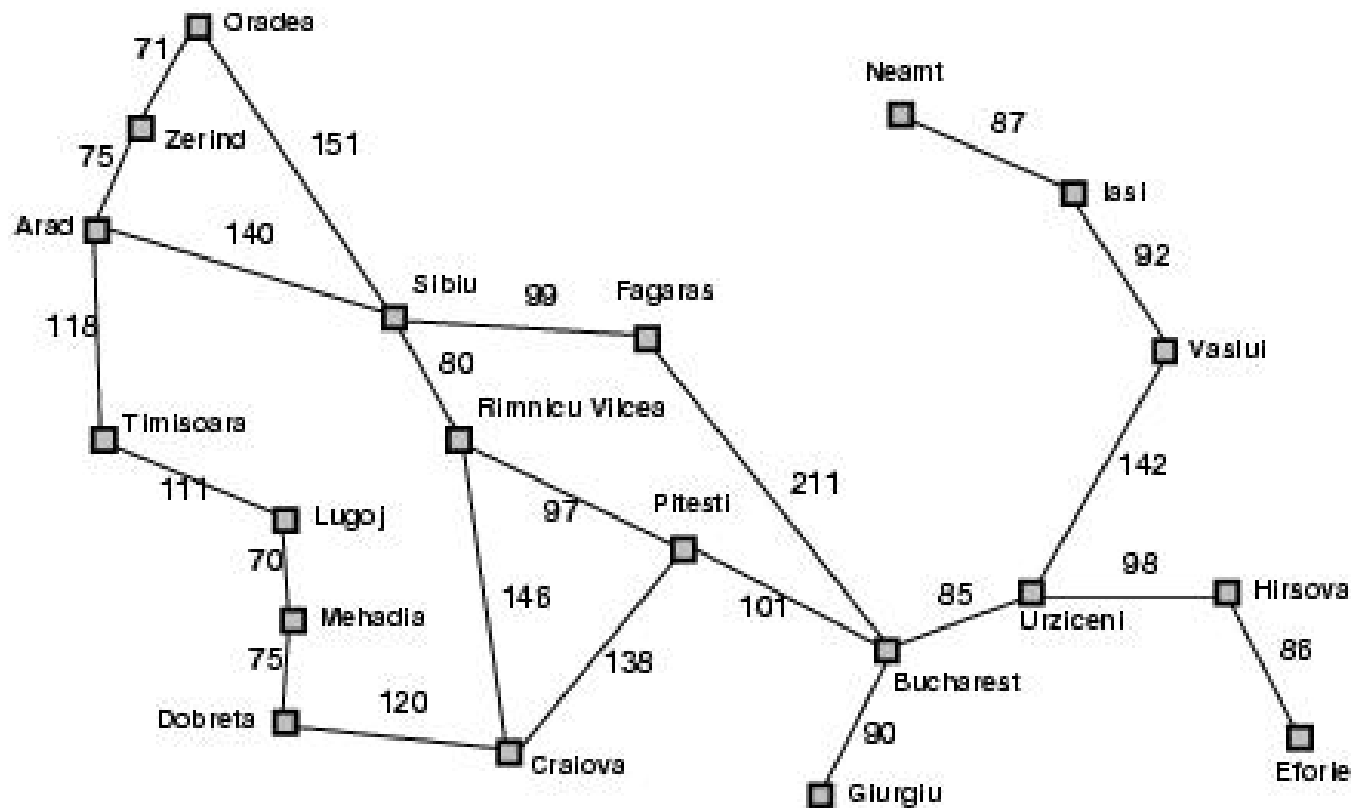
Outline

- Best-first search
- Greedy best-first search
- A^* search
- Heuristics
- Local search algorithms
- Hill-climbing search
- Simulated annealing search
- Local beam search
- Genetic algorithms
- Summary

Best-first search

- Idea: use an **evaluation function** $f(n)$ for each node
 - estimate of "desirability"
Expand most desirable unexpanded node
e.g. the node with the lowest evaluation expanded first
- Implementation:
Order the nodes in fringe in *decreasing* order of desirability
- Special cases:
 - greedy best-first search
 - A* search

Romania with step costs in km



Straight-line distance
to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	10
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

Greedy best-first search

- Evaluation function $f(n) = ?$

$$f(n) = h(n) \text{ (huristic)}$$

– to estimate of cost from n to *goal*


e.g., $h_{SLD}(n)$ = straight-line distance from n to Bucharest

- Greedy best-first search expands the node that **appears** to be closest to goal


Greedy best-first search example



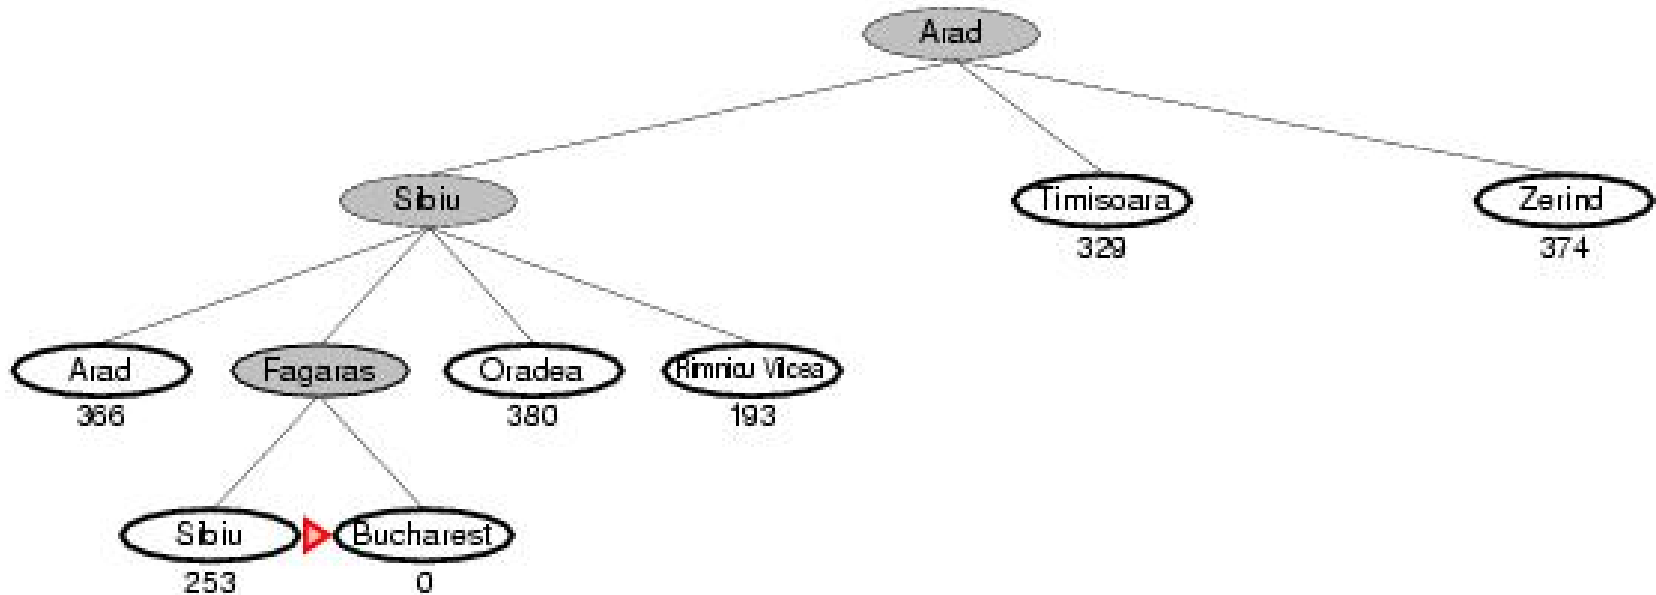
Greedy best-first search example




Greedy best-first search example



Greedy best-first search example

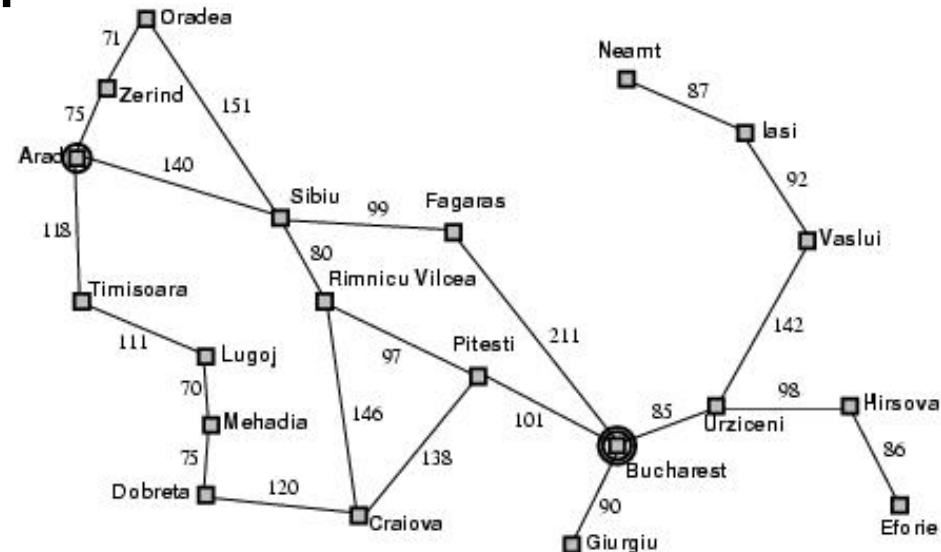


Result based on Greedy BFS



Properties of greedy best-first search

- Complete? No – can get stuck in loops
e.g., *Iasi* \rightarrow *Bucharest*?
- Time? $O(b^m)$, but a good heuristic can give dramatic improvement
- Space? $O(b^m)$ -- keeps all nodes in memory
- Optimal? No



A* search

- **Idea:** to avoid expanding paths that are already expensive, how to improve $f(n)$?
- Evaluation function

$$f(n) = g(n) + h(n)$$

$g(n)$ = cost so far to reach n


$h(n)$ = estimated cost from n to goal

$f(n)$ = estimated total cost of path
through n to goal


A* search example



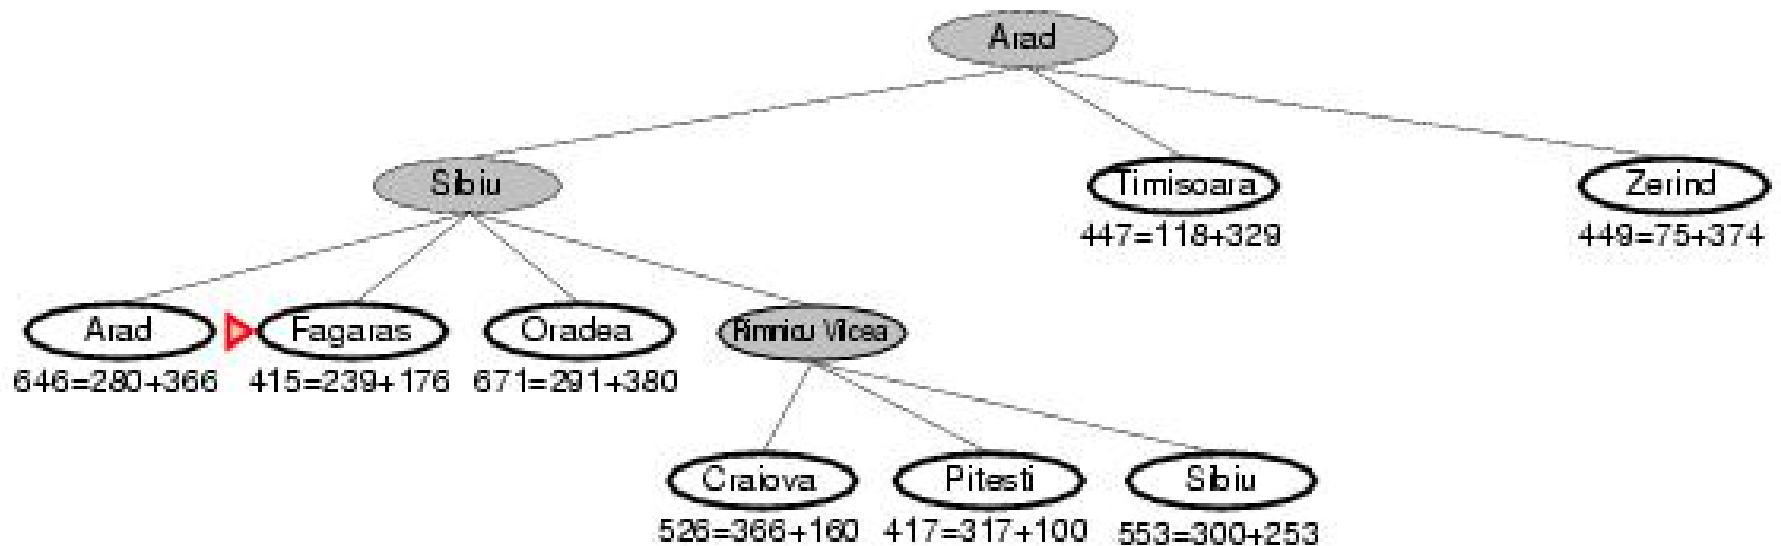
A* search example



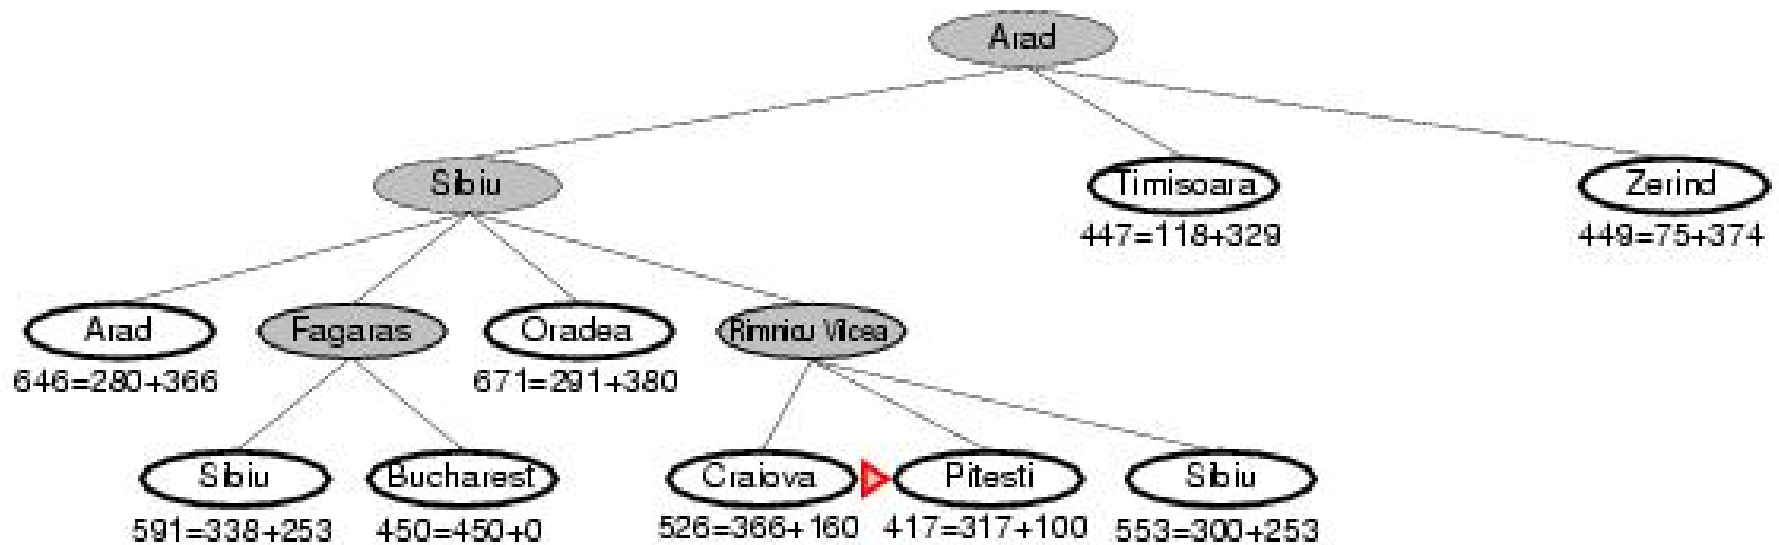
A* search example



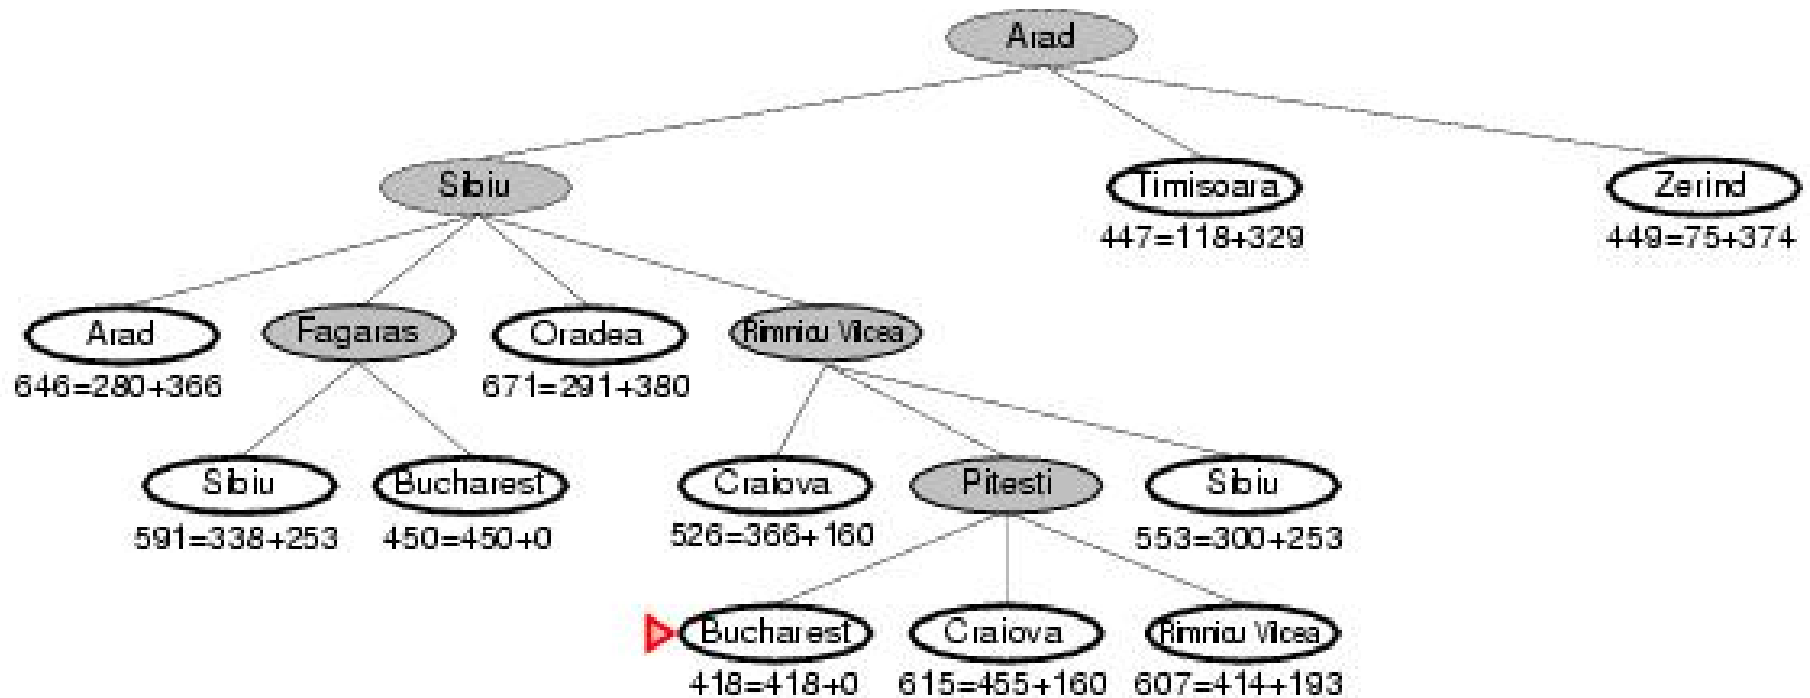
A* search example




A* search example



A* search example



Result based on A*




Admissible heuristics

- A heuristic $h(n)$ is **admissible** if for every node n , $h(n) \leq h^*(n)$, where $h^*(n)$ is the **true** cost to reach the goal state from n .
- An admissible heuristic **never overestimates** the cost to reach the goal, i.e., it is **optimistic**
- Example: $h_{SLD}(n)$ (never overestimates the actual road distance)
- **Theorem:** If $h(n)$ is admissible, A^* using **TREE-SEARCH** is optimal

Optimality of A^* (proof)


- Suppose some suboptimal goal G_2 has been generated and is in the fringe. Let n be an unexpanded node in the fringe such that n is on a shortest path to an optimal goal G .



- $f(G_2) = g(G_2)$ since $h(G_2) = 0$
- $g(G_2) > g(G)$ since G_2 is suboptimal
- $f(G) = g(G)$ since $h(G) = 0$
- $f(G_2) > f(G)$ from above

Optimality of A^* (proof)

- Suppose some suboptimal goal G_2 has been generated and is in the fringe. Let n be an unexpanded node in the fringe such that n is on a shortest path to an optimal goal G .



- $f(G_2) > f(G)$ from above
- $h(n) \leq h^*(n)$ since h is admissible
- $g(n) + h(n) \leq g(n) + h^*(n)$
- $f(n) \leq f(G)$

Hence $f(G_2) > f(n)$, and A^* will never select G_2 for expansion

Consistent heuristics


- Assume: A heuristic is **consistent** if for every node n , every successor n' of n generated by any action a ,

$$h(n) \leq c(n,a,n') + h(n')$$

- Proof: If h is consistent, we have


$$\begin{aligned} f(n') &= g(n') + h(n') \\ &= g(n) + c(n,a,n') + h(n') \\ &\geq g(n) + h(n) \\ &= f(n) \end{aligned}$$

- i.e., $f(n)$ is non-decreasing along any path.
- Theorem:** If $h(n)$ is consistent, A^* using GRAPH-SEARCH is optimal



Optimality of A^*

- A^* expands nodes in order of increasing f value
Gradually adds " f -contours" of nodes
- Contour i has all nodes with $f=f_i$, where $f_i < f_{i+1}$



Properties of A*

- Complete? Yes (unless there are infinitely many nodes with $f \leq f(G)$)
- Time? Exponential
- Space? Keeps all nodes in memory
- Optimal? Yes

Admissible heuristics

E.g., for the 8-puzzle:

- $h_1(n)$ = number of misplaced tiles
- $h_2(n)$ = total Manhattan distance
(i.e., no. of squares from desired location of each tile)

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- $h_1(S) = ?$
- $h_2(S) = ?$

Admissible heuristics

E.g., for the 8-puzzle:

- $h_1(n)$ = number of misplaced tiles
- $h_2(n)$ = total Manhattan distance

(i.e., no. of squares from desired location of each tile)

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- $h_1(S) = ?$ 8
- $h_2(S) = ?$ $3+1+2+2+2+3+3+2 = 18$

Dominance

- If $h_2(n) \geq h_1(n)$ for all n (both admissible)
then h_2 **dominates** h_1
 h_2 is better for search
- Typical search costs (average number of nodes expanded):
- $d=12$ IDS = 3,644,035 nodes
 $A^*(h_1) = 227$ nodes
 $A^*(h_2) = 73$ nodes
- $d=24$ IDS = too many nodes
 $A^*(h_1) = 39,135$ nodes
 $A^*(h_2) = 1,641$ nodes

Relaxed problems

- A problem with fewer restrictions on the actions is called a **relaxed problem**
- The cost of an optimal solution to a relaxed problem is an admissible heuristic for the original problem
- If the rules of the 8-puzzle are relaxed so that a tile can move **anywhere**, then $h_1(n)$ gives the shortest solution
- If the rules are relaxed so that a tile can move to **any adjacent square**, then $h_2(n)$ gives the shortest solution

Summary

- Greedy Best-first Search
- A* search
- Properties of search
 - Complete
 - Optimal
 - Local or global
- Heuristic $h(n)$
 - Admissible
 - Consistent

Assignment

– Chap 3: exercise 3.14, 3.21, 3.23

*Handed in next Tuesday

To be continued

Informed search algorithms

- Greedy Best-first Search
- A* search
 - Heuristic $h(n)$
- Questions and Answers
 - Algorithms

Lecture 6: Informed search algorithms(2)

Chapter 4

Outline

- Best-first search
- Greedy best-first search
- A^* search
- Heuristics
- Local search algorithms
- Hill-climbing search
- Simulated annealing search
- Local beam search
- Genetic algorithms
- Summary

Local search algorithms

- In many optimization problems, the **path** to the goal is irrelevant; the goal state itself is the solution
State space = set of "complete" configurations
- Find configuration satisfying constraints
- **local search algorithms**
 - keep a single "current" state, try to improve it
 - *Gradient Descent algorithm*




Fig. A one-dimensional state-space landscape

Different strategies for hill-climbing search:

- Gradient Descent algorithm
- Simulated Annealing algorithm
- GA

Example: n -queens

- Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal



Hill-climbing search


- "Like climbing Everest in thick fog with amnesia"

```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  inputs: problem, a problem
  local variables: current, a node
                  neighbor, a node

  current ← MAKE-NODE(INITIAL-STATE[problem])
  loop do
    neighbor ← a highest-valued successor of current
    if VALUE[neighbor] ≤ VALUE[current] then return STATE[current]
    current ← neighbor
```

Hill-climbing search

- Problem: depending on initial state, can get stuck in local maxima




Hill-climbing search: 8-queens problem

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♚	13	16	13	16
♚	14	17	15	♚	14	16	16
17	♚	16	18	15	♚	15	♚
18	14	♚	15	15	14	♚	16
14	14	13	17	12	14	12	18

- h = number of pairs of queens that are attacking each other, either directly or indirectly
- $h = ?$ for the above state **17**

Hill-climbing search: 8-queens problem



- A local minimum with $h = 1$

Simulated annealing search

- Idea: escape local maxima by allowing some "bad" moves but **gradually decrease** their frequency

function SIMULATED-ANNEALING(*problem*, *schedule*) **returns** a solution state

inputs: *problem*, a problem

schedule, a mapping from time to "temperature"

local variables: *current*, a node

next, a node

T, a "temperature" controlling prob. of downward steps

current \leftarrow MAKE-NODE(INITIAL-STATE[*problem*])

for *t* \leftarrow 1 **to** ∞ **do**

T \leftarrow *schedule*[*t*]

if *T* = 0 **then return** *current*

next \leftarrow a randomly selected successor of *current*

$\Delta E \leftarrow$ VALUE[*next*] - VALUE[*current*]

if $\Delta E > 0$ **then** *current* \leftarrow *next*

else *current* \leftarrow *next* only with probability $e^{\Delta E/T}$

Properties of simulated annealing search

- Annealing
 - is actually a term used in the metallurgical industry
 - probably by heating the crystal to a very high temperature and then slowly cooling down, reducing defects in the crystal (reaching the most stable state with the lowest energy)
- One can prove: If T decreases slowly enough, then simulated annealing search will find a global optimum with probability approaching 1
- Questions:
 - What will be if T remain a big value?
 - What will be if T decrease extremely quickly?

Local beam search

- Keep track of k states rather than just one
- Start with k randomly generated states
- At each iteration, all the successors of all k states are generated
- If any one is a goal state, stop; else select the k best successors from the complete list and repeat.

Genetic algorithms

- Ideas

- A successor state is generated by combining two parent states

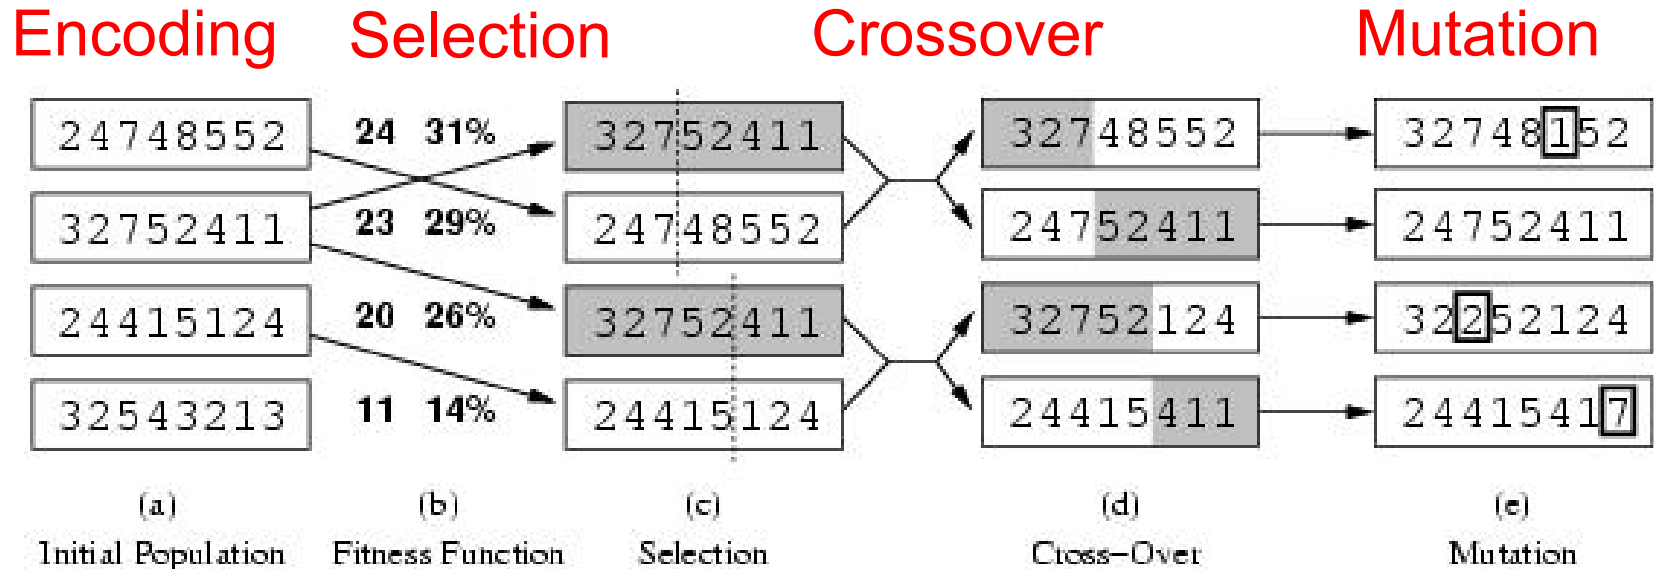
Start with k randomly generated states (**population**)

A state is represented as a string over a finite alphabet (often a string of 0s and 1s) --**encoding**

Evaluation function (**fitness function**). Higher values for better states.


Produce the next generation of states by **selection**, **crossover**, and **mutation**

Genetic algorithms




- Encoding: 8-bit string
- Fitness function: number of non-attacking pairs of queens (min = 0, max = $8 \times 7/2 = 28$)
- $24/(24+23+20+11) = 31\%$
- $23/(24+23+20+11) = 29\%$ etc

Genetic algorithms



“Grand” evolutionary scheme

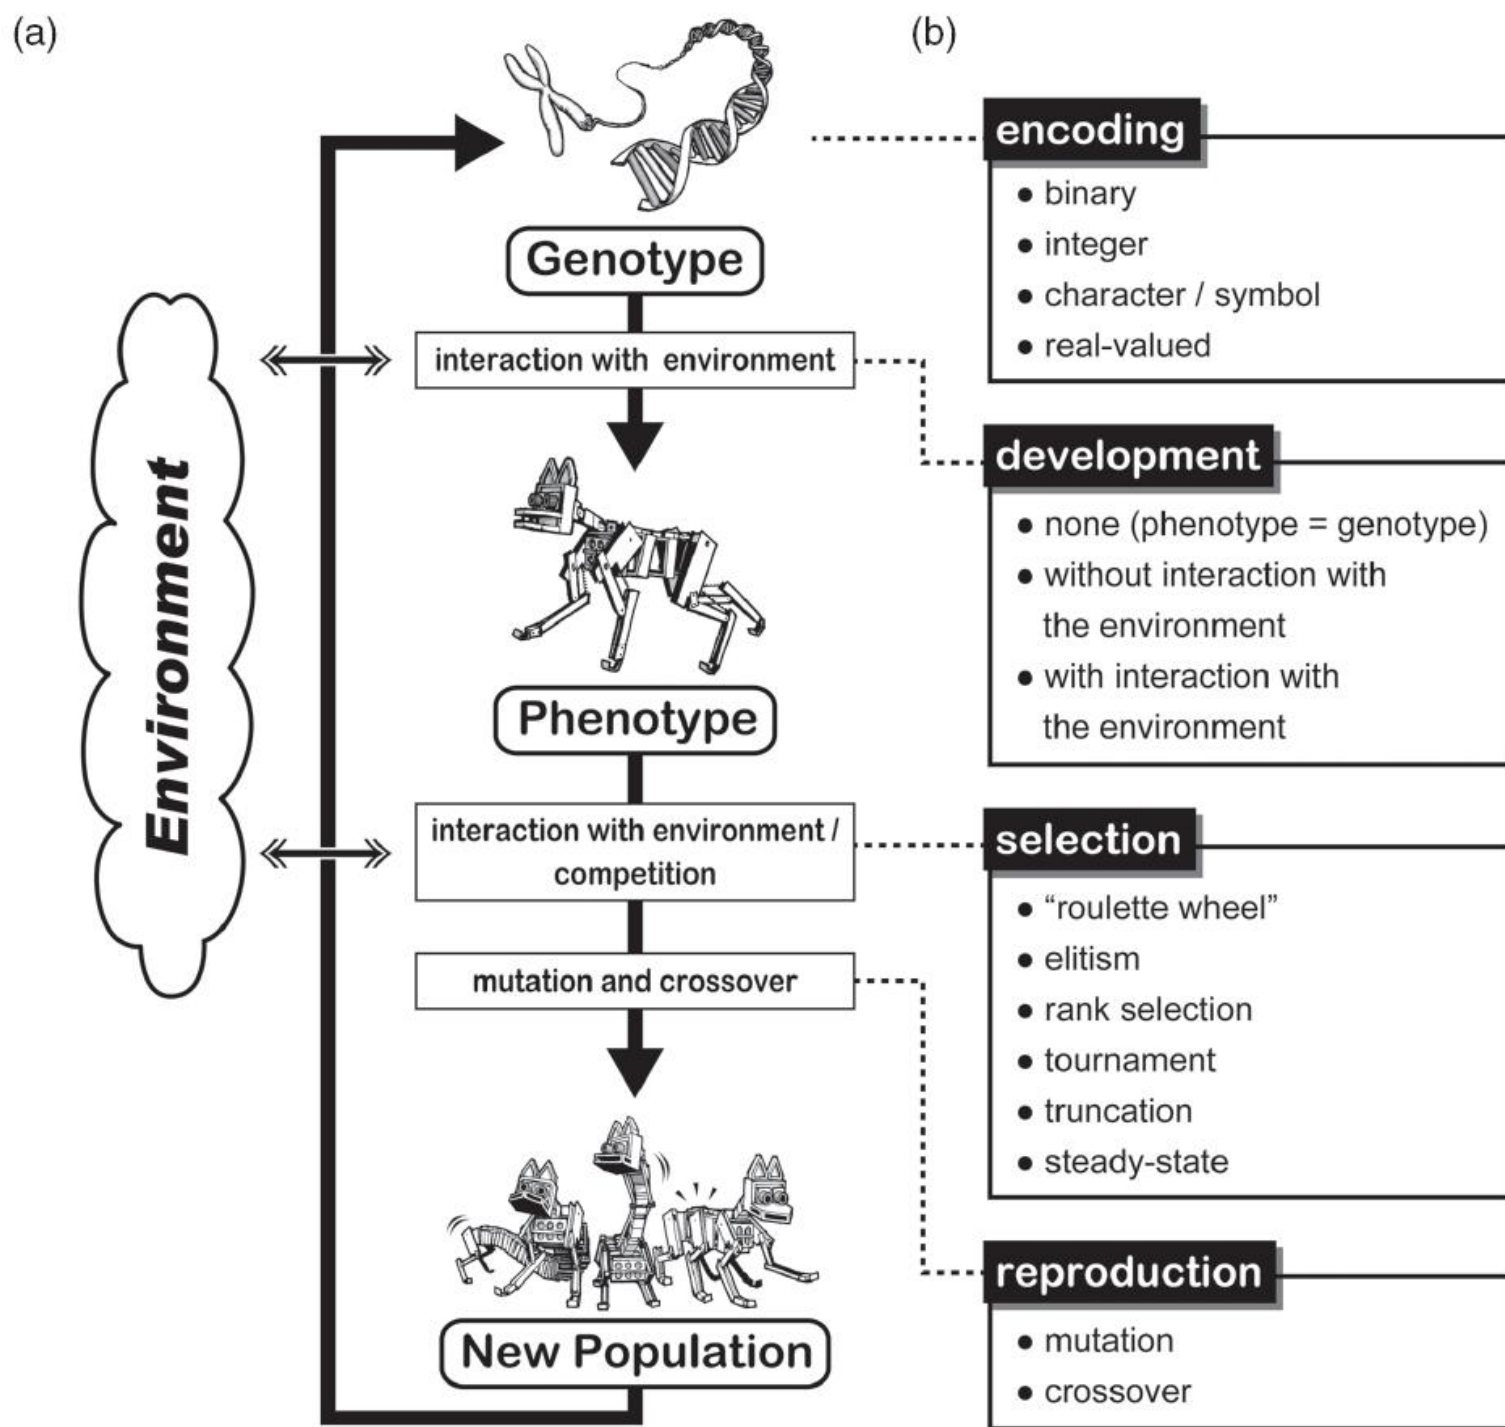


•Q & A :

Exploration
vs
Exploitation



encoding	development	selection	reproduction
<ul style="list-style-type: none"> • binary • many-character • real-valued 	<ul style="list-style-type: none"> • no development (phenotype = genotype) • development with and without interaction with the environment 	<ul style="list-style-type: none"> • “roulette wheel” • elitism • rank selection • tournament • truncation • steady-state 	<ul style="list-style-type: none"> • mutation • crossover

Cycle for artificial evolution




from:
“How the
body...”

Evolving a neural controller for an agent



how to proceed?

Encoding in genome



b.

initial genome	1101	0110	0001	0011	1010	1100
encodes weights (numbers)	1	2	3	4	5	6

c.

initial weights after "development"	.37	-.1	-.43	-.3	.16	.3
--	-----	-----	------	-----	-----	----


Encoding in genome “development”

1. take the one single individual with the highest fitness
2. choose another individual from the population at random, irrespective of fitness, for sexual reproduction
3. add the fittest individual to the new population


	fittest individual (highest rank)						other individual					
	1	2	3	4	5	6	1	2	3	4	5	6
initial genome	0011	0101	0010	0111	1010	1010	1010	1011	0000	1001	1101	1100
encoded weights	-.3	-.17	-.37	.03	.17	.17	.17	.23	-.5	.1	.37	.3

encoding	development	selection	reproduction
<ul style="list-style-type: none"> • binary • many-character • real-valued 	<ul style="list-style-type: none"> • no development (phenotype = genotype) • development with and without interaction with the environment 	<ul style="list-style-type: none"> • “roulette wheel” • elitism • rank selection • tournament • truncation • steady-state 	<ul style="list-style-type: none"> • mutation • crossover

Reproduction: Crossover and mutation “development”



“Grand” evolutionary scheme



encoding	development	selection	reproduction
<ul style="list-style-type: none"> • binary • many-character • real-valued 	<ul style="list-style-type: none"> • no development (phenotype = genotype) • development with and without interaction with the environment 	<ul style="list-style-type: none"> • “roulette wheel” • elitism • rank selection • tournament • truncation • steady-state 	<ul style="list-style-type: none"> • mutation • crossover

Summary

- Informed Search Algorithms
- Beyond classical search algorithms
 - Local search algorithms
 - Simulated annealing search
 - Genetic algorithms
- Still challenging
 - Neural Networks (NN, ANN)
 - EA (GA, EP, ES, GP)
 - Deep Learning
 -

Assignment

- Readings: Chap 4
- Chap 4: exercise 4.1

*Handed in next Tuesday