Operating System

Chapter 3: Processes

Chapter 3: Processes

- Process Concept
- Process Scheduling
- Operations on Processes
- Interprocess Communication
- Examples of IPC Systems
- Communication in Client-Server Systems

Objectives

- To introduce the notion of a process
- To describe the various features of processes
- To explore inter-process communication

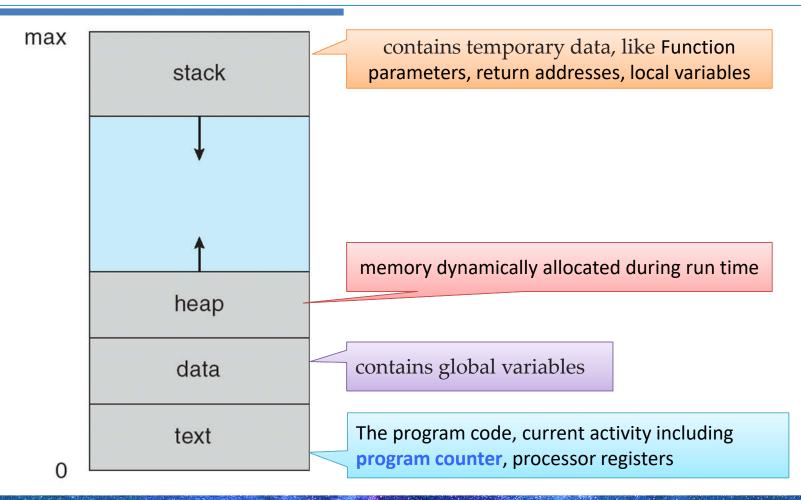
Process Concept



Process Definition

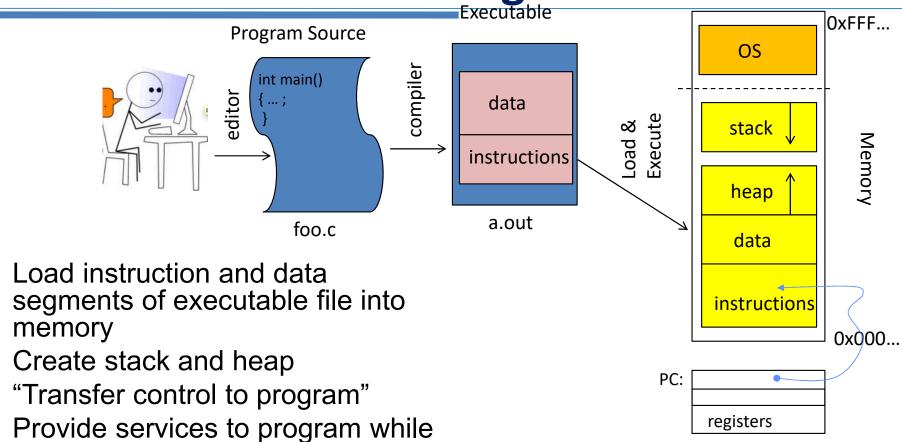
- Process
 - A program in execution; process execution must progress in sequential fashion
 - In time-sharing sys: unit of work
 - All processes are executed concurrently
- Process vs. Job?
 - Passive: program
 - Active: process
 - Program becomes process when executable file loaded into memory
 - One program can be several processes
 - Question?
 - java program

Process in memory



OS Bottom Line: Run Programs

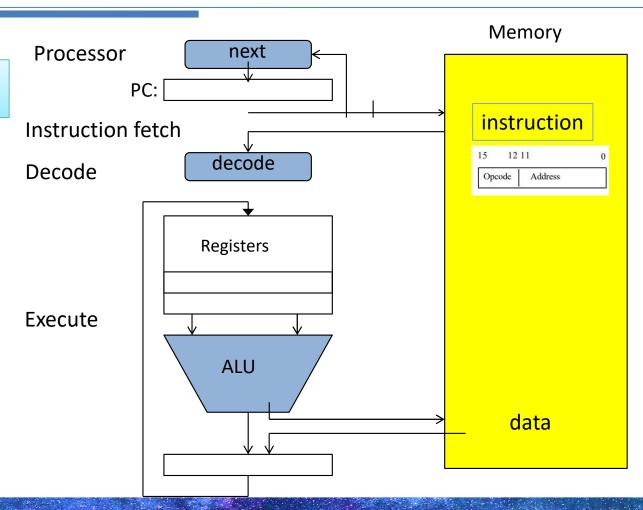
protecting OS and program



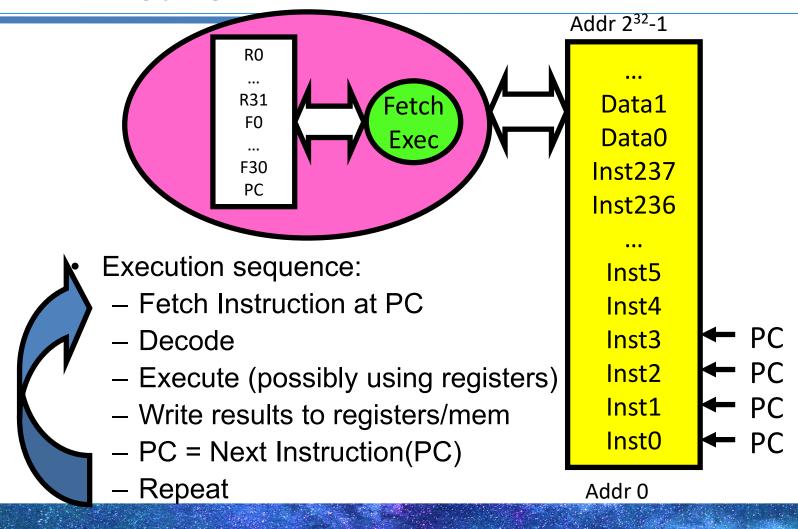
Processor

Instruction Fetch/Decode/Execute

The instruction cycle

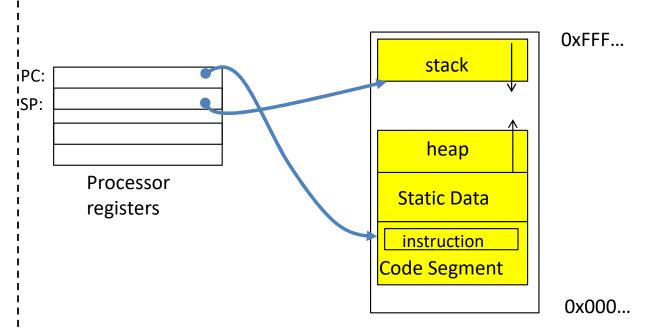


What happens during program execution?

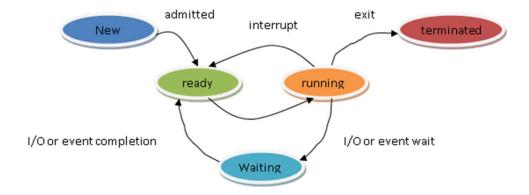


Address Space: In a Picture

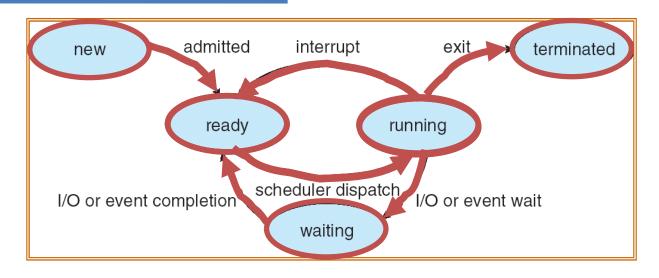
- What's in the code segment?
 Static data segment?
- What's in the Stack Segment?
 - How is it allocated? How big is it?
- What's in the Heap Segment?
 - How is it allocated? How big?



Process State



Lifecycle of a Process



- As a process executes, it changes state:
 - new: The process is being created
 - ready: The process is waiting to run
 - running: Instructions are being executed
 - waiting: Process waiting for some event to occur
 - terminated: The process has finished execution

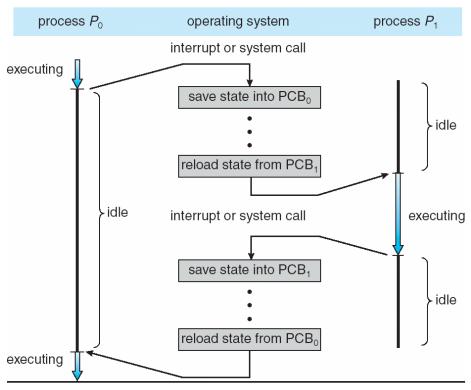
Process Control Block (PCB)

- How to manage processes?
- Information associated with each process
 (also task control block)
 - Process state
 - Program counter
 - CPU registers contents of all process-centric registers
 - CPU scheduling info. priorities, scheduling queue pointers
 - Memory-management info. memory allocated to the process
 - Accounting info. CPU used, clock time elapsed since start, time limits
 - I/O status info. I/O devices allocated to process, list of open files

process state
process number
program counter
registers
memory limits
list of open files

Context Switch

• Switching the CPU to another process requires performing a state save of the current process and a state restore of a different process.

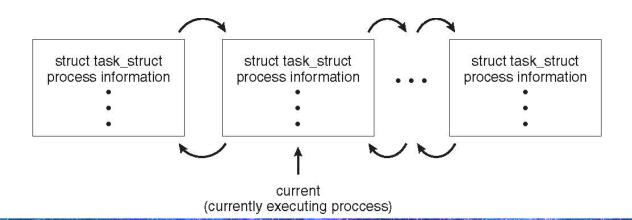


- The context is represented in the PCB of the process
- Context-switch time is pure overhead

Process Representation in Linux

Represented by the C structure task struct

```
pid t_pid; /* process identifier */
long state; /* state of the process */
unsigned int time_slice /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm; /* address space of this process */
```



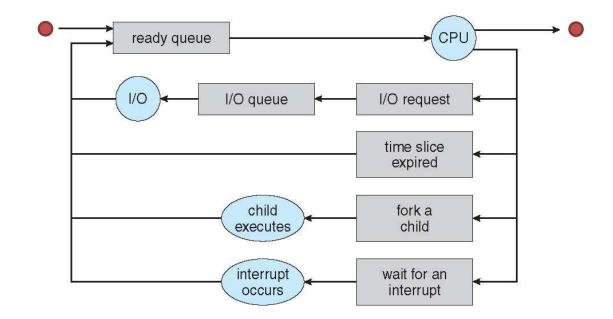
Scheduler

```
if ( readyProcesses(PCBs) ) {
    nextPCB = selectProcess(PCBs);
    run( nextPCB );
} else {
    run_idle_process();
}
```

- Scheduling: Mechanism for deciding which processes receive the CPU
- Lots of different scheduling policies provide ...
 - Fairness or
 - Realtime guarantees or
 - Latency optimization or ...

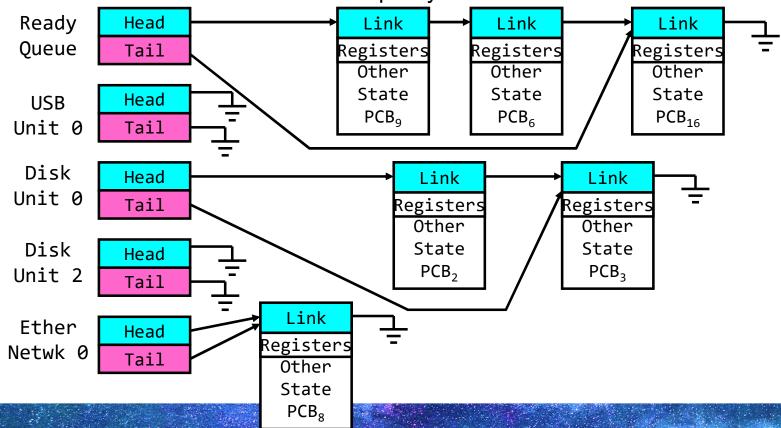
Diagram representation of process scheduling

Queueing diagram represents queues, resources, flows



Ready Queue And Various I/O Device Queues

- Process not running ⇒ PCB is in some scheduler queue
 - Separate queue for each device/signal/condition
 - Each queue can have a different scheduler policy



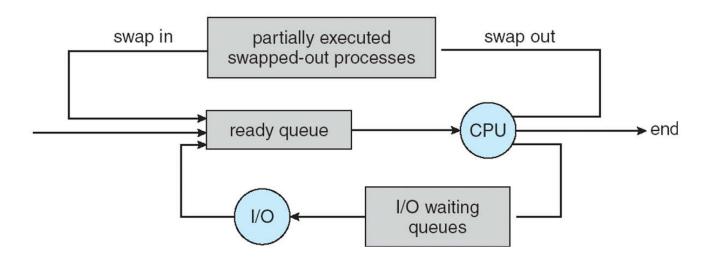
Schedulers

- Short-term scheduler (or CPU scheduler)
 - selects which process should be executed next and allocates CPU
 - Sometimes the only scheduler in a system
 - Short-term scheduler is invoked frequently (milliseconds) ⇒ (must be fast)
- Long-term scheduler (or job scheduler)
 - selects which processes should be brought into the ready queue
 - Long-term scheduler is invoked infrequently (seconds, minutes) ⇒ (may be slow)
 - The long-term scheduler controls the degree of multiprogramming
- Processes:
 - I/O-bound
 - spends more time doing I/O than computations, many short CPU bursts
 - CPU-bound
 - spends more time doing computations; few very long CPU bursts
- Long-term scheduler strives for good process mix

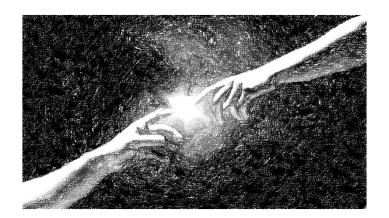
Example of standard API

Medium-term scheduler

- Can be added if degree of multiple programming needs to decrease
- Remove process from memory, store on disk, bring back in from disk to continue execution:
 swapping



Process Creation



Can a process create a process?

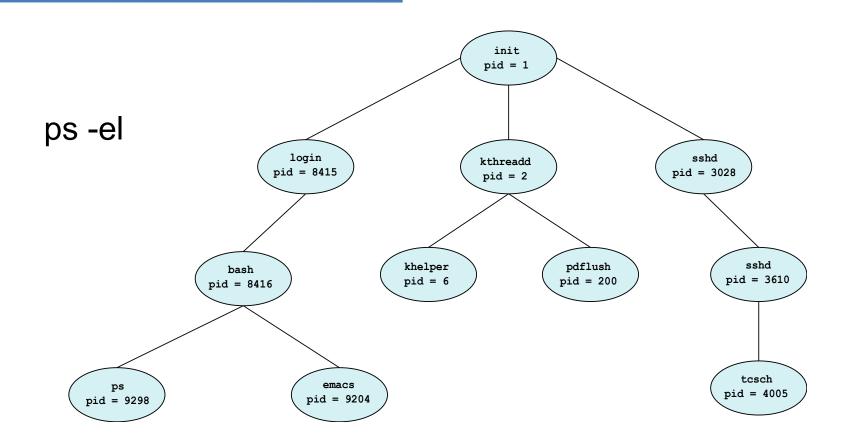
- Yes! Unique identity of process is the "process ID" (or PID)
- fork() system call creates a copy of current process with a new PID
- Return value from fork(): integer
 - When > 0:
 - Running in (original) Parent process
 - return value is pid of new child
 - When = 0:
 - Running in new Child process
 - When < 0:
 - Error! Must handle somehow
 - Running in original process

Process creation

- Parent vs. Child
- Generally, process identified and managed via a process identifier (pid)
- Resource sharing options
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Execution options
 - Parent and children execute concurrently
 - Parent waits until children terminate

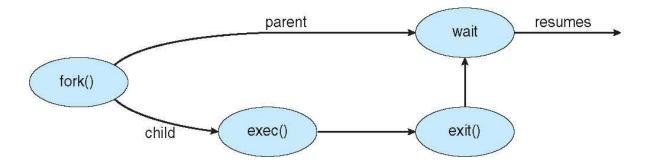
State of original process duplicated in *both* Parent and Child! Memory, File Descriptors (next topic), etc...

A Tree of Processes in Linux



Process Creation (Cont.)

- Address space
 - Child duplicate of parent
 - Child has a program loaded into it
- UNIX examples
 - fork () system call creates new process
 - exec() system call used after a fork() to replace the process' memory space with a new program



Process creation with C

```
#include <sys/types.h>
#include <stdio.h>
                                          POSIX
#include <unistd.h>
int main()
pid_t pid;
   /* fork a child process */
   pid = fork();
   if (pid < 0) { /* error occurred */
      fprintf(stderr, "Fork Failed");
      return 1;
   else if (pid == 0) { /* child process */
      execlp("/bin/ls","ls",NULL);
   else { /* parent process */
      /* parent will wait for the child to complete */
      wait(NULL);
      printf("Child Complete");
   return 0;
```

Why pid has two values when fork() return?

EX 1

```
Parent

main()
{
    fork();
    pid = ...;
}
```

```
This line is from pid 95978, value = 1
This line is from pid 95978, value = 2
This line is from pid 95978, value = 3
This line is from pid 95979, value = 1
This line is from pid 95978, value = 4
This line is from pid 95979, value = 2
This line is from pid 95978, value = 5
This line is from pid 95979, value = 3
This line is from pid 95979, value = 6
This line is from pid 95979, value = 4
This line is from pid 95979, value = 7
This line is from pid 95979, value = 7
```

```
#include <stdio.h>
#include <string.h>
#include <sys/types.h>

#define MAX_COUNT 200
#define BUF_SIZE 50

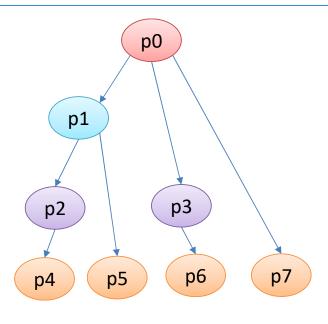
void main(void){

   pid_t pid;
   int i;
   char buf[BUF_SIZE];

   fork();
   pid = getpid();
   for(i=1;i<=MAX_COUNT;i++){
        sprintf(buf, "This line is from pid %d, value = %d\n", pid, i);
        write(1, buf, strlen(buf));
   }
}</pre>
```

Question?

```
#include <stdio.h>
#include <unistd.h>
int main()
   /* fork a child process */
   fork();
   /* fork another child process */
   fork();
   /* and fork another */
   fork();
   return 0;
```



How many processes are created?

Ex 2

```
main() pid = 3456
{
    pid=fork();
    if (pid == 0)
        ChildProcess();
    else
        ParentProcess();
}

void ChildProcess()
{
    ....
}

void ParentProcess()
{
    ....
}
```

```
main() pid = 0
{
    pid=fork();
    if (pid == 0)
        ChildProcess();
    else
        ParentProcess();
}

void ChildProcess()
{
    .....
}

void ParentProcess()
{
    .....
}
```

```
void main(void)
    pid_t pid;
    pid = fork();
    if (pid == 0)
         ChildProcess();
    else
         ParentProcess();
void ChildProcess(void)
    int i;
    int mypid;
    mypid = getpid();
    for (i = 1; i \leftarrow MAX COUNT; i++)
         printf(" This line is from child, value = %d\n", i);
    printf(" *** Child process %d is done ***\n",mypid);
    exit(0);
void ParentProcess(void)
    int i, status;
    int got_pid,mypid;
    mypid = getpid();
    for (i = 1; i \leftarrow MAX COUNT; i++)
          printf("This line is from parent, value = %d\n", i);
    printf("*** Parent %d is done ***\n",mypid);
    got pid = wait(&status);
    printf("[%d] bye %d (%d)\n", mypid, got_pid, status);
```

EX2 (cont.)

```
Child
      Parent
main()
                          main()
                                   pid = 0
        pid = 3456
                                                                   This line is from parent, value = 1
                           pid=fork();
 pid=fork();
                                                                   This line is from parent, value = 2
  if (pid == 0)
                             if (pid == 0)
    ChildProcess():
                               ChildProcess();
                                                                    *** Parent 118633 is done ***
    ParentProcess();
                               ParentProcess();
                                                                       This line is from child, value = 1
                                                                       This line is from child, value = 2
                          void ChildProcess()
void ChildProcess()
                                                                       *** Child process 118634 is done ***
                                                                    [118633] bye 118634 (0)
                          void ParentProcess()
void ParentProcess()
                                                                            Child
                                                Parent
                                         main()
                                                  pid = 3456
                                                                     main()
                                                                              pid = 0
                                            pid=fork();
                                                                       pid=fork();
                                            if (pid == 0)
                                                                       if (pid == 0)
                                              ChildProcess();
                                                                          ChildProcess();
                                              ParentProcess();
                                                                          ParentProcess();
                                         void ChildProcess()
                                                                    void ChildProcess()
                                        void ParentProcess()
                                                                     void ParentProcess()
                                            ....
                                                              Operating Systems Fall 2020
```

Process termination

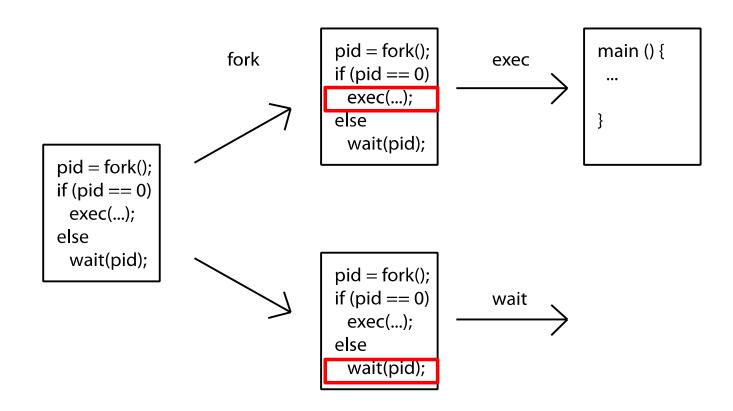
- Child → Parent
 - Process' resources are deallocated when:
 - exit(n)
 - return() in main()
 - Catch exit status → wait()
 - -pid = wait(&status);
- Parent → Child
 - terminate child
 - Why?
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - The parent is exiting and the operating systems does not allow a child to continue if its parent terminates

Problems of process termination

- zombie process
 - No parent waiting
- orphan process
 - Parent termination without wait
- Multi process example: Chrome Browser
 - Browser, Renderer, Plugins, etc



UNIX Process Management



Interprocess communication (IPC)



Interprocess communication (IPC)

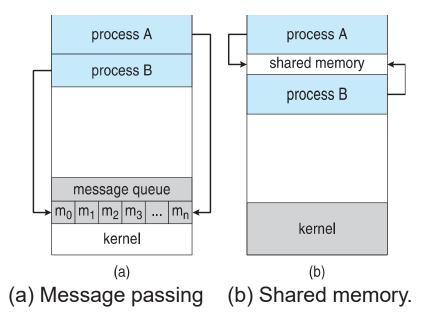
Process:

independent vs. cooperating

- Cooperating process:
 - ➤ Shared memory
 - Message passing

Reasons for process cooperation:

- ➤ Information sharing
- Computation speedup
- ➤ Modularity
- convenience



Circular buffer & producer-consumer problem

```
#define BUFFER_SIZE 16
typedef struct {
    . . .
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

```
item next_consumed;
while (true) {
          while (in == out) ; /* do
nothing */

          next_consumed = buffer[out];
          out = (out + 1) % BUFFER_SIZE;

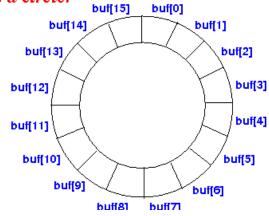
          /* consume the item in next
consumed */
}
```

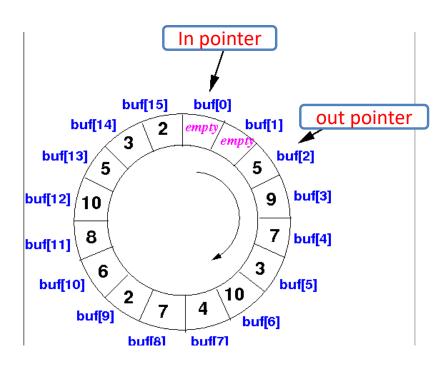
circular array

Array:



Pretend array is a circle:





POSIX examples of shared memory: (sender->receiver)

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
                                     Producer
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>
int main()
/* the size (in bytes) of shared memory object */
const int SIZE = 4096;
/* name of the shared memory object */
const char *name = "OS";
/* strings written to shared memory */
const char *message_0 = "Hello";
const char *message_1 = "World!";
/* shared memory file descriptor */
int shm fd;
/* pointer to shared memory obect */
void *ptr;
   /* create the shared memory object */
   shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666)
   /* configure the size of the shared memory object */
   ftruncate(shm_fd, SIZE);
   /* memory map the shared memory object */
   ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);
   /* write to the shared memory object */
   sprintf(ptr, "%s", message_0);
   ptr += strlen(message_0);
   sprintf(ptr,"%s",message_1);
   ptr += strlen(message_1);
   return 0;
```

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
                                        Consumer
#include <sys/shm.h>
#include <sys/stat.h>
int main()
/* the size (in bytes) of shared memory object */
const int SIZE = 4096:
/* name of the shared memory object */
const char *name = "OS";
/* shared memory file descriptor */
int shm_fd:
/* pointer to shared memory obect */
void *ptr;
   /* open the shared memory object */
   shm_fd = shm_open(name, O_RDONLY, 0666);
   /* memory map the shared memory object */
   ptr = mmap(0, SIZE, PROT READ, MAP SHARED, shm fd, 0);
   /* read from the shared memory object */
   printf("%s",(char *)ptr);
   /* remove the shared memory object */
   shm_unlink(name);
   return 0;
```

Interprocess Communication – Message Passing

- Mechanism for processes to communicate and to synchronize their actions
- IPC facility provides two operations:
 - send(message) message size fixed or variable
 - receive(message)
- If P and Q wish to communicate, they need to:
 - establish a communication link between them
 - exchange messages via send/receive
- Implementation of communication link
 - physical (e.g., shared memory, hardware bus)
 - logical (e.g., direct or indirect, synchronous or asynchronous, automatic or explicit buffering)

Message Passing (Cont.)

- If processes *P* and *Q* wish to communicate, they need to:
 - Establish a communication link between them
 - Exchange messages via send/receive
- Implementation issues:
 - How are links established?
 - Can a link be associated with more than two processes?
 - > How many links can there be between every pair of communicating processes?
 - What is the capacity of a link?
 - > Is the size of a message that the link can accommodate fixed or variable?
 - Is a link unidirectional or bi-directional?

Message passing

- Direct communication (unidirectional)
 - send (P, message) send a message to process P
 - receive(Q, message) receive a message from process Q
- Indirect communication (uni & bidirectional)
 - Messages are directed and received from mailboxes (or ports)
 - Can be used by multiple processes
 - Primitives are defined as:
 - send(A, message) send a message to mailbox A
 - receive (A, message) receive a message from mailbox A

Direct Communication

- Processes must name each other explicitly:
 - send (P, message) send a message to process P
 - receive(Q, message) receive a message from process Q
- Properties of communication link
 - Links are established automatically
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
 - The link may be unidirectional, but is usually bi-directional

Indirect Communication

- Messages are directed and received from mailboxes (also referred to as ports)
 - Each mailbox has a unique id
 - Processes can communicate only if they share a mailbox
- Properties of communication link
 - ➤ Link established only if processes share a common mailbox
 - > A link may be associated with many processes
 - ➤ Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional

Indirect Communication

- Operations
 - create a new mailbox (port)
 - send and receive messages through mailbox
 - destroy a mailbox
- Primitives are defined as:

send(A, message) - send a message to mailbox A
receive(A, message) - receive a message from mailbox A

Indirect Communication

- Mailbox sharing
 - $-P_1$, P_2 , and P_3 share mailbox A
 - $-P_1$, sends; P_2 and P_3 receive
 - Who gets the message?
- Solutions
 - Allow a link to be associated with at most two processes
 - Allow only one process at a time to execute a receive operation
 - Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.

Synchronization

- Blocking vs. non-blocking
- Blocking is considered synchronous
 - Blocking send
 - Blocking receive
- Non-blocking is considered asynchronous
 - Non-blocking send
 - Non-blocking receive
 - The receiver receives
 - √ A valid message
 - √ Null message
- Different combinations possible
 - If both send and receive are blocking, we have a rendezvous

Synchronization (Cont.)

Producer-consumer becomes trivial

```
message next_produced;
while (true) {
   /* produce an item in next produced */
   send(next_produced);
}
```

```
message next_consumed;
while (true) {
   receive(next_consumed);

   /* consume the item in next consumed */
}
```

Buffering

- Queue of messages attached to the link.
- implemented in one of three ways
 - 1. Zero capacity no messages are queued on a link.
 - Sender must wait for receiver (rendezvous)
 - 2. Bounded capacity finite length of *n* messages
 - Sender must wait if link full
 - 3. Unbounded capacity infinite length
 - Sender never waits

Examples of IPC Systems - Mach

- Mach communication is message based
 - Even system calls are messages
 - Each task gets two mailboxes at creation- Kernel and Notify
 - Only three system calls needed for message transfer
 msg_send(), msg_receive(), msg_rpc()
 - Mailboxes needed for commuication, created via port_allocate()
 - Send and receive are flexible, for example four options if mailbox full:
 - Wait indefinitely
 - Wait at most n milliseconds
 - Return immediately
 - Temporarily cache a message

Pipes

- Acts as a conduit allowing two processes to communicate
- Issues:
 - Is communication unidirectional or bidirectional?
 - In the case of two-way communication, is it half or full-duplex?
 - Must there exist a relationship (i.e., *parent-child*) between the communicating processes?
 - Can the pipes be used over a network?

Ordinary pipes

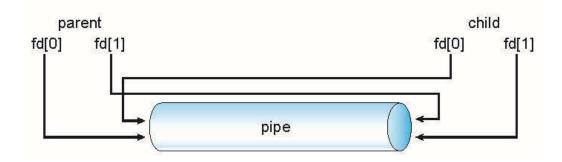
cannot be accessed from outside the process that created it. Typically, a parent process creates a
pipe and uses it to communicate with a child process that it created.

Named pipes

can be accessed without a parent-child relationship.

Ordinary Pipes

- Ordinary Pipes allow communication in standard producer-consumer style
- Producer writes to one end (the write-end of the pipe)
- Consumer reads from the other end (the read-end of the pipe)
- Ordinary pipes are therefore unidirectional
- Require parent-child relationship between communicating processes



■ Windows calls these anonymous pipes

Ordinary pipe (POSIX), parent-child

```
#include <sys/types.h>
#include <stdio.h>
#include <string.h>
#include <unistd.h>
#define BUFFER_SIZE 25
#define READ END 0
#define WRITE_END 1
int main(void)
char write_msg[BUFFER_SIZE] = "Greetings";
char read_msg[BUFFER_SIZE];
int fd[2];
pid_t pid;
    /* create the pipe */
   if pipe(fd) == -1) {
      fprintf(stderr, "Pipe failed");
      return 1;
   /* fork a child process */
   pid = fork();
   if (pid < 0) { /* error occurred */
      fprintf(stderr, "Fork Failed");
      return 1;
```

```
if (pid > 0) { /* parent process */
  /* close the unused end of the pipe */
  close(fd[READ_END]);
  write(fd[WRITE_END], write_msg, strlen(write_msg)+1);
  /* close the write end of the pipe */
  close(fd[WRITE_END]);
else { /* child process */
  /* close the unused end of the pipe */
  close(fd[WRITE_END]);
  /* read from the pipe */
  read(fd[READ_END], read_msg, BUFFER_SIZE);
  printf("read %s",read_msg);
  /* close the write end of the pipe */
  close(fd[READ_END]);
return 0;
```

Named pipes

- Named Pipes are more powerful than ordinary pipes (?)
- Communication is bidirectional
- No parent-child relationship is necessary between the communicating processes
- Several processes can use the named pipe for communication
- Provided on both UNIX and Windows systems

Creating a FIFO file:

- In order to create a FIFO file, a function calls i.e. mkfifo is used.

int mkfifo(const char *pathname, mode_t mode);

mkfifo() makes a FIFO special file with name *pathname*. Here *mode* specifies the FIFO's permissions.

Using FIFO: As named pipe(FIFO) is a kind of file, we can use all the system calls associated with it i.e. *open, read, write, close*.

EX: Named pipe

```
int main()
   int fd;
   char * myfifo = "/tmp/myfifo"; // FIFO file path
   // Creating the named file(FIFO)
   // mkfifo(<pathname>, <permission>)
   mkfifo(myfifo, 0666);
   char arr1[80], arr2[80];
   while (1)
       // Open FIFO for write only
       fd = open(myfifo, O_WRONLY);
       // Take an input arr2ing from user.
       // 80 is maximum length
       fgets(arr2, 80, stdin);
       // Write the input arr2ing on FIFO
       // and close it
       write(fd, arr2, strlen(arr2)+1);
       close(fd);
       // Open FIFO for Read only
       fd = open(myfifo, O_RDONLY);
       // Read from FIFO
       read(fd, arr1, sizeof(arr1));
       // Print the read message
       printf("User2: %s\n", arr1);
       close(fd);
   return 0;
```

```
#include <stdio.h>
#include <string.h>
#include <fcntl.h>
#include <sys/stat.h>
#include <sys/types.h>
#include <unistd.h>
int main()
    int fd1;
    char * myfifo = "/tmp/myfifo"; // FIFO file path
    // Creating the named file(FIFO)
    // mkfifo(<pathname>,<permission>)
    mkfifo(myfifo, 0666);
    char str1[80], str2[80];
    while (1)
        // First open in read only and read
        fd1 = open(myfifo,O RDONLY);
        read(fd1, str1, 80);
        // Print the read string and close
        printf("User1: %s\n", str1);
        close(fd1);
        // Now open in write mode and write
        // string taken from user.
        fd1 = open(myfifo,O_WRONLY);
        fgets(str2, 80, stdin);
        write(fd1, str2, strlen(str2)+1);
        close(fd1);
    return 0;
```

Communications in client-server systems

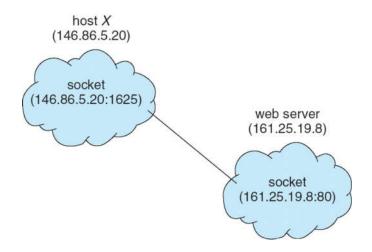
- Sockets
- Remote Procedure Calls (windows)
- Pipes
- Remote Method Invocation (Java)

Sockets

- A socket is defined as an endpoint for communication
- Concatenation of IP address and port a number included at start of message packet to differentiate network services on a host
- The socket 161.25.19.8:1625 refers to port 1625 on host 161.25.19.8
- Communication consists between a pair of sockets
- All ports below 1024 are well known, used for standard services
- Special IP address 127.0.0.1 (loopback) to refer to system on which process is running

Socket Communication

- Three types of sockets
 - Connection-oriented (TCP)
 - Connectionless (UDP)
 - MulticastSocket class- data can be sent to multiple recipients



Consider this "Date" server:

```
import java.net.*;
import java.io.*;
public class DateClient
  public static void main(String[] args) {
    try {
       /* make connection to server socket */
       Socket sock = new Socket("127.0.0.1",6013);
       InputStream in = sock.getInputStream();
       BufferedReader bin = new
          BufferedReader(new InputStreamReader(in));
       /* read the date from the socket */
       String line;
       while ((line = bin.readLine()) != null)
          System.out.println(line);
       /* close the socket connection*/
       sock.close();
     catch (IOException ioe)
       System.err.println(ioe);
```

```
import java.net.*;
import java.io.*;
public class DateServer
  public static void main(String[] args) {
     try {
       ServerSocket sock = new ServerSocket(6013);
       /* now listen for connections */
       while (true) {
          Socket client = sock.accept();
         PrintWriter pout = new
           PrintWriter(client.getOutputStream(), true);
          /* write the Date to the socket */
         pout.println(new java.util.Date().toString());
          /* close the socket and resume */
          /* listening for connections */
          client.close();
     catch (IOException ioe) {
       System.err.println(ioe);
```

End of Chapter 3

