




# Operating System

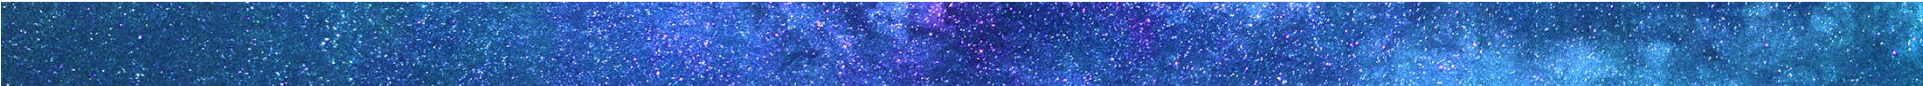
## Chapter 5 synchronization





# Objectives

---

- To introduce the critical-section problem, whose solutions can be used to ensure the consistency of shared data
  - To present both software and hardware solutions of the critical-section problem
  - To examine several classical process-synchronization problems
  - To explore several tools that are used to solve process synchronization problems
- 



# Motivation

---

- Cooperating process/thread:
  - the one that can affect or be affected by other processes executing in system.
  - Processes, threads
- Processes can execute concurrently
  - May be interrupted at any time, partially completing execution
- Problem: Data inconsistency
  - It may occur in the case of concurrent access to shared data
- How to solve?
  - Orderly execution of cooperating processes that share a logical address space

# Circular buffer & producer-consumer problem

```
#define BUFFER_SIZE 16
typedef struct {
    . . .
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

```
item next_produced;
while (true) {
    /* produce an item in next produced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```


```
item next_consumed;
while (true) {
    while (in == out) ; /* do
nothing */

    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;



    /* consume the item in next
consumed */
}
```

# circular array

Array:



*Pretend array is a circle:*








# One solution!

---

- A solution to **consumer-producer** problem that fills **all** the buffers.
- We can have an integer **counter** that keeps track of the **number of full buffers**.
- Initially, **counter** is **set to 0**.
- It is **incremented** by the **producer** after it produces a new buffer
- It is **decremented** by the **consumer** after it consumes a buffer.



# Circular buffer & producer-consumer problem

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;
item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
int counter = 0;
```

## Producer

```
item next_produced;
while (true) {
    /* produce an item in next produced */
    while (counter == BUFFER_SIZE) ;
    /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
    counter ++;
}
```

## Consumer

```
item next_consumed;
while (true) {
    while (counter == 0) ;
    /* do nothing */

    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    counter --;

    /* consume the item in next
consumed */
}
```

# Race condition

- **counter++** could be implemented as

<code>register1 = counter</code>	<code>MOV AX, [100]</code>
<code>register1 = register1 + 1</code>	<code>ADD AX, 1</code>
<code>counter = register1</code>	<code>MOV [100], AX</code>

- **counter--** could be implemented as

<code>register2 = counter</code>	<code>MOV BX, [100]</code>
<code>register2 = register2 - 1</code>	<code>ADD BX, 1</code>
<code>counter = register2</code>	<code>MOV [100], BX</code>

- Consider this execution interleaving with “count = 5” initially:

S0: producer execute <code>register1 = counter</code>	{register1 = 5}
S1: producer execute <code>register1 = register1 + 1</code>	{register1 = 6}
S2: consumer execute <code>register2 = counter</code>	{register2 = 5}
S3: consumer execute <code>register2 = register2 - 1</code>	{register2 = 4}
S4: producer execute <code>counter = register1</code>	{counter = 6}
S5: consumer execute <code>counter = register2</code>	{ <b>counter = 4</b> }



# Another Race condition



- Invoking *echo()* procedure:

```
void echo()  
{  
    chin = getchar();  
    chout = chin;  
    putchar(chout);  
}
```

Process P1	Process P2
•	•
chin = getchar();	•
•	chin = getchar();
chout = chin;	chout = chin;
putchar(chout);	•
•	putchar(chout);
•	•

- Same problem exists on:
  - Multiprogramming environment
  - Multiprocessing environment
  - Distributed processing environment

# Problem is at the Lowest Level

- Most of the time, threads are working on separate data, so scheduling doesn't matter:

Thread A

$x = 1;$

Thread B

$y = 2;$

- However, what about (Initially,  $y = 12$ ):

Thread A

$x = 1;$

$x = y + 1;$

Thread B

$y = 2;$

$y = y * 2;$

- What are the possible values of  $x$ ?

- Or, what are the possible values of  $x$  below?

Thread A

$x = 1;$

Thread B

$x = 2;$

- $X$  could be 1 or 2 (non-deterministic!)

# Atomic Operations

---

- To understand a concurrent program, we need to know what the underlying indivisible operations are!
- **Atomic Operation**: an operation that always runs to completion or not at all
  - It is *indivisible*: it cannot be stopped in the middle and state cannot be modified by someone else in the middle
  - Fundamental building block – if no atomic operations, then have no way for threads to work together
- On most machines, memory references and assignments (i.e. loads and stores) of words are atomic
- Many instructions are not atomic
  - Double-precision floating point store often not atomic
  - VAX and IBM 360 had an instruction to copy a whole array



# Another Concurrent Program Example

- Two threads, A and B, compete with each other
  - One tries to increment a shared counter
  - The other tries to decrement the counter

## Thread A

```
i = 0;  
while (i < 10)  
    i = i + 1; i = i - 1;  
printf("A wins!");
```

## Thread B

```
i = 0;  
while (i > -10)  
    printf("B wins!");
```


- Assume that **memory loads and stores are atomic**, but incrementing and decrementing are *not* atomic
- Who wins? Could be either
- Is it guaranteed that someone wins? Why or why not?
- What if both threads have their own CPU running at same speed? Is it guaranteed that it goes on forever?



# Other examples?

---

Have you ever seen other examples?



# Motivating Example: “Too Much Milk”

- Great thing about OS's – analogy between problems in OS and problems in real life
  - Help you understand real life problems better
  - But, computers are much stupider than people
- Example: People need to coordinate:



Time	Person A	Person B
3:00	Look in Fridge. Out of milk	
3:05	Leave for store	
3:10	Arrive at store	Look in Fridge. Out of milk
3:15	Buy milk	Leave for store
3:20	Arrive home, put milk away	Arrive at store
3:25		Buy milk
3:30		Arrive home, put milk away



# Definitions

---

- **Synchronization**: using atomic operations to ensure cooperation between threads
  - For now, only loads and stores are atomic
  - We are going to show that its hard to build anything useful with only reads and writes
- **Mutual Exclusion**: ensuring that only one thread does a particular thing at a time
  - One thread *excludes* the other while doing its task
- **Critical Section**: piece of code that only one thread can execute at once. Only one thread at a time will get into this section of code
  - Critical section is the result of mutual exclusion
  - Critical section and mutual exclusion are two ways of describing the same thing

# More Definitions

- **Lock**: prevents someone from doing something
  - Lock before entering critical section and before accessing shared data
  - Unlock when leaving, after accessing shared data
  - Wait if locked
    - Important idea: all synchronization involves waiting
- For example: fix the milk problem by putting a key on the refrigerator
  - Lock it and take key if you are going to go buy milk
  - Fixes too much: roommate angry if only wants OJ




- Of Course – We don't know how to make a lock yet



# Too Much Milk: Correctness Properties

---

- Need to be careful about correctness of concurrent programs, since non-deterministic
  - Instead, think first, then code
  - Always write down behavior first
- What are the correctness properties for the “Too much milk” problem???
- Never more than one person buys
- Someone buys if needed
- Restrict ourselves to use only **atomic load and store** operations as building blocks






# Too Much Milk: Solution #1

- Use **a note** to avoid buying too much milk:
  - Leave a note before buying (kind of “lock”)
  - Remove note after buying (kind of “unlock”)
  - Don’t buy if note (wait)
- Suppose a computer tries this (remember, only memory read/write are atomic):

```
if (noMilk) {  
    if (noNote) {  
        leave Note;  
        buy milk;  
        remove note;  
    }  
}
```



# Too Much Milk: Solution #1

- Use a note to avoid buying too much milk:
  - Leave a note before buying (kind of “lock”)
  - Remove note after buying (kind of “unlock”)
  - Don’t buy if note (wait)
- Suppose a computer tries this (remember, only memory read/write are atomic):

Thread A

```
if (noMilk) {  
  
    if (noNote) {  
        leave Note;  
    }  
    buy Milk;  
    remove Note;  
}
```


Thread B

```
if (noMilk) {  
    if (noNote) {  
  
        leave Note;  
        buy Milk;  
        remove Note;  
    }  
}
```

# Too Much Milk: Solution #1

- Use a note to avoid buying too much milk:
  - Leave a note before buying (kind of “lock”)
  - Remove note after buying (kind of “unlock”)
  - Don’t buy if note (wait)
- Suppose a computer tries this (remember, only memory read/write are atomic):

```
if (noMilk) {  
    if (noNote) {  
        leave Note;  
        buy milk;  
        remove note;  
    }  
}
```



- Result?
  - Still too much milk **but only occasionally!**
  - Thread can get context switched after checking milk and note but before buying milk!
- Solution makes problem worse since fails **intermittently**
  - Makes it really hard to debug...
  - Must work despite what the dispatcher does!



# Too Much Milk: Solution #1½

- Clearly the Note is not quite blocking enough
  - Let's try to fix this by placing note first
- Another try at previous solution:

```
leave Note;  
if (noMilk) {  
    if (noNote) {  
        buy milk;  
    }  
}  
remove Note;
```



- What happens here?
  - Well, with human, probably nothing bad
  - With computer: no one ever buys milk

## Too Much Milk Solution #2


- How about labeled notes?
  - Now we can leave note before checking
- Algorithm looks like this:

```
Thread A
leave note A;
if (noNote B) {
    if (noMilk) {
        buy Milk;
    }
}
remove note A;
```

```
Thread B
leave note B;
if (noNoteA) {
    if (noMilk) {
        buy Milk;
    }
}
remove note B;
```

- Does this work?
- Possible for neither thread to buy milk
  - Context switches at exactly the wrong times can lead each to think that the other is going to buy
- Really insidious:
  - **Extremely unlikely** this would happen, but will at worse possible time
  - Probably something like this in UNIX

# Too Much Milk Solution #2: problem!



- *I'm* not getting milk, *You're* getting milk
- This kind of lockup is called “starvation!”



# Too Much Milk Solution #3

- Here is a possible two-note solution:

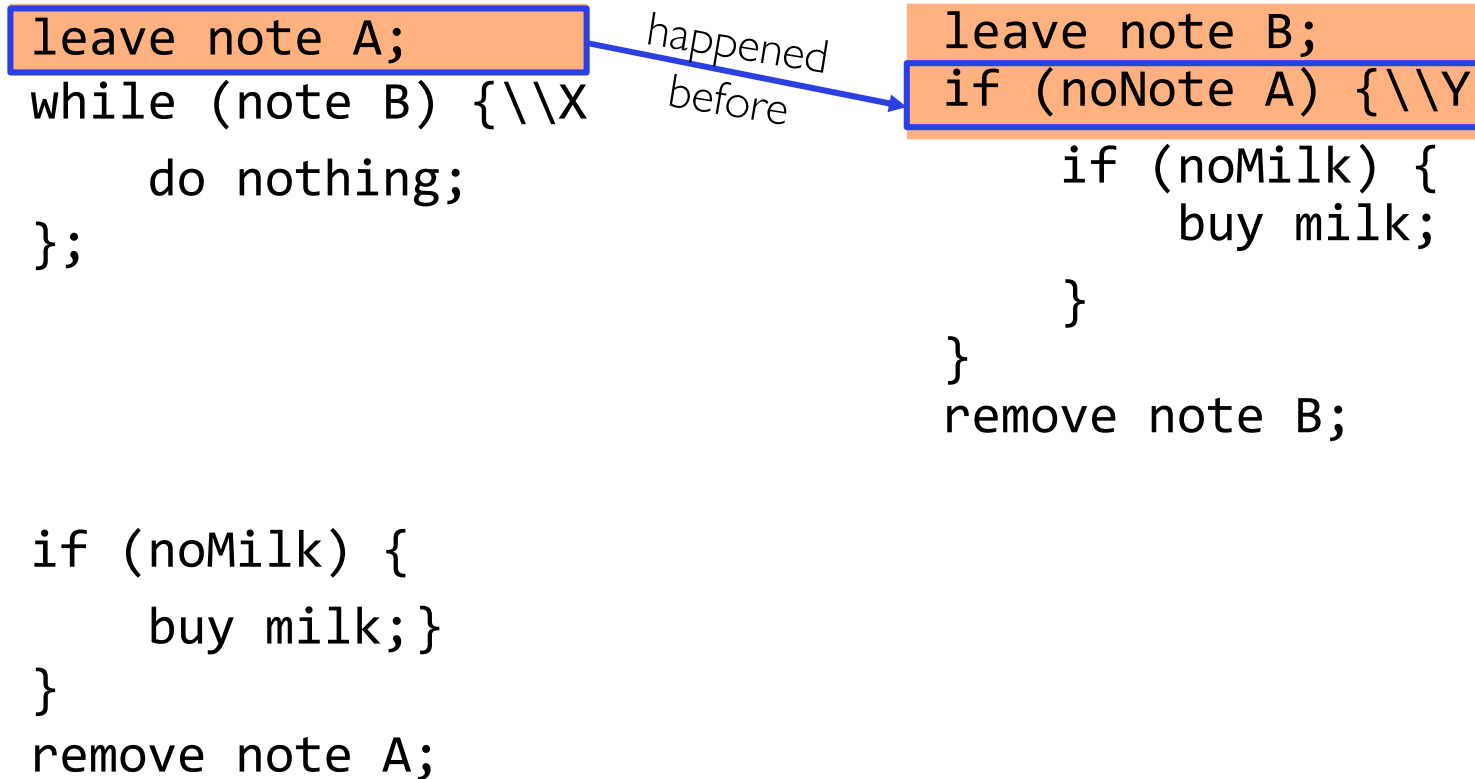
```
Thread A
leave note A;
while (note B) {\\X
    do nothing;
}
if (noMilk) {
    buy milk;
}
remove note A;
```

```
Thread B
leave note B;
if (noNote A) {\\Y
    if (noMilk) {
        buy milk;
    }
}
remove note B;
```

- Does this work? **Yes**. Both can guarantee that:
  - It is safe to buy, or
  - Other will buy, ok to quit
- At X:
  - If no note B, safe for A to buy,
  - Otherwise wait to find out what will happen
- At Y:
  - If no note A, safe for B to buy
  - Otherwise, A is either buying or waiting for B to quit

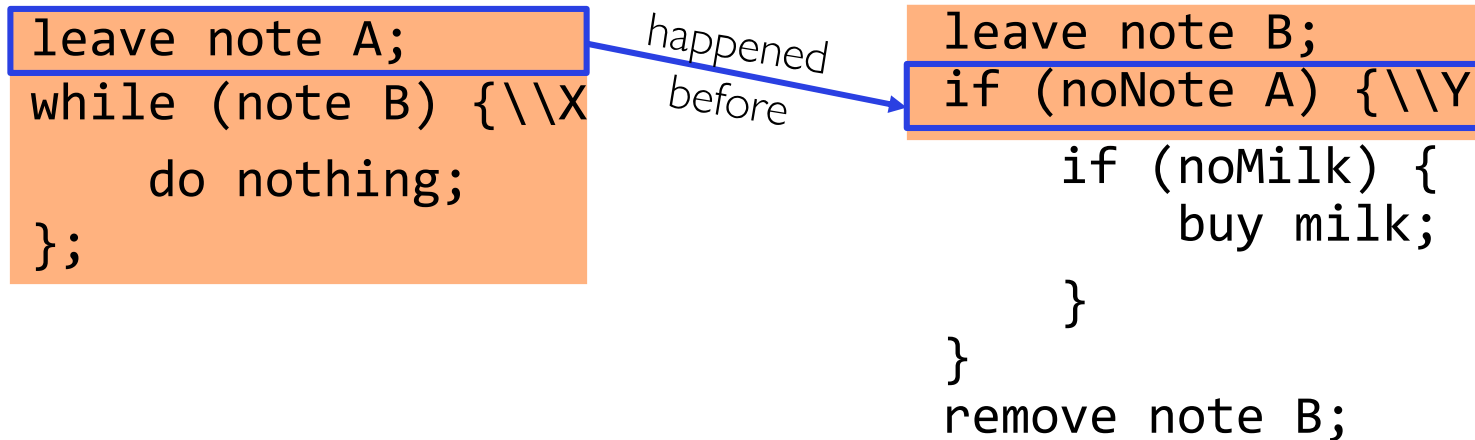
# Case 1

- “leave note A” happens before “if (noNote A)”



# Case 1

- “leave note A” happens before “if (noNote A)”

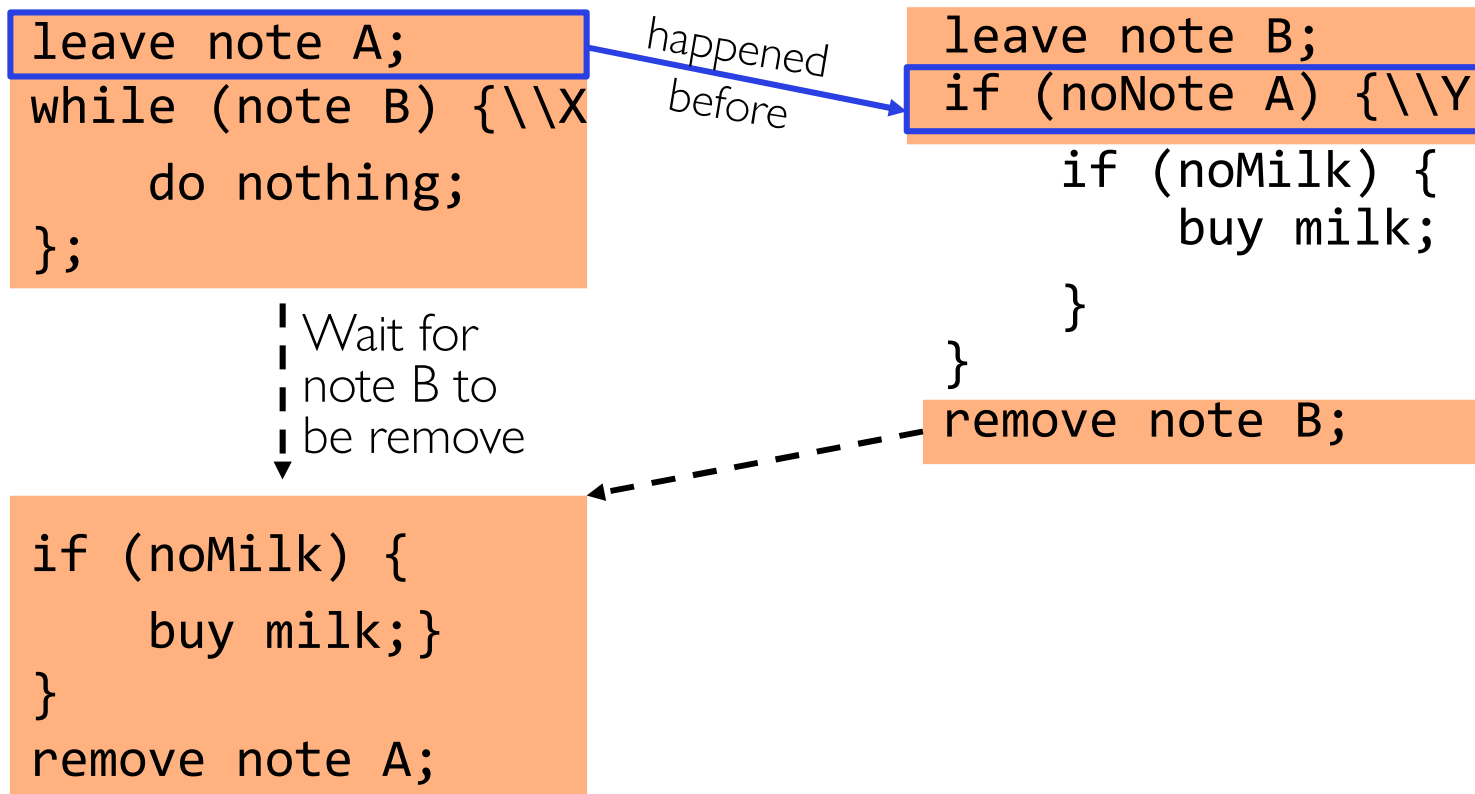


```
if (noMilk) {  
    buy milk;  
}  
remove note A;
```



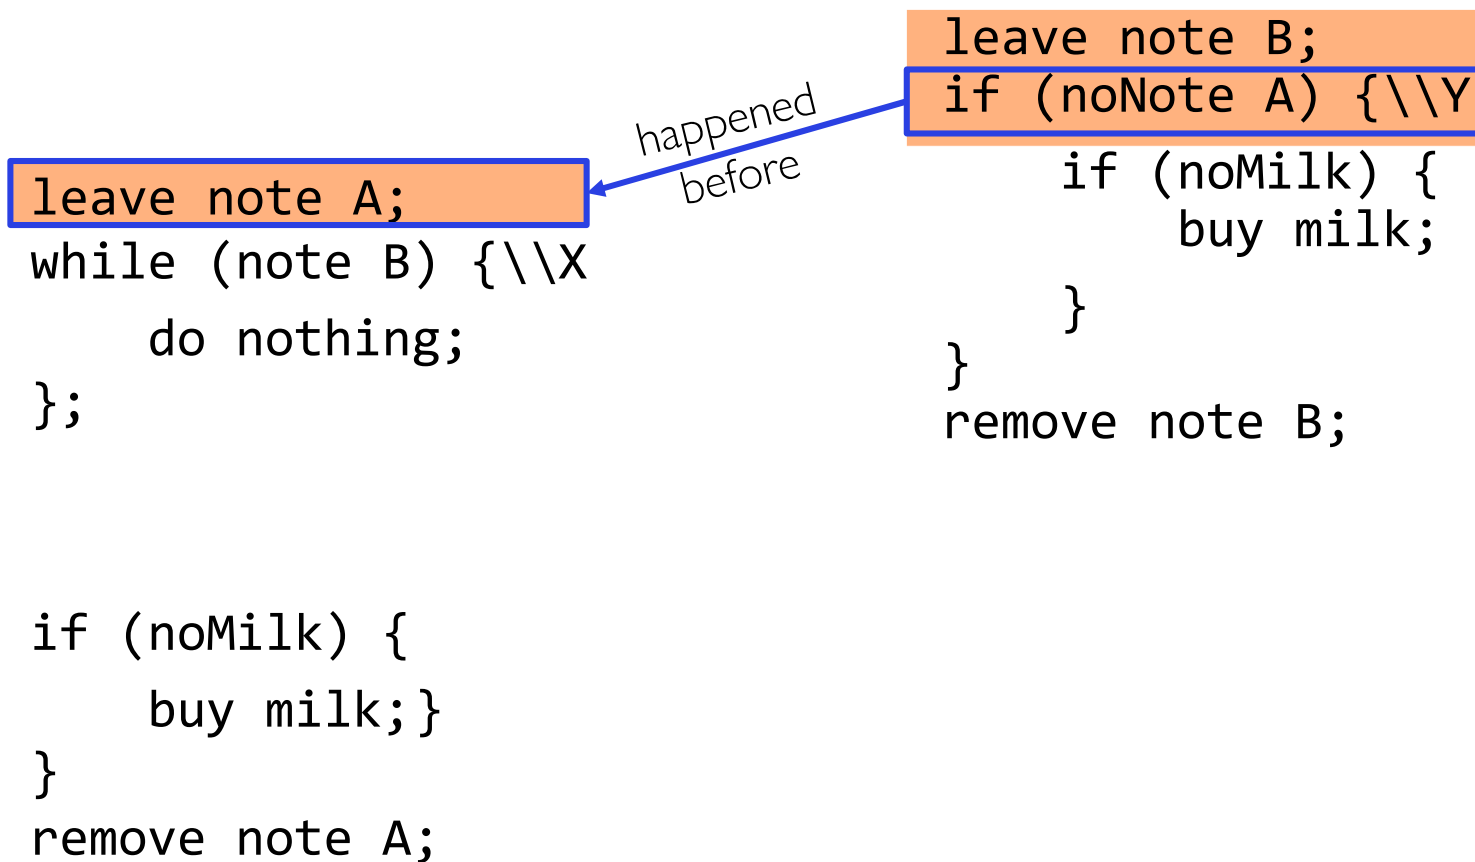
# Case 1

- “leave note A” happens before “if (noNote A)”




## Case 2

- “if (noNote A)” happens before “leave note A”



## Case 2

- “if (noNote A)” happens before “leave note A”

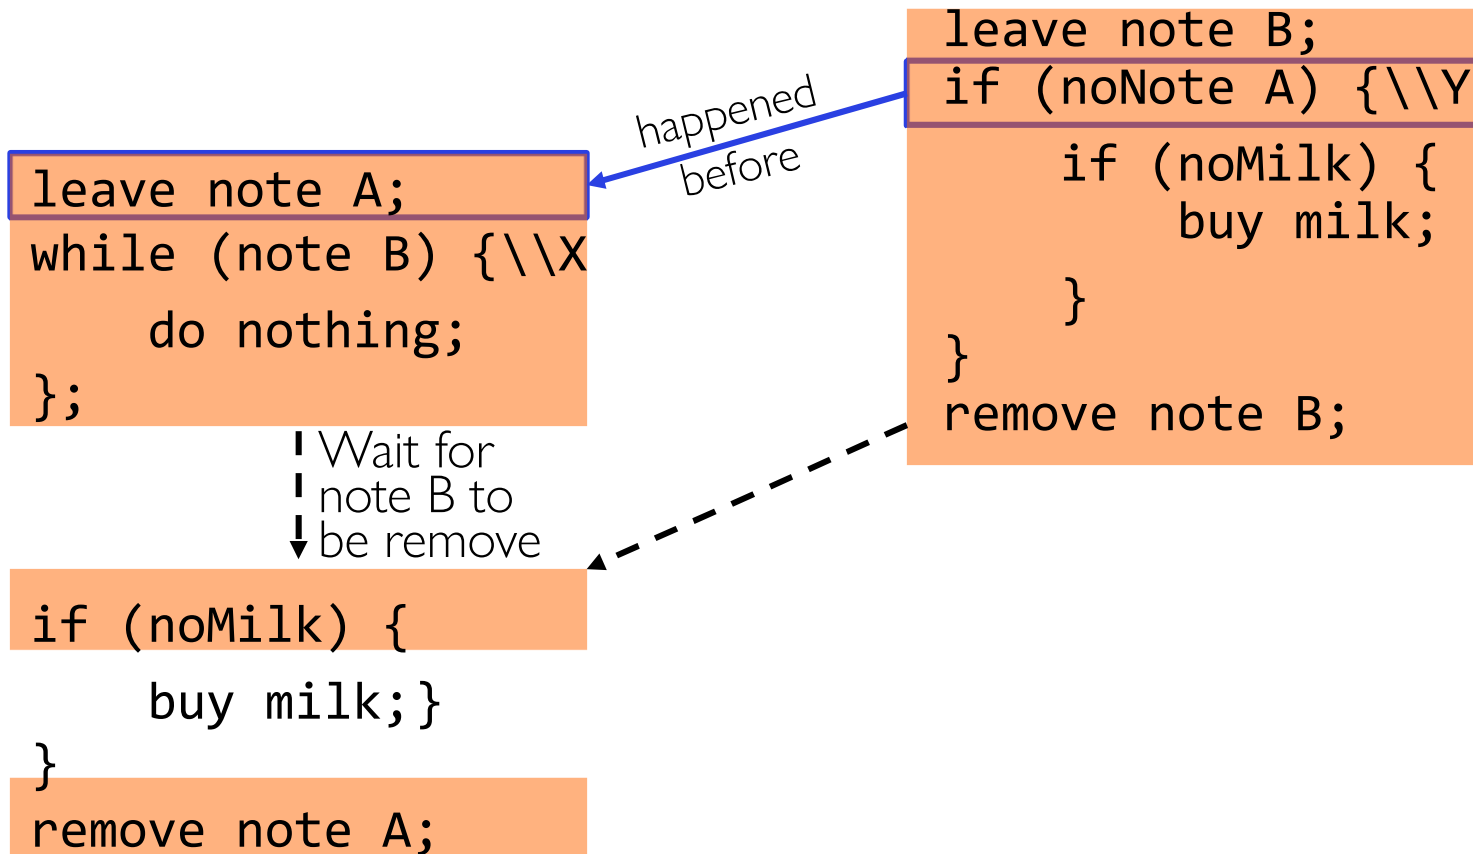


```
if (noMilk) {  
    buy milk;}  
}  
remove note A;
```



## Case 2

- “if (noNote A)” happens before “leave note A”



# Solution #3 discussion

---

- Our solution protects a single “Critical-Section” piece of code for each thread:

```
if (noMilk) {  
    buy milk;  
}
```

- Solution #3 works, but it's really unsatisfactory
  - Really complex – even for this simple an example
    - Hard to convince yourself that this really works
  - A's code is different from B's – what if lots of threads?
    - Code would have to be slightly different for each thread
  - While A is waiting, it is consuming CPU time
    - This is called “busy-waiting”
- There's a better way
  - Have hardware provide higher-level primitives than atomic load & store
  - Build even higher-level programming abstractions on this hardware support

# Too Much Milk: Solution #4

---

- Suppose we have some sort of implementation of a lock
  - `lock.Acquire()` – wait until lock is free, then grab
  - `lock.Release()` – Unlock, waking up anyone waiting
  - These must be **atomic operations** – if two threads are waiting for the lock and both see it's free, only one succeeds to grab the lock
- Then, our milk problem is easy:

```
    milklock.Acquire();  
    if (nomilk)  
        buy milk;  
    milklock.Release();
```

- Once again, section of code between `Acquire()` and `Release()` called a “**Critical Section**”
- Of course, you can make this even simpler: suppose you are out of ice cream instead of milk
  - Skip the test since you always need more ice cream ;-)



# Where are we going with synchronization?

Programs	Shared Programs
Higher-level API	Locks Semaphores Monitors
Hardware	Load/Store Disable Ints Test&Set Compare&Swap

- We are going to implement various higher-level synchronization primitives using atomic operations
  - Everything is pretty painful if only atomic primitives are load and store
  - Need to provide primitives useful at user-level



# Critical Section Problem






# Definition

---

- Race condition
  - Several processes access and manipulate the same data **concurrently**
  - Outcomes of the execution **depends** on the **order** in which the access take place
- How to remove **Race Condition**?
  - **Serial execution**





# Critical section problem


---

- Consider system of  $n$  processes  $\{p_0, p_1, \dots, p_{n-1}\}$
- Each process has **critical section** segment of code
  - Process may be changing common variables, updating table, writing file, etc
  - When one process in critical section, no other may be in its critical section
- **Critical section problem** is to design protocol to solve this
- Each process **must ask permission** to enter critical section in **entry section**, may follow critical section with **exit section**, then **remainder section**

# Critical section

---

- General structure of process  $P_i$



# Requirements to solutions

---

- **Mutual exclusion**

- If process  $P_i$  is executing in its critical section, then **no other processes** can be executing in their critical sections

- **Progress**

- If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the **selection of the processes** that will enter the critical section next **cannot be postponed indefinitely**

- **Bounded waiting**

- A **bound** must exist on the **number of times that other processes are allowed** to enter their critical sections after a process has made a request to enter its critical section and before that request is granted
  - Assume that each process executes at a nonzero speed
  - No assumption concerning **relative speed** of the  $n$  processes






# Preemption definition

---

- Preemption
  - The act of temporarily interrupting a [task](#) being carried out by a [computer system](#), without requiring its cooperation, and with the intention of resuming the task at a later time [[wiki](#)]



# Handling critical-section by OS

---

- Two approaches, depend on type of OS kernels
  - **Preemptive**
    - Allows preemption of process when running in kernel mode
    - Difficult to design in SMP architectures (why?)
  - **Non-preemptive**
    - Runs until exits kernel mode, blocks, or voluntarily yields CPU
    - Essentially free of race conditions in kernel mode (why?)
- Which one
  - is responsive?
  - is suitable for real-time programming?

# 1) Peterson's solution

---

- A classic SW solution
- No guarantees in correct working of the method
  - Correctness depends on computer architecture
  - Atomic instructions are needed (which & where?)
- Good algorithm!
- Shared variables
  - `int turn; /* whose turn is */`
  - `Boolean flag[2] /* who enters the critical-section */`



# Peterson algorithm for $P_i$

$(P_i, P_j) = (P_0, P_1)$

```
do {  
    flag[i] = true;  
    turn = j;  
    while (flag[j] && turn == j);  
    critical section  
    flag[i] = false;  
    remainder section  
} while (true);
```

How requirements are satisfied?

Mutual exclusion (?)

Progress (?)

Bounded waiting (?)

the algorithm uses **two variables**, *flag* and *turn*.

- A *flag[n]* value of *true* indicates that the process *n* wants to enter the **critical section**.
- *turn* resolves simultaneously conflicts
  - entrance to the critical section is granted for process *P0* if *P1* does not want to enter its critical section or if *P1* has given priority to *P0* by setting *turn* to 0.

# Proof

---

- **Mutual Exclusion**

- assume both P0 and P1 are in their CS
- then  $\text{flag}[0] = \text{flag}[1] = \text{true}$
- the test for entry cannot have been true for both processes at the same time (because turn favors one);
- therefore one process must have entered its CS first; the other process will be blocked.

- **Progress**

- **Case I: (Stuck)**
- P1 is not interested in entering its CS
- then  $\text{flag}[1] = \text{false}$ ; hence the while loop is false for P0 and it can go
- **Case II: (Deadlock)**
- P1 is also blocked at the while loop
- impossible, because  $\text{turn} = 0$  or  $1$ ; hence the while loop is false for some process and it can go

## 2) Hardware solution

---

- Some hardware support implementing the critical section code!
- All solutions are based on idea of **locking**
  - Protecting critical regions via locks
- **Uniprocessors** – **could disable interrupts**
  - Currently running code would execute without preemption
  - Generally too inefficient on multiprocessor systems
    - Operating systems using this not broadly scalable
- **Multiprocessors** – **provide special atomic hardware instructions**
  - **Atomic** = non-interruptible
  - either
    - test memory word and set value
    - swap contents of two memory words



# Hardware solution for critical section

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (TRUE);
```

How requirements are satisfied?

Mutual exclusion (?)

Progress (?)

Bounded waiting (?)

# *test\_and\_set* instruction

Definition:

```
boolean test_and_set (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;    /* old value */
}
```

1. Executed **atomically**
2. Returns the original value of passed parameter
3. Set the new value of passed parameter to "TRUE".

# Hardware solution using *test\_and\_set()*

- Shared Boolean variable lock, initialized to FALSE

```
do {  
    while (test_and_set(&lock))  
        ; /* do nothing */  
    /* critical section */  
    lock = false;  
    /* remainder section */  
} while (true);
```



# compare\_and\_swap instruction

## Definition:

```
int compare_and_swap(int *value, int expected, int
new_value) {
    int temp = *value;

    if (*value == expected)
        *value = new_value;
    return temp;    /* old value */
}
```

1. Executed **atomically**
2. Returns the original value of passed parameter “value”
3. Set the variable “value” the value of the passed parameter “new\_value” but only if “value” == “expected”. That is, the swap takes place only under this condition.

# Hardware solution using *compare\_and\_swap()*

➤ Shared integer “lock” initialized to 0;

```
do {  
    while (compare_and_swap(&lock, 0, 1) != 0)  
        ; /* do nothing */  
  
    /* critical section */  
  
    lock = 0;  
  
    /* remainder section */  
  
} while (true);
```

How requirements are satisfied?

Mutual exclusion (?)


Progress (?)

Bounded waiting (?)

## Bounded-waiting mutual exclusion with test\_and\_set

```
do {
    waiting[i] = true;
    key = true;
    while (waiting[i] && key)
        key = test_and_set(&lock);
    waiting[i] = false;
    /* critical section */
    j = (i + 1) % n;
    while ((j != i) && !waiting[j])
        j = (j + 1) % n;
    if (j == i)
        lock = false;
    else
        waiting[j] = false;
    /* remainder section */
} while (true);
```

- Key indicates who can enter CS
- Use waiting array indicates who is waiting to enter CS




Waiting array



## Case :

- $P_i$  can enter its critical section only if either
  - $Waiting[i] == false$  or  $key == false$
- When  $P_i$  leaves its critical section, it scans the array waiting in the cyclic order
  - $i+1, i+2, \dots, n-1, 0, \dots, i-1$
  - find the first process in waiting array



Waiting array

### 3) OS solution!: Mutex locks

---

- Previous solutions are **complicated** and generally **inaccessible** to application programmers
- OS designers build software tools to solve critical section problem
- Simplest is **mutex** lock (*mutual exclusions*)
- Protect a critical section by first **acquire()** a lock then **release()** the lock
  - Boolean variable indicating if lock is available or not
- Calls to **acquire()** and **release()** must be **atomic**
  - Usually implemented via **hardware atomic instructions**
- But this solution requires **busy waiting**
  - This lock therefore called a **spinlock**

# acquire() and release()

➤ `acquire() {`  
    `while (!available)`  
        `; /* busy wait */`  
    `available = false;;`  
`}`

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (true);
```

➤ `release() {`  
    `available = true;`  
`}`

How requirements are satisfied?

Mutual exclusion (?)

Progress (?)

Bounded waiting (?)



# Naïve use of Interrupt Enable/Disable

- How can we build multi-instruction atomic operations?
  - Recall: dispatcher gets control in two ways.
    - Internal: Thread does something to relinquish the CPU
    - External: Interrupts cause dispatcher to take CPU
  - On a uniprocessor, can avoid context-switching by:
    - Avoiding internal events (although virtual memory tricky)
    - Preventing external events by disabling interrupts
- Consequently, naïve Implementation of locks:

```
LockAcquire { disable Ints; }  
LockRelease { enable Ints; }
```

- Problems with this approach:
  - **Can't let user do this!** Consider following:

```
LockAcquire();  
While(TRUE) {;
```

- **Real-Time system**—no guarantees on timing!
  - Critical Sections might be arbitrarily long



## Better Implementation of Locks by Disabling Interrupts

- Key idea: maintain a lock variable and impose mutual exclusion only during operations on that variable

```
int value = FREE;
```



```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
        // Enable interrupts?  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

```
Release() {  
    disable interrupts;  
    if (anyone on wait queue) {  
        take thread off wait queue  
        Place on ready queue;  
    } else {  
        value = FREE;  
    }  
    enable interrupts;  
}
```



What is the main problem of all mentioned methods?

**Busy waiting!**








## 4) Semaphores

- Semaphores are a kind of generalized lock
  - First defined by **Dijkstra** in late 60s
  - Main synchronization primitive used in original UNIX
- **Definition:** a Semaphore has a non-negative integer value and supports the following two operations:
  - **P()**: an atomic operation that waits for semaphore to become positive, then decrements it by 1
    - Think of this as the **wait()** operation
  - **V()**: an atomic operation that increments the semaphore by 1, waking up a waiting P, if any
    - Think of this as the **signal()** operation
  - Note that **P()** stands for “*proberen*” (to test) and **V()** stands for “*verhogen*” (to increment) in Dutch

# Semaphores Like Integers Except

- Semaphores are like integers, except
  - No negative values
  - Only operations allowed are P and V – can't read or write value, except to set it initially
  - Operations must be atomic
    - Two P's together can't decrement value below zero
    - Similarly, thread going to sleep in P won't miss wakeup from V – even if they both happen at same time
- Semaphore from railway analogy
  - Here is a semaphore initialized to 2 for resource control:



# Semaphore(cont.)

- Synchronization tool that provides more sophisticated ways (than **Mutex** locks) for process to synchronize their activities
  - Semaphore **S** – integer variable
  - Can only be accessed via two indivisible (**atomic**) operations

**P()** (**wait()**) and **V()** (**signal()**)

**S. P()**

```
wait(S)
{
    while (S < 0)
        ; // busy wait
    S--;
}
```

**S. V()**

```
signal(S)
{
    S++;
}
```



# The implementation of semaphore

```
Class Semaphore {  
    int sem;  
    WaitQueue q;  
}
```

```
Semaphore::P() {  
    sem--;  
    if (sem < 0) {  
        Add this thread t to q;  
        block(p);  
    }  
}
```

```
Semaphore::V() {  
    sem++;  
    if (sem <= 0) {  
        Remove a thread t from q;  
        wakeup(t);  
    }  
}
```

Note that in this implementation, semaphore values **may be negative**, whereas semaphore values are never negative under the classical definition of semaphores with busy waiting.

# Types of semaphore

- Types
  - Binary semaphore (same as mutex lock)
  - Counting semaphore (suitable for managing number of resources)
- Can solve various synchronization problems
- Example:
  - Consider  $P_1$  and  $P_2$  that require  $S_1$  to happen before  $S_2$

Create a semaphore “**synch**” initialized to **zero**

P1 :

```
S1 ;  
signal (synch) ;
```

P2 :

```
wait (synch) ;  
S2 ;
```

# Two Uses of Semaphores: Mutual Exclusion

**Mutual Exclusion** (initial value = 1)

- Also called “Binary Semaphore”.
- Can be used for mutual exclusion:

```
semaphore.P();  
    // Critical section goes here  
semaphore.V();
```

```
mutex = new Semaphore(1);
```

```
//Thread A:  
mutex.P();  
    // Critical section  
mutex.V();
```

```
//Thread B:  
mutex.P();  
    // Critical section  
mutex.V();
```



# Two Uses of Semaphores: **synchronization**

Scheduling Constraints (initial value = 0)

- Allow thread 1 to wait for a signal from thread 2
  - thread A **schedules** thread B when a given **event** occurs


```
condition = new Semaphore(0);
```

线程A

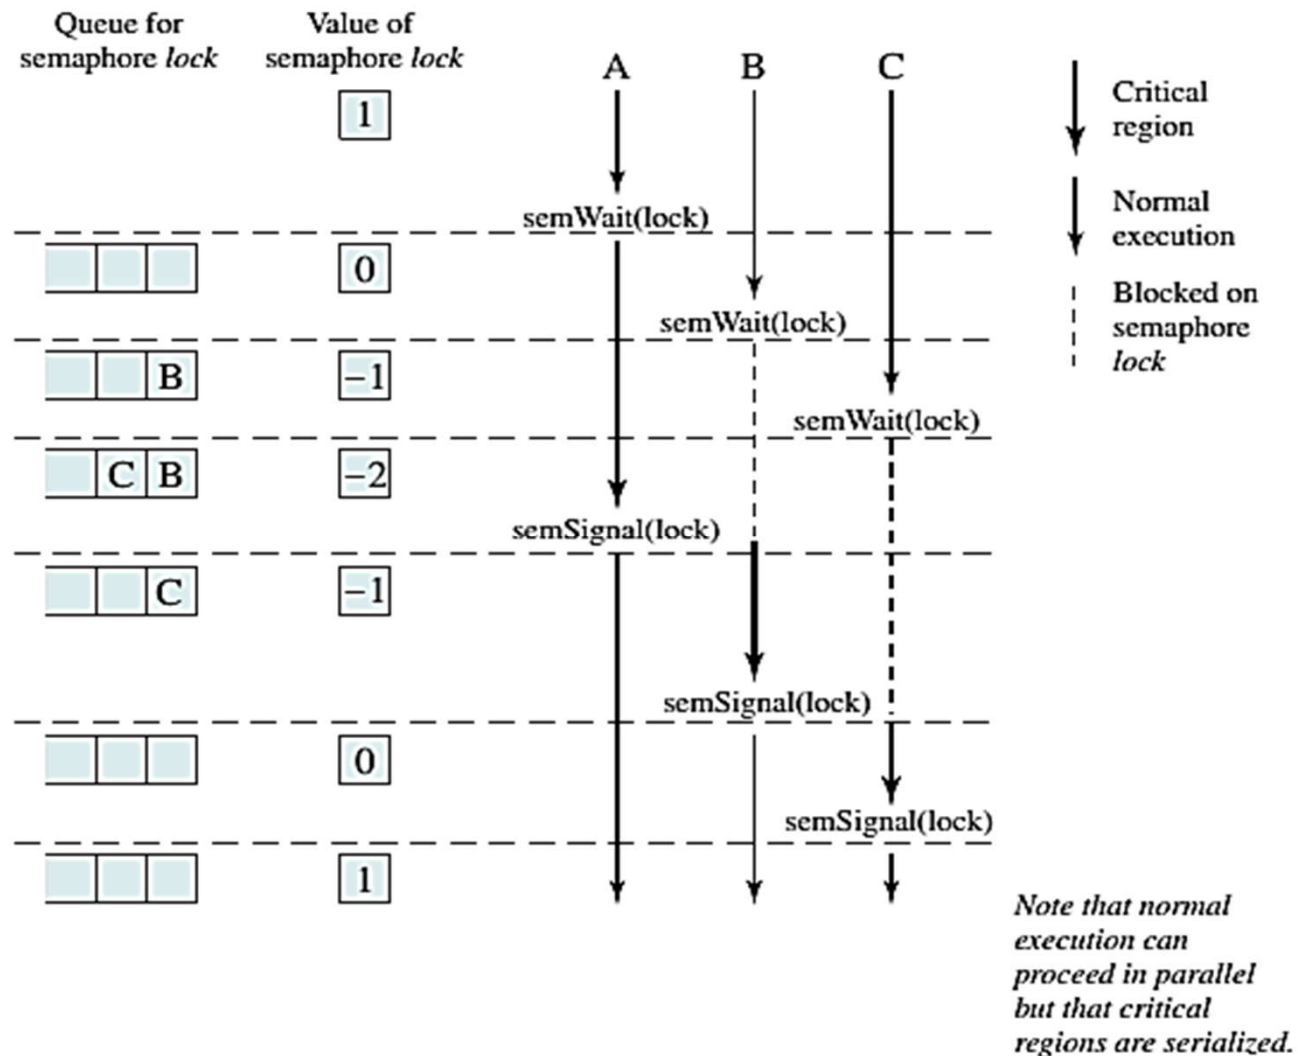
```
... M ...  
condition->P();  
... N ...
```

线程B

```
... X ...  
condition->V();  
... Y ...
```



# Accessing shared data by Semaphore



# Semaphore points

---

- Must guarantee that no two processes can execute the **wait()** and **signal()** on the same semaphore at the same time (**why?**)
  - wait() and signal() must be **atomic!**
  - wait() and signal() generate a **Critical Section Problem!**
  - **How to solve?**
    - Uniprocessors
      - Disabling interrupts
    - SMP (Multiprocessors)
      - Disabling interrupts (**bad performance effect**)
      - Other methods: **compare\_and\_swap()** and **spinlock** (**is it good to have busy waiting?**)



# Two implementations of semaphores

```
semWait(s)
{
    while (compare_and_swap(s.flag, 0, 1) == 1)
        /* do nothing */;
    s.count--;
    if (s.count < 0) {
        /* place this process in s.queue*/;
        /* block this process (must also set
s.flag to 0) */;
    }
    s.flag = 0;
}

semSignal(s)
{
    while (compare_and_swap(s.flag, 0, 1) == 1)
        /* do nothing */;
    s.count++;
    if (s.count <= 0) {
        /* remove a process P from s.queue */;
        /* place process P on ready list */;
    }
    s.flag = 0;
}
```

(a) Compare and Swap Instruction

```
semWait(s)
{
    inhibit interrupts;
    s.count--;
    if (s.count < 0) {
        /* place this process in s.queue */;
        /* block this process and allow inter-
rupts*/;
    }
    else
        allow interrupts;
}

semSignal(s)
{
    inhibit interrupts;
    s.count++;
    if (s.count <= 0) {
        /* remove a process P from s.queue */;
        /* place process P on ready list */;
    }
    allow interrupts;
}
```

(b) Interrupts

# Problems with semaphores

- Be careful in the usage
  - Deadlock, Starvation, Priority inversion

$P_0$

```
wait(S) ;  
wait(Q) ;  
...  
signal(S) ;  
signal(Q) ;
```

$P_1$

```
wait(Q) ;  
wait(S) ;  
...  
signal(Q) ;  
signal(S) ;
```

- Starvation
  - LIFO queue
- Priority Inversion – Scheduling problem when lower-priority process holds a lock needed by higher-priority process
  - Example:  $L(R) < M < H(R)$
  - Solved via priority-inheritance protocol




# Classic synchronization problems

---

- The bounded-buffer problem
- The readers-writers problem
- The dining-philosophers problem


How can semaphore solve these problems?





# Producer-Consumer with a Bounded Buffer

- Problem Definition
  - Producer puts things into a shared buffer
  - Consumer takes them out
  - Need synchronization to coordinate producer/consumer
- Don't want producer and consumer to have to work in lockstep, so put a fixed-size buffer between them
  - Need to synchronize access to this buffer
  - Producer needs to wait if buffer is full
  - Consumer needs to wait if buffer is empty



# Correctness constraints for solution

- Correctness Constraints:
  - Consumer must wait for producer to fill buffers, if none full (**scheduling constraint**)
  - Producer must wait for consumer to empty buffers, if all full (**scheduling constraint**)
  - Only one thread can manipulate buffer queue at a time (**mutual exclusion**)
- Remember why we need mutual exclusion
  - Because computers are stupid
- General rule of thumb:  
**Use a separate semaphore for each constraint**
  - Semaphore fullBuffer; // consumer's constraint
  - Semaphore emptyBuffer; // producer's constraint
  - Semaphore mutex; // mutual exclusion

# Full Solution to Bounded Buffer

```
Semaphore mutex = 1;  
Semaphore fullBuffer = 0; //empty  
Semaphore emptyBuffer = bufSize;
```

```
Producer(item) {  
    emptyBuffer.P();  
    mutex.P();  
    Add item to the buffer;  
    mutex.V();  
    fullBuffer.V();  
}
```

```
Consumer(item) {  
    fullBuffers.P();  
    mutex.P();  
    Remove item from buffer;  
    mutex.V();  
    emptyBuffers.V();  
}
```

Does it matter about the order of P()、V();



# Discussion about Solution

Decrease # of  
empty buffer

Increase # of  
occupied buffer

- Why asymmetry?
  - Producer does: `emptyBuffer.P()`, `fullBuffer.V()`
  - Consumer does: `fullBuffer.P()`, `emptyBuffer.V()`

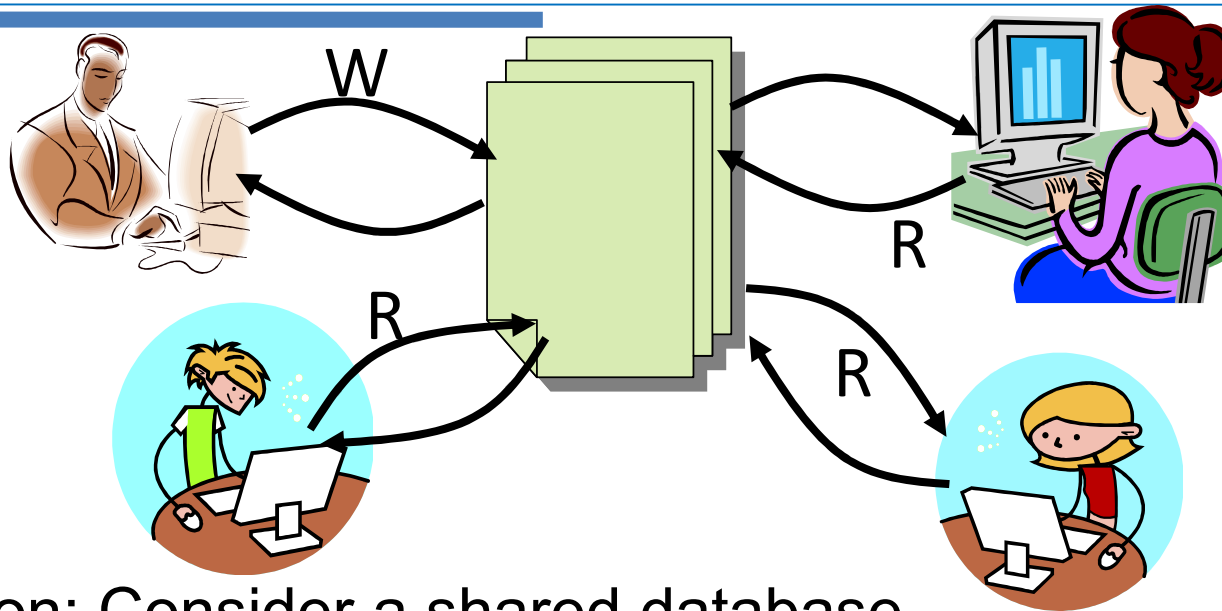
Decrease # of  
occupied buffer

Increase # of  
empty buffer

- Is order of P's important?
- Is order of V's important?
  - No, except that it might affect scheduling efficiency
- What if we have 2 producers or 2 consumers?
  - Do we need to change anything?

```
Producer(item) {  
    mutex.P();  
    emptyBuffer.P();  
    Enqueue(item);  
    mutex.V();  
    fullBuffer.V();  
}  
Consumer() {  
    fullBuffer.P();  
    mutex.P();  
    item = Dequeue();  
    mutex.V();  
    emptyBuffer.V();  
    return item;  
}
```

# Readers/Writers Problem



- Motivation: Consider a shared database
  - Two classes of users:
    - Readers – never modify database
    - Writers – read and modify database
  - Is using a single lock on the whole database sufficient?
    - Like to have many readers at the same time
    - Only one writer at a time

# Basic Readers/Writers Solution

- Correctness Constraints:
  - Readers can access database when no writers
  - Writers can access database when no readers or writers
  - Only one thread manipulates state variables at a time
- Basic structure of a solution:

- **Reader ()**

- Wait until no writers

- Access data base

- Check out - wake up a waiting writer

- **Writer ()**

- Wait until no active readers or writers

- Access database


- Check out - wake up waiting readers or writer

- Semaphore and variable:

- ```
semaphore rw_mutex = 1; //read and write semaphore
```

- ```
semaphore mutex = 1; // mutual exclusion for read_count
```

- ```
int read_count = 0; // currently reading
```





# Solution 1: using semaphore

---

Writer

`write;`

Reader

`read;`

# Solution 1: using semaphore(cont.)

Writer

```
P(rw_mutex);  
  
write;  
  
V(rw_mutex);
```

Reader

```
P(rw_mutex);  
  
read;  
  
V(rw_mutex);
```

# Solution 1: using semaphore(cont.)

Writer

```
P(rw_mutex);  
  
write;  
  
V(rw_mutex);
```

Reader

```
if (read_count == 0)  
    P(rw_mutex);  
    ++read_count;  
  
read;  
  
V(rw_mutex);
```



# Solution 1: using semaphore(cont.)

Writer

```
P(rw_mutex);  
  
write;  
  
V(rw_mutex);
```

Reader

```
if (read_count == 0)  
    P(rw_mutex);  
    ++read_count;  
  
read;  
  
--read_count;  
if (read_count == 0)  
    V(rw_mutex);
```

# Solution 1: using semaphore(cont.)

Writer

```
P(rw_mutex);  
  
write;  
  
V(rw_mutex);
```

Reader

```
P(mutex);  
if (read_count == 0)  
    P(rw_mutex);  
    ++read_count;  
V(mutex);  
  
read;  
  
--read_count;  
if (read_count == 0)  
    V(rw_mutex);
```

# Solution 1: using semaphore(cont.)

Writer

```
P(rw_mutex);  
  
write;  
  
V(rw_mutex);
```

Reader first


Reader

```
P(mutex);  
if (read_count == 0)  
    P(rw_mutex);  
    ++read_count;  
V(mutex);  
  
read;  
  
P(mutex);  
--read_count;  
if (read_count == 0)  
    V(rw_mutex);  
V(mutex)
```



# Dining Philosophers Problem

- *Five* philosophers seated around a circular table
  - There is one chopstick between each philosopher
  - A philosopher must pick up its two nearest chopsticks in order to eat
  - A philosopher must pick up first one chopstick, then the second one, not both at once
- Devise an algorithm for allocating these limited resources (chopsticks) among several processes (philosophers) in a manner that is
  - deadlock-free, and
  - starvation-free



# The dining-philosophers problem

- Five philosophers are in a thinking - eating cycle.
- When a philosopher gets hungry, he sits down, picks up two nearest chopsticks, and eats.
- A philosopher can eat only if he has both chopsticks.
- After eating, he puts down both chopsticks and thinks.
- This cycle continues.




# The dining-philosophers problem(cont.)

- **Chopsticks are shared items** (by two philosophers) and must be protected.
- Each chopstick has a semaphore with **initial value 1**.
- A philosopher calls wait() before picks up a chopstick and calls signal() to release it.





# Solution 1:



Thinking and eating alternatively

```
#define N 5                                     // number of philosopher
semaphore chopstick[N];                         // semaphore initialize 1
void philosopher(int i)                        // philosopher id: 0 - 4
{
    while (TRUE)
    {
        think( );                             // thinking
        P(chopstick[i]);                       // get left chopstick
        P(chopstick[(i + 1) % N]);             // get right chopstick
        eat( );                                // eating....
        V(chopstick[i]);                       // put left chopstick
        V(chopstick[(i + 1) % N]);             // put right chopstick
    }
}
```

Any problem?  
**starvation**

## Solution 2:

---

```
#define    N    5                // number of philosopher
semaphore chopstick [5];        // semaphore initialize 1
semaphore  mutex;               // mutex, initial value 1
```

```
#define      N      5
semaphore chopstick [5];
semaphore mutex;
void philosopher(int i)
    while(TRUE) {
        think( );

        eat( );

    }

// number of philosopher
// semaphore initialize 1
// mutex, initial value 1
// philosopher Id: 0 - 4
// thinking
// eating ...
```



# Solution 2:

```
#define    N    5                // number of philosopher
semaphore chopstick [5];        // semaphore initialize 1
semaphore  mutex;               // mutex, initial value 1
void  philosopher(int    i)     // philosopher Id: 0 - 4
    while(TRUE) {
        think( );              // thinking
        P(mutex);              // enter CS

        eat( );                // eating ...

        V(mutex);              // leaving CS

    }
```

# Solution 2:

```
#define    N    5                // number of philosopher
semaphore chopstick [5];        // semaphore initialize 1
semaphore  mutex;                // mutex, initial value 1
void  philosopher(int    i)      // philosopher Id: 0 - 4
{
    while(TRUE) {
        think( );                // thinking
        P(mutex);                // enter CS

        P(chopstick[i]);          // get left chopstick
        P(chopstick[(i + 1) % N]); //get right chopstick

        eat( );                  // eating ...

        V(mutex);                // leaving CS
    }
}
```

## Solution 2:

```
#define    N    5                // number of philosopher
semaphore chopstick [5];        // semaphore initialize 1
semaphore  mutex;                // mutex, initial value 1
void  philosopher(int  i)        // philosopher Id: 0 - 4
{
    while(TRUE) {
        think( );                // thinking
        P(mutex);                // enter CS

        P(chopstick[i]);          // get left chopstick
        P(chopstick[(i + 1) % N]); //get right chopstick

        eat( );                  // eating ...

        V(chopstick[i]);          // put left chopstick
        V(chopstick[(i + 1) % N]); // put right chopstick
        V(mutex);                // leaving CS
    }
}
```

Any problem? It's correct! But just one philosopher can eat!




# Dining Philosophers Problem

---

- Some deadlock-free solutions:
  - allow at most 4 philosophers at the same table when there are 5 resources
  - odd philosophers pick first left then right, while even philosophers pick first right then left
  - allow a philosopher to pick up chopsticks only if both are free. This requires protection of critical sections to test if both chopsticks are free before grabbing them.
    - ✓ we'll see this solution next using monitors
- A deadlock-free solution is not necessarily starvation-free
  - for now, we'll focus on breaking deadlock

## Solution 3:

- **asymmetric solution**
- breaking the circular wait
- **an** odd philosopher pick up first her left chopstick and then picks up her right chopstick ,
- **an** even philosopher picks up her right chopstick and then her left chopstick
- when 1,3,5 got left chopstick , 2,4 cannot get their right chopstick
- Now, 1,3 got their right chopstick too and got executed. After that 2,4 can participate.



# Solution 3:

```
#define    N    5                // number of philosopher
semaphore chopstick [5];        // semaphore initialize 1
semaphore  mutex;                // mutex, initial value 1
void  philosopher(int  i)        // philosopher Id: 0 - 4
    while(TRUE) {
        think( );                // thinking

        if (i%2 == 0) {
            P(chopstick[i]);      // get left chopstick
            P(chopstick[(i + 1) % N]); //get right chopstick
        else{
            P(chopstick[i+1]%N);  // get right chopstick
            P(chopstick[i % N]);  //get left chopstick
        }

        eat( );                  // eating ...

        V(chopstick[i]);          // put left chopstick
        V(chopstick[(i + 1) % N]); // put right chopstick
    }
```



# Other problems with semaphore

- Problems with **bad** usage

```
signal(mutex);  
...  
critical section  
...  
wait(mutex);
```

```
wait(mutex);  
...  
critical section  
...  
wait(mutex);
```

```
...  
critical section  
...  
wait(mutex);
```

```
wait(mutex);  
...  
critical section  
...
```

- Deadlock** and **starvation** are possible.

# 5) Monitor


- A **high-level abstraction** that provides a convenient and effective mechanism for process synchronization
- Only **one process** may be active within the monitor at a time

```
monitor monitor-name
{
    // shared variable
    declarations

    procedure P1 (...) { ... }

    procedure Pn (...) {.....}

    Initialization_Code (...) { ... }
}
```



# Monitors and Condition Variables

- Example:

```
monitor sharedcounter {  
    int counter;  
    function add() { counter++;}  
    function sub() { counter--;}  
    init() { counter=0; }  
}
```

- If two processes want to access this *sharedcounter* monitor, then access is mutually exclusive and only one process at a time can modify the value of counter
  - if a write process calls `sharedcounter.add()`, then it has exclusive access to modifying counter until it leaves `add()`. No other process, e.g. a read process, can come in and call `sharedcounter.sub()` to decrement counter while the write process is still in the monitor



# Monitors and Condition Variables

---

- In the previous *sharedcounter* example, a writer process may be interacting with a reader process via a bounded buffer
  - like the solution to the bounded buffer producer/consumer problem, the writer should signal blocked reader processes when there are no longer zero elements in the buffer
  - monitors alone don't provide this signalling synchronization capability
- In general, there may be times when one process wishes to signal another process based on a condition, much like semaphores.
  - Thus, monitors alone are insufficient.
  - Augment monitors with *condition variables*.

# Monitor (with *condition variables*)

The only operations that can be invoked on a condition variable are **wait()** and **signal()**.

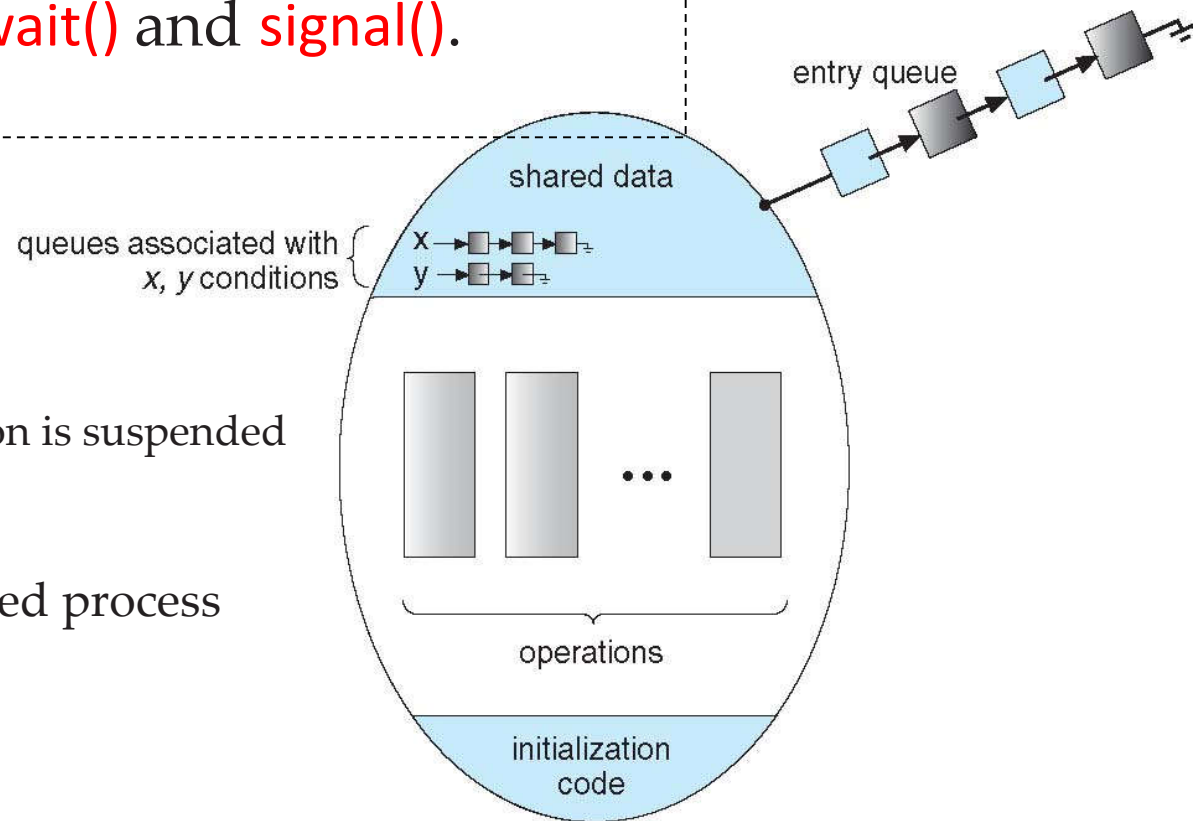
**condition x, y;**

// the process invoking this operation is suspended

**x.wait();**

// resumes exactly one suspended process

**x.signal();**

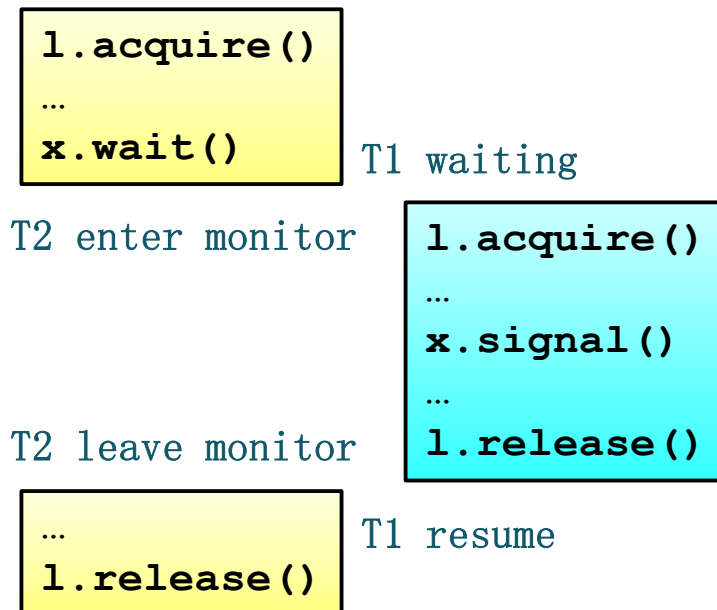


# Monitors and Condition Variables

Semantics concerning what happens just after `x.signal()` is called by a process P in order to wake up a process Q waiting on this CV x

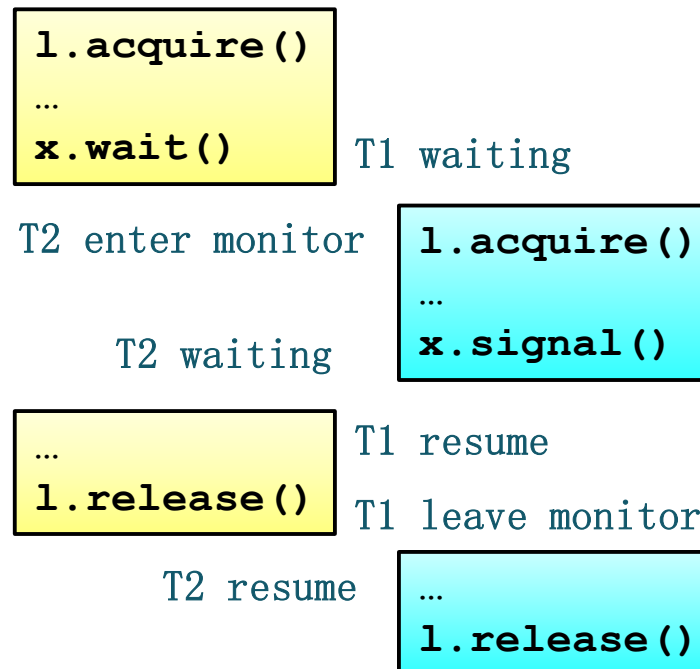
## ■ Hansen( singal-and-wait)

### ■ Applied in OS and Java



## ■ Hoare(signal-and-continue)

### ■ text book





# Monitor-based Solution to Dining Philosophers

---

- Key insight: pick up 2 chopsticks only if both are free
  - this avoids deadlock
  - a philosopher moves to his/her eating state only if both neighbors are not in their eating states
    - thus, need to define a state for each philosopher
  - if one of my neighbors is eating, and I'm hungry, ask them to signal() me when they're done
    - thus, states of each philosopher are: thinking, hungry, eating
    - thus, need condition variables to signal() waiting hungry philosopher(s)
  - Also, need to Pickup() and Putdown() chopsticks

# Monitor-based Solution to Dining Philosophers

- Some basic pseudo-code for monitor

```
Monitor
DiningPhilosophers {
    // THINKING; HUNGRY, EATING
    status state[5];
    condition self[5];
    Pickup(int i);
    Putdown(int i);
}
```

- Each philosopher  $i$  runs pseudo-code:

```
DP.Pickup( $i$ );
...
DP.Putdown( $i$ );
```

# The dining-philosophers problem

```
monitor DiningPhilosophers
{
    enum { THINKING; HUNGRY, EATING) state
[5] ;
    condition self [5];

    void pickup (int i) {
        state[i] = HUNGRY;
        test(i);
        if (state[i] != EATING)
            self[i].wait();
    }

    void putdown (int i) {
        state[i] = THINKING;
        // test left and right neighbors
        test((i + 4) % 5);
        test((i + 1) % 5);
    }
}
```

```
void test (int i) {
    if ((state[(i + 4) % 5] != EATING)
    &&
        (state[i] == HUNGRY) &&
        (state[(i + 1) % 5] != EATING) ) {
        state[i] = EATING;
        self[i].signal();
    }
}

initialization_code() {
    for (int i = 0; i < 5; i++)
        state[i] = THINKING;
}
}
```



# The dining-philosophers problem



```
procedure philosopher(i)
{
  while TRUE do
  {
    THINKING;
    DiningPhilosophers.pickup(i);
    EATING;
    DiningPhilosophers.putdown(i);
  }
}
```

Any problem?

No deadlock



# semaphores in C language



# Semaphores in C

- the POSIX system in Linux presents its own built-in semaphore library. To use it, we have to :

1. Include semaphore.h

2. Compile the code by linking with -lpthread -lrt

To declare a semaphore

```
sem_t sem;
```

To lock a semaphore or wait we can use the sem\_wait function:

```
int sem_wait(sem_t *sem);
```

To release or signal a semaphore, we use the **sem\_post** function:

```
int sem_post(sem_t *sem);
```

A semaphore is initialised by using **sem\_init**(for processes or threads) or **sem\_open** (for IPC).

```
sem_init(sem_t *sem, int pshared, unsigned int value);
```

To destroy a semaphore

```
sem_destroy(sem_t *mutex);
```



# EX:1

```
// C program to demonstrate working of Semaphores
#include <stdio.h>
#include <pthread.h>
#include <semaphore.h>
#include <unistd.h>

sem_t mutex;

void* thread(void* arg)
{
    //wait
    sem_wait(&mutex);
    printf("\nEntered..\n");

    //critical section
    sleep(4);

    //signal
    printf("\nJust Exiting...\n");
    sem_post(&mutex);
}

int main()
{
    sem_init(&mutex, 0, 1);
    pthread_t t1,t2;
    pthread_create(&t1,NULL,thread,NULL);
    sleep(2);
    pthread_create(&t2,NULL,thread,NULL);
    pthread_join(t1,NULL);
    pthread_join(t2,NULL);
    sem_destroy(&mutex);
    return 0;
}
```

gcc a.c -lpthread -lrt

## Ex 2:

```
#include <pthread.h>
#include <semaphore.h>
#include <stdio.h>
#include <stdlib.h>

#define NITER 1000000

int cnt = 0;
sem_t mutex;

void * Count(void * a)
{
    int i, tmp;
    for(i = 0; i < NITER; i++)
    {
        sem_wait(&mutex);
        tmp = cnt;      /* copy the global cnt locally */
        tmp = tmp+1;    /* increment the local copy */
        cnt = tmp;      /* store the local value into the global */
        sem_post(&mutex);
    }
}

int main(int argc, char * argv[])
{
    pthread_t tid1, tid2;
    sem_init(&mutex, 0, 1); // mutex
```

```
int main(int argc, char * argv[])
{
    pthread_t tid1, tid2;
    sem_init(&mutex, 0, 1); // mutex

    if(pthread_create(&tid1, NULL, Count, NULL))
    {
        printf("\n ERROR creating thread 1");
        exit(1);
    }

    if(pthread_create(&tid2, NULL, Count, NULL))
    {
        printf("\n ERROR creating thread 2");
        exit(1);
    }

    if(pthread_join(tid1, NULL)) /* wait for th
    {
        printf("\n ERROR joining thread");
        exit(1);
    }

    if(pthread_join(tid2, NULL)) /* wait fo
    {
        printf("\n ERROR joining thread");
        exit(1);
    }
```

# Points to monitor

---


- Monitors can be implemented by semaphores
- OSes support
  - Monitor, semaphore, spinlock, mutex
  - Examples
    - Solaris
    - Windows
    - Linux
    - Pthreads
- Alternative approaches
  - Transactional Memory
  - OpenMP
  - Functional Programming Languages

```
void update(int value)
{
    #pragma omp critical
    {
        count += value
    }
}
```



# Semaphore use: Exercise 1

- Obtain the following precedence graph



`init (S, 0);`

Process i


A  
wait (S);  
C

Process j

B  
signal (S);

# Semaphore use: Exercise 2

```
init (S1, 0);  
init (S2, 0);
```



Process 0

```
...  
i=1  
while (i<=n) {  
    signal (S1);  
    i++;  
}  
...
```

Process i

```
wait (S1);  
...  
signal (S2);  
...
```

Process n+1

```
...  
i=1;  
while (i<=n) {  
    wait (S2);  
    i++;  
}  
...
```

# Semaphore use: Exercise 3

- Realize the represented precedence graph using semaphores


```
init (s1, 0);  
init (s2, 0);
```

```
...  
wait (s1);  
B  
signal (s2);  
...
```

```
A  
signal (s1);  
signal (s1);  
...
```

```
...  
wait (s1);  
B  
signal (s2);  
...
```

```
...  
wait (s2);  
wait (s2);  
D
```





# Summary (1/2)

---

- Important concept: **Atomic Operations**
  - An operation that runs to completion or not at all
  - These are the primitives on which to construct various synchronization primitives
- Talked about hardware atomicity primitives:
  - Disabling of Interrupts, test&set, swap, compare&swap, load-locked & store-conditional
- Showed several constructions of Locks
  - Must be very careful not to waste/tie up machine resources
    - Shouldn't disable interrupts for long
    - Shouldn't spin wait for long
  - Key idea: Separate lock variable, use hardware mechanisms to protect modifications of that variable

# Summary (2/2)

---

- **Semaphores**: Like integers with restricted interface
  - Two operations:
    - **P()**: Wait if zero; decrement when becomes non-zero
    - **V()**: Increment and wake a sleeping task (if exists)
    - Can initialize value to any non-negative value
  - Use separate semaphore for each constraint
- **Monitors**: A lock plus one or more condition variables
  - Always acquire lock before accessing shared data
  - Use condition variables to wait inside critical section
    - Three Operations: **Wait()**, **Signal()**

# Questions?

---

