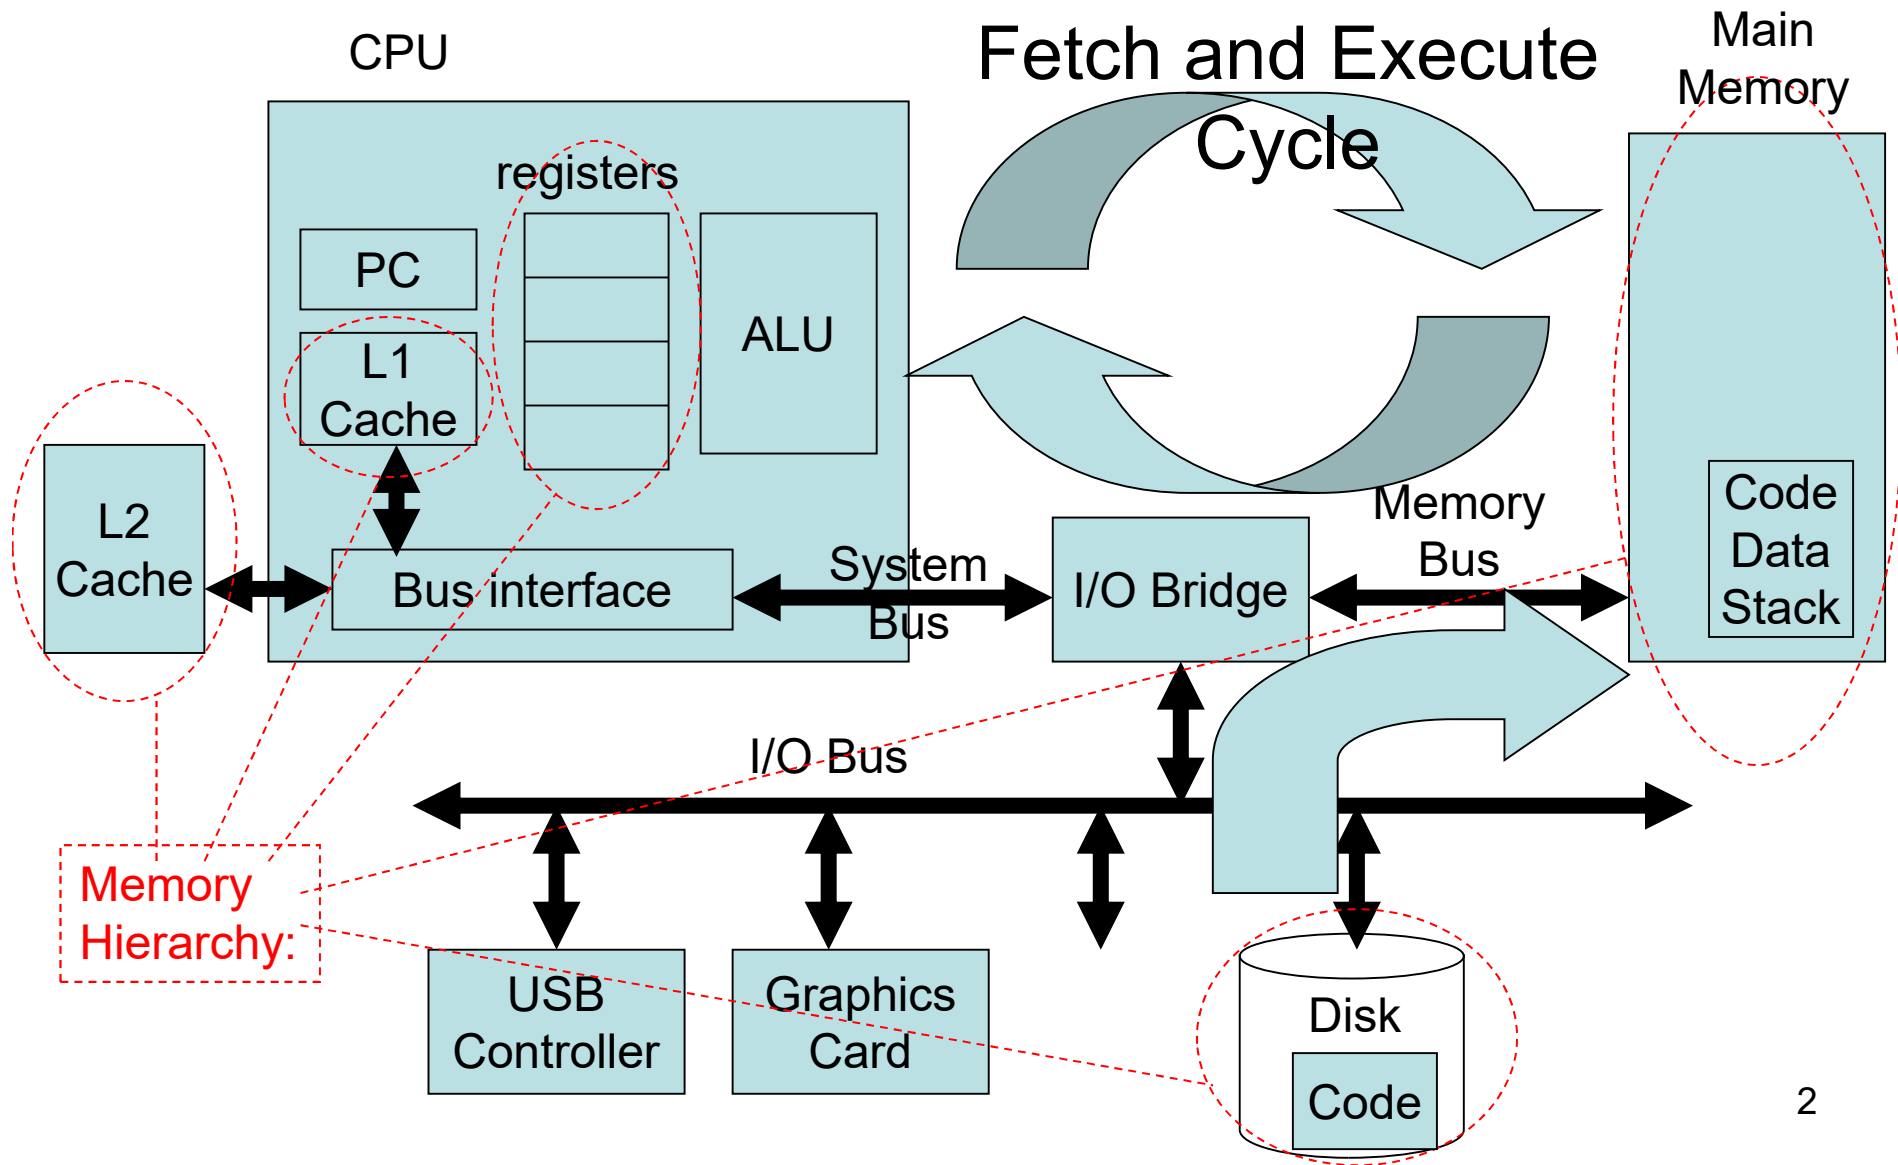




Operating System


Main Memory Management

Memory Management

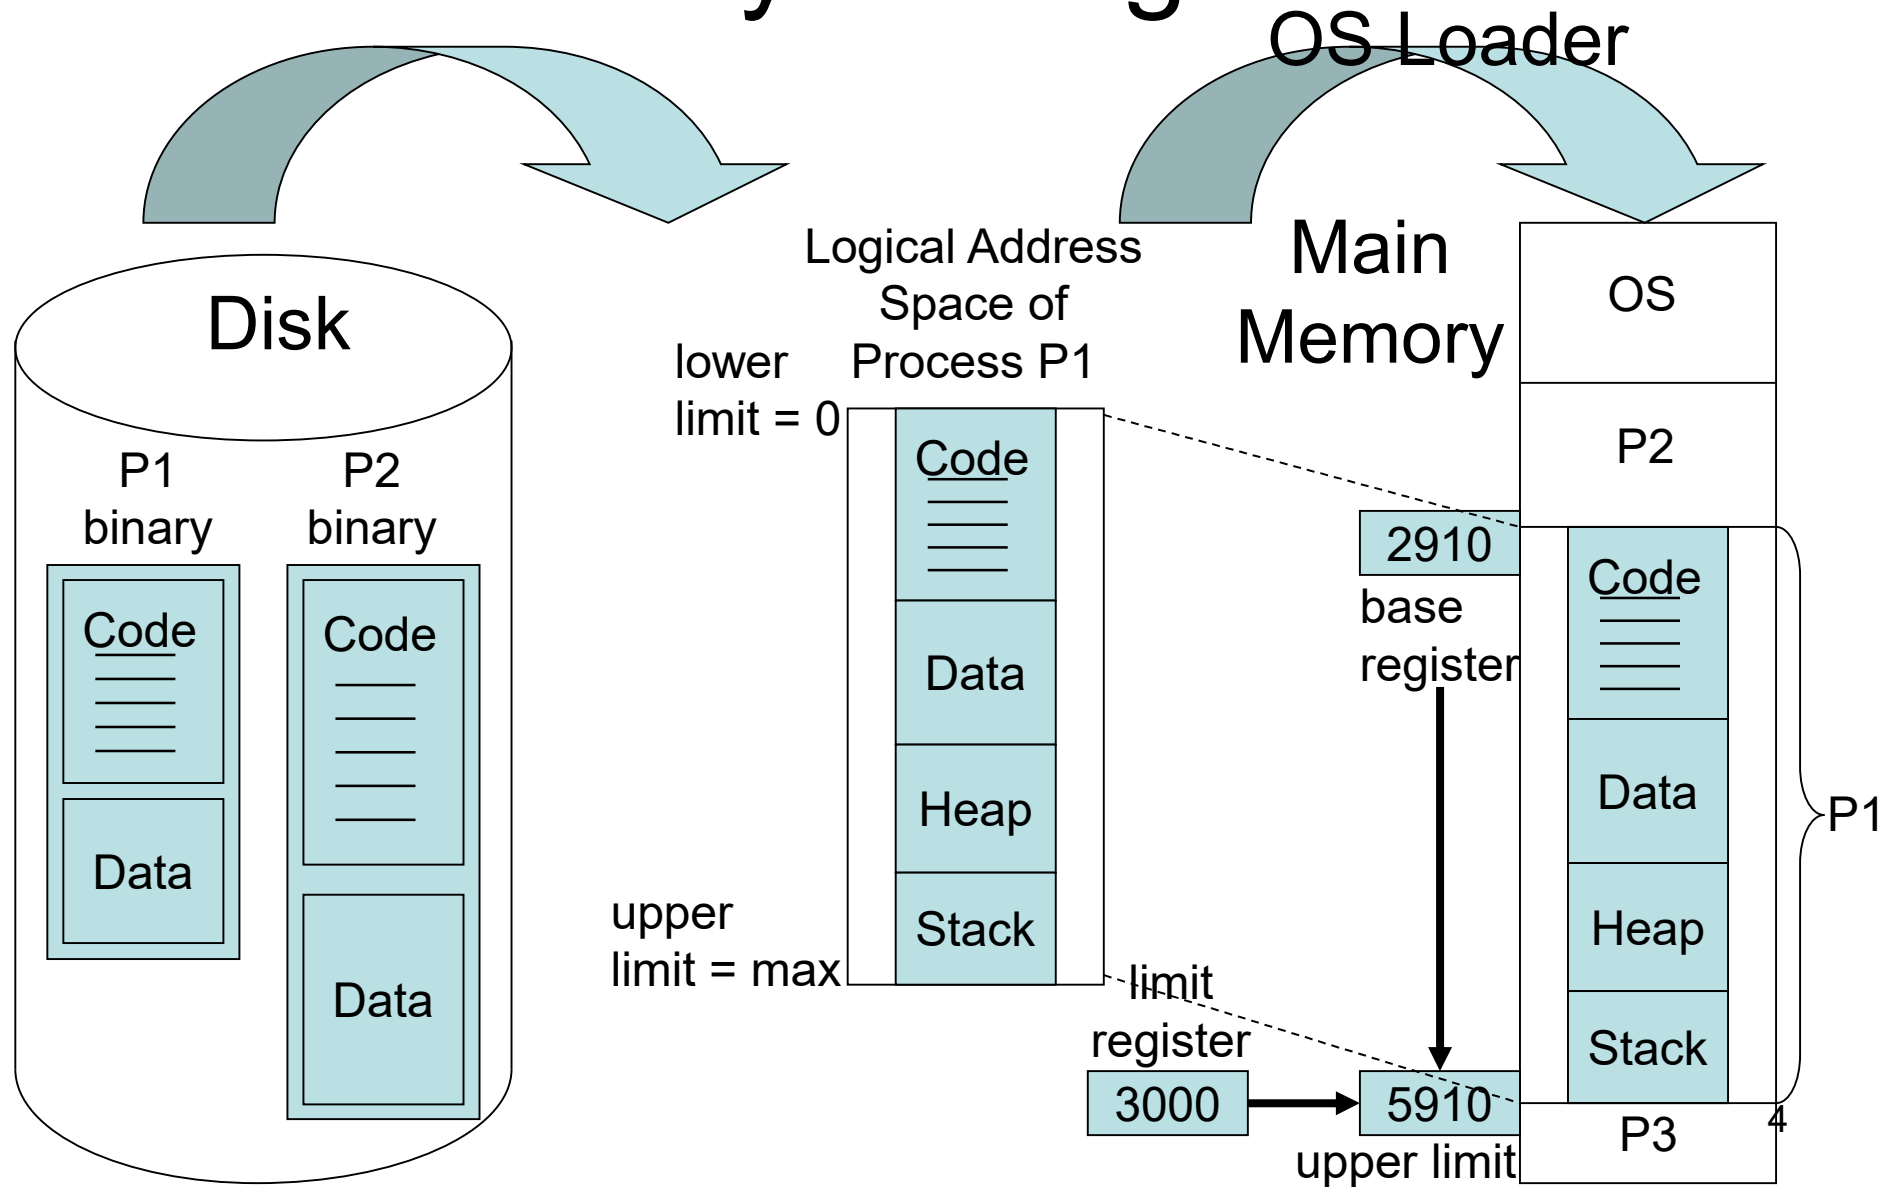


Memory Management

- Memory Hierarchy
 - cache frequently accessed instructions and/or data in local memory that is faster but also more expensive



Memory Management

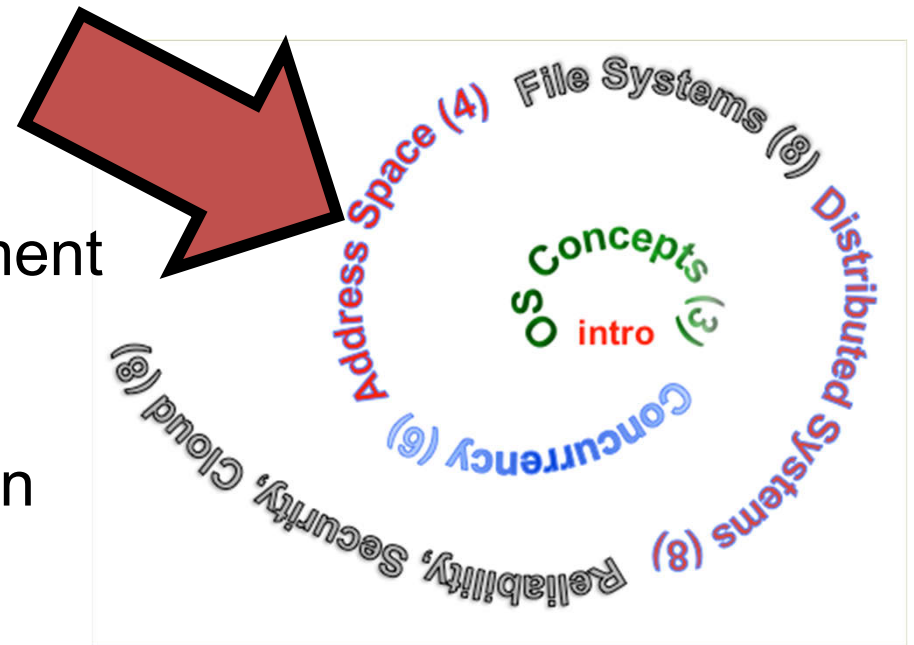


Why memory management?

- CPU utilization and increase of performance
- → need of scheduler to schedule multiple processes
- → they should be resident in RAM (physical memory)
- Why need of memory management?
 - Having multiple processes
 - Protection (processes from one another, processes from OS)
- Protection should be done by Hardware than OS (why?)
 - Performance penalty and latency


Next Objective

- Dive deeper into the concepts and mechanisms of memory sharing and address translation
- Enabler of many key aspects of operating systems
 - Protection
 - Multi-programming
 - Isolation
 - Memory resource management
 - I/O efficiency
 - Sharing
 - Inter-process communication
 - Debugging
 - Demand paging
- Today: Translation

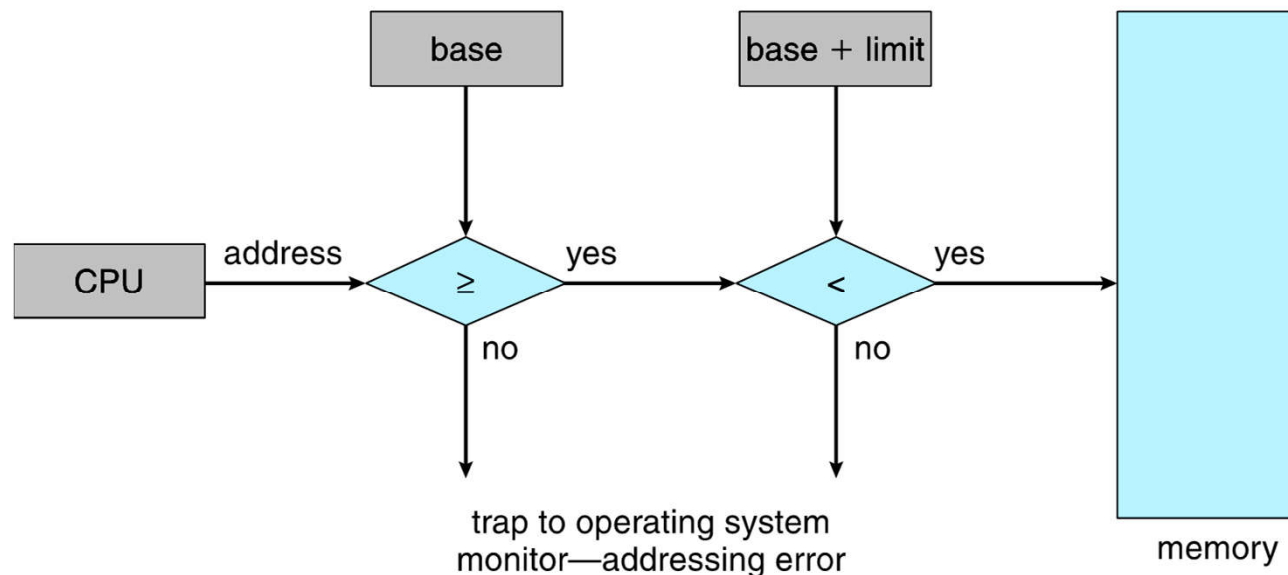


How to protect process memory space?

- Each process has a **separate memory space**
- To protect processes' spaces
 - Determining **legal address**
 - **Base register**: smallest legal physical memory address
 - **Limit register**: size of the range
 - Example:
 - Base register = 300040
 - Limit register = 120900
 - **Legal address space:**
 $\text{Base} \leq (\text{any address}) < \text{Base} + \text{Limit}$
 - Can easily be checked by hardware




HW address protection (base & limit regs.)



Any illegal address generates a **trap exception** which is known as **fatal error**

Who loads Base and Limit registers?

Address binding



Address binding

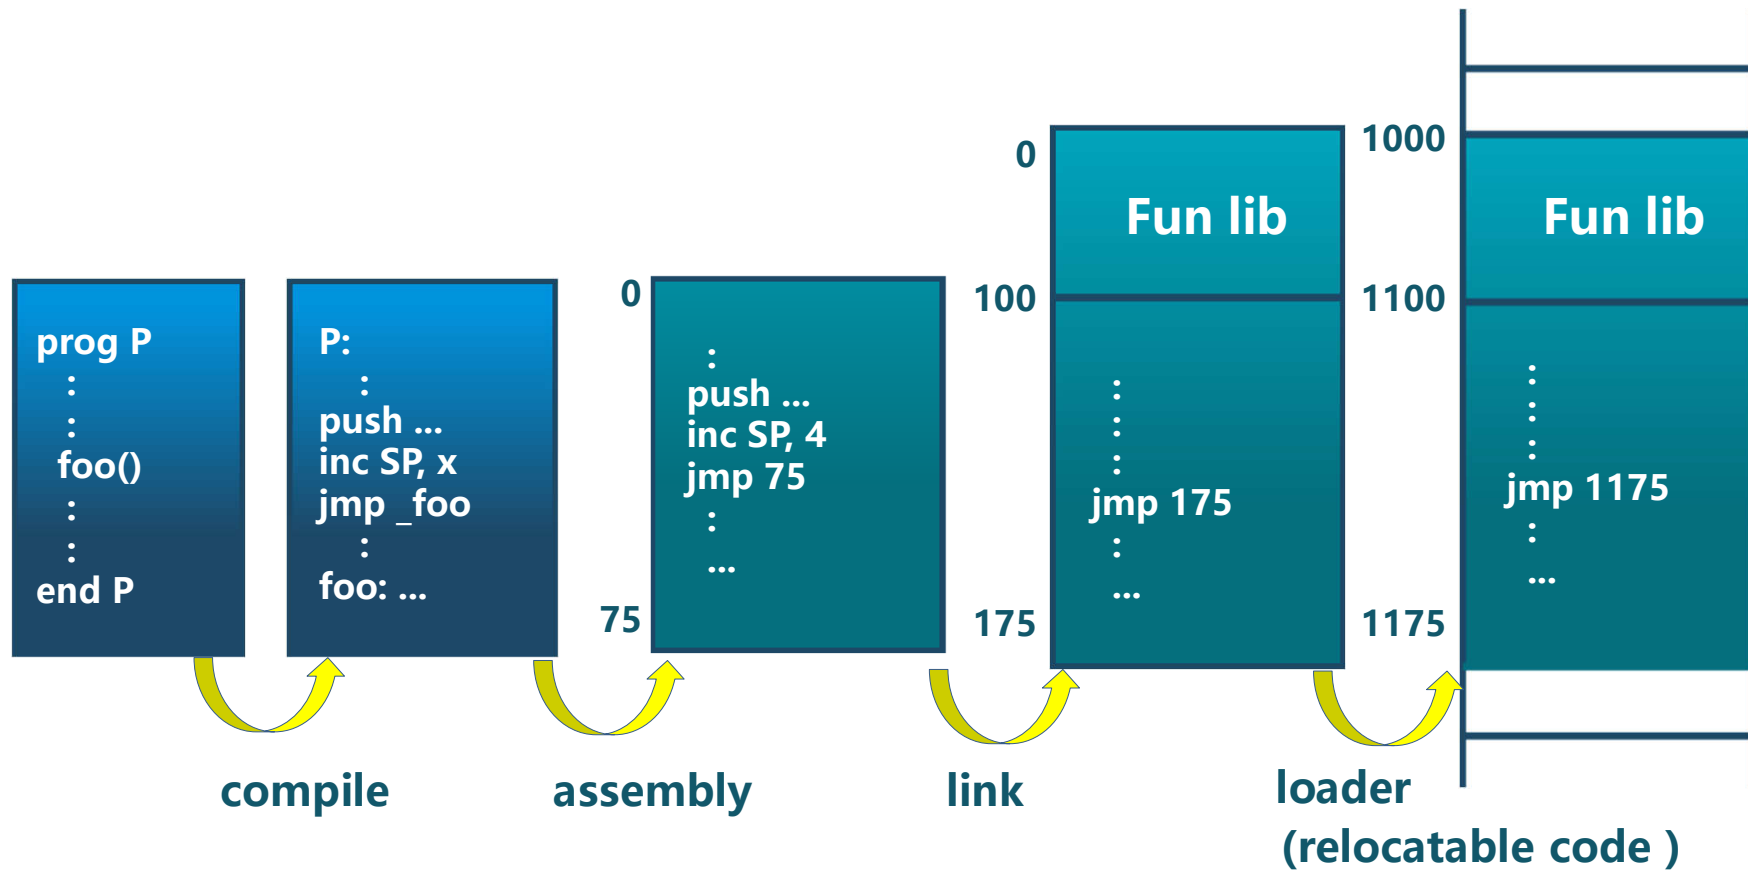
- **Input Queue**

- Processes on disk that are waiting to be brought into memory during execution (Part of **ready queue** which is on **disk**)

- **How to put a process in a physical address?**

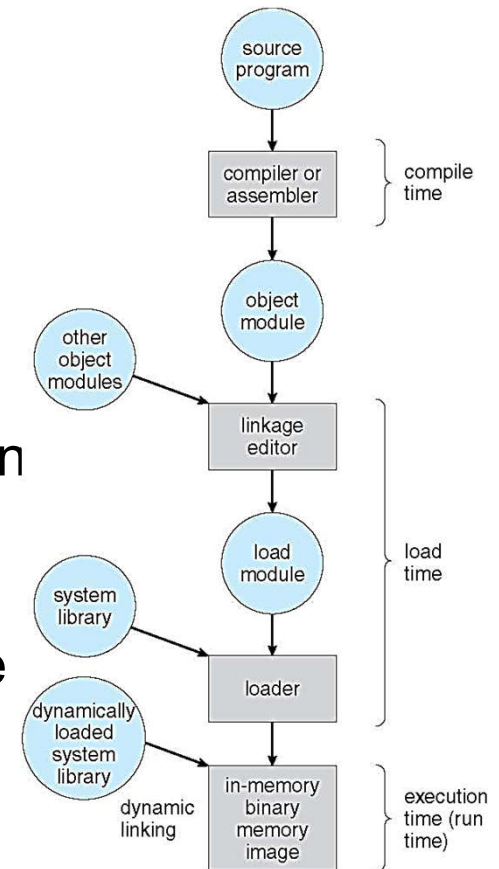
- Addresses in **source** program are **symbolic**
 - Example: **count** variable
- A **compiler** binds them to **relocatable** addresses
 - Example: 14 bytes from beginning of this module
- **Linker** and **loader** bind them next to **absolute** addresses
 - Example: 74014

Generate logic address



How to bind inst./data to mem. address?

- **Compile time**
 - If memory location known **a priori**, **absolute code** can be generated; must recompile code if starting location changes
 - Example: **COM files** in MS-DOS
- **Load time**
 - Must generate **relocatable code** if memory location is not known at compile time
- **Execution time**
 - Binding delayed until run time if the process can be moved during its execution from one memory segment to another
 - **Need hardware support for address maps**
 - Example: Base-Limit registers





Compile Time	Load Time	Execution Time
Code is translated into machine readable format	Program code is loaded into memory	Program code is executed
Completed with a compiler	Completed with a loader	Completed with run-time libraries

Logical vs. physical address space

- **Logical address (CPU address)**
 - CPU logically sees addresses
 - **Logical address space**: set of all **logical addresses** generated by a program
- **Physical address (Memory address)**
 - Address of memory line
 - **Physical address space**: a physical address (also real address, or binary address)
 - The set of all physical addresses corresponding to these logical addresses is a physical address space.
- May be equal or not
 - **Equal**: **compile-time** and **load-time** address bindings
 - **Not equal**: **execution-time** address binding
 - In this case, **logical address** is said **virtual address**

What is MMU (Memory Management Unit)?

- MMU is responsible of mapping **virtual (logical) address** to **physical address**
- Simple version: **Base-Limit** registers
- Here **Base register** is called **relocation register**





Dynamic linking and loading



Dynamic loading

- **Dynamic loading**
 - A routine is **not loaded** until it is **called**
 - All routines except *main()* are kept on disk
 - In the case of call, if it is not loaded, linking loader first loads it to the memory and update program's address table to reflect this change

Pros:

- **Better** space utilization
- Some routines are **infrequently** needed: *error functions*

Dynamic linking


- **Dynamic linking (Dynamically linked libraries)**
 - System libraries that are **linked** to user programs when they are **run**
 - **No waste** of memory and disk space
 - Example: **Language subroutines**
 - **Stub**: small piece of code used to locate dynamic linked libraries
 - Replace itself with the address of the routine & executes the routine
 - Also known as **Shared Libraries**
- **Static linking**
 - System libraries are linked to user programs during **compile time**



Swapping

Swapping

- **Swapping** is performed temporarily between memory and backing store
- **Backing store** — fast disk large enough to accommodate copies of all memory images for all users; must provide direct access to these memory images
- Possible of total process memory spaces exceeds real physical memory
→ Use of Backing store



Swapping issues

- Does the swapped out process need to swap back in to same physical addresses?
 - Depends on address binding method
 - Plus consider pending I/O to/from process memory space
- Pending IO: cannot swap out as IO would occur to wrong swapped in process!
 - Or always transfer IO to kernel space, then to IO device; this is known as **double buffering**; adds **overheads**

Swapping and modern OSs

- Standard swapping is not used in Linux & Windows!
 - Modified version is used:
 - Swap only when free memory is extremely low (less than threshold)
 - Disable swapping when free memory is more than threshold

Swapping cost

- Major part of swap time is **transfer time**
- **Total transfer time** is **proportional** to the **amount** of memory swapped.
- Example:
 - **100 MB** process swapping to hard disk with transfer rate of **50 MB/sec**
 - Swap out time of **2(sec)+swap in** of same size process
 - Total context switch swapping component time of **4 sec.**



Memory allocation



Memory allocation

- Different types

1. Contiguous

- Each process is in a **single section** of memory that is **contiguous** to sections of **others**

2. Segmentation

- Each process is divided into different **segments**; each one is located in **different** part

3. Paging

- Each process is divided into **same–small–size pages**; **some** of them are swapped in/out

Criteria and problems

- CPU Utilization

- Utilization = Percent of time a CPU is busy

$$= \frac{\text{CPU time used}}{\text{Total time system is busy}}$$

Some times is $\frac{\text{CPU time used}}{\text{CPU time used} + 2 * \text{swapping time}}$

Some times is $\frac{\text{CPU time used}}{\max(\text{CPU time used}, 2 * \text{swapping time})}$

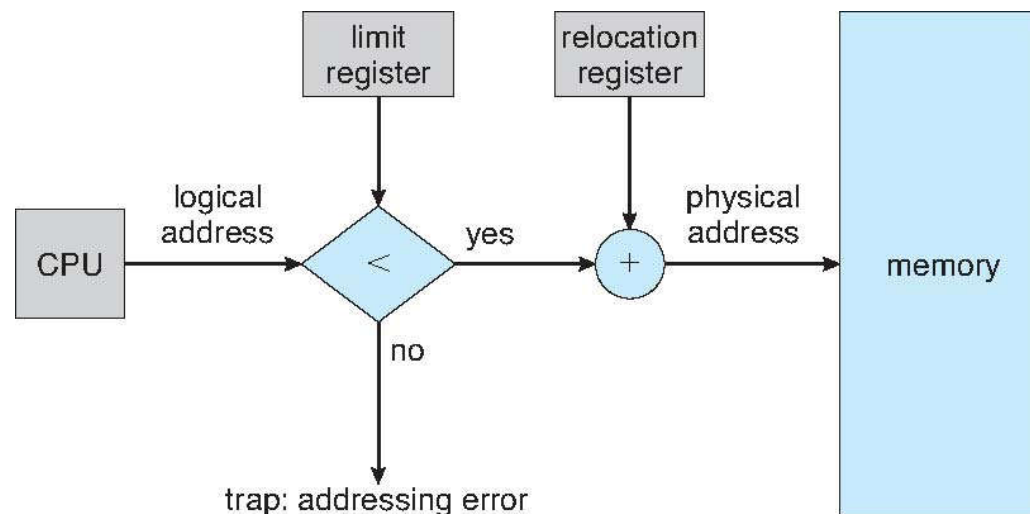
- Fragmentation

- External

- Internal

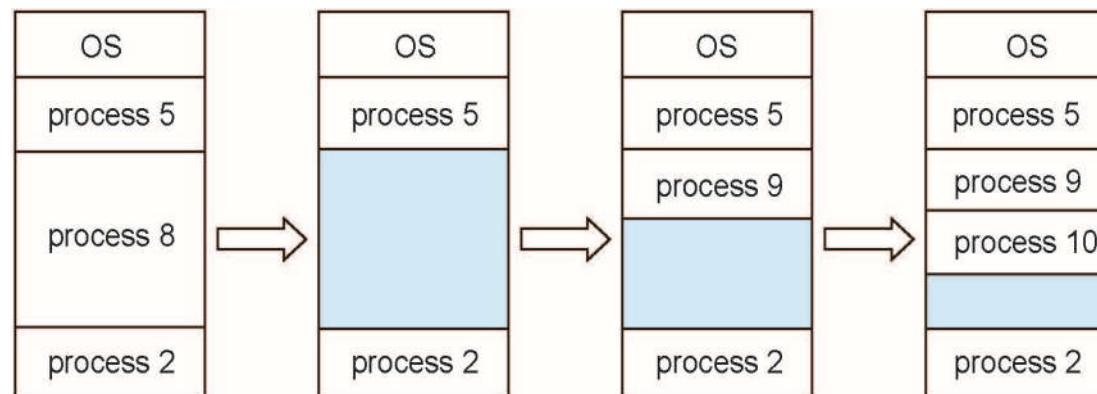
Memory Protection

- **Relocation registers** used to protect user processes from each other, and from changing operating-system code and data
 - **Base register** contains value of **smallest physical address**
 - **Limit register** contains **range of logical addresses**
 - Each logical address must be less than the limit register



1) Contiguous allocation

- Multiple-partition memory allocation
 - Degree of multiprogramming limited by number of partitions
 - **Variable-partition** sizes for efficiency (sized to a given process' needs)
 - **Hole** – block of available memory; holes of various size are scattered throughout memory
 - When a process arrives, it is allocated memory from a hole large enough to accommodate it
 - Operating system maintains information about:
 - a) allocated partitions
 - b) free partitions (**hole**)



Dynamic storage-allocation problem

- How to satisfy a request of size *n* from a list of free holes?
 - **First-fit**: Allocate the *first* hole that is big enough
 - **Best-fit**: Allocate the *smallest* hole that is big enough; must search entire list, unless ordered by size
 - Produces the smallest leftover hole
 - **Worst-fit**: Allocate the *largest* hole; must also search entire list
 - Produces the largest leftover hole

First-fit and best-fit better than worst-fit in terms of speed and storage utilization

Fragmentation problem

- **External Fragmentation**

- Total memory space exists to satisfy a request, but it is not contiguous

- **Internal Fragmentation**

- Not whole of allocated memory is filled by the process memory

- **First fit** analysis reveals that given N blocks allocated, $0.5 N$ blocks lost to fragmentation

- $1/3$ may be unusable -> **50-percent rule**

Fragmentation solution

- **Compaction**


- Shuffle memory contents to place all free memory together in one large block
- Compaction is possible *only* if relocation is *dynamic*, and is done at *execution time*
- I/O problem
 - Latch job in memory while it is involved in I/O
 - Do I/O only into OS buffers

- Now consider that backing store has same fragmentation problems

solution to the external-fragmentation problem is to permit the logical address space of the processes to be noncontiguous


2) Segmentation

- Memory-management scheme that supports **user view of memory**
- A program is a collection of **segments**
 - A segment is a logical unit such as:
 - main program
 - procedure
 - function
 - method
 - object
 - local variables, global variables
 - common block
 - stack
 - symbol table
 - arrays



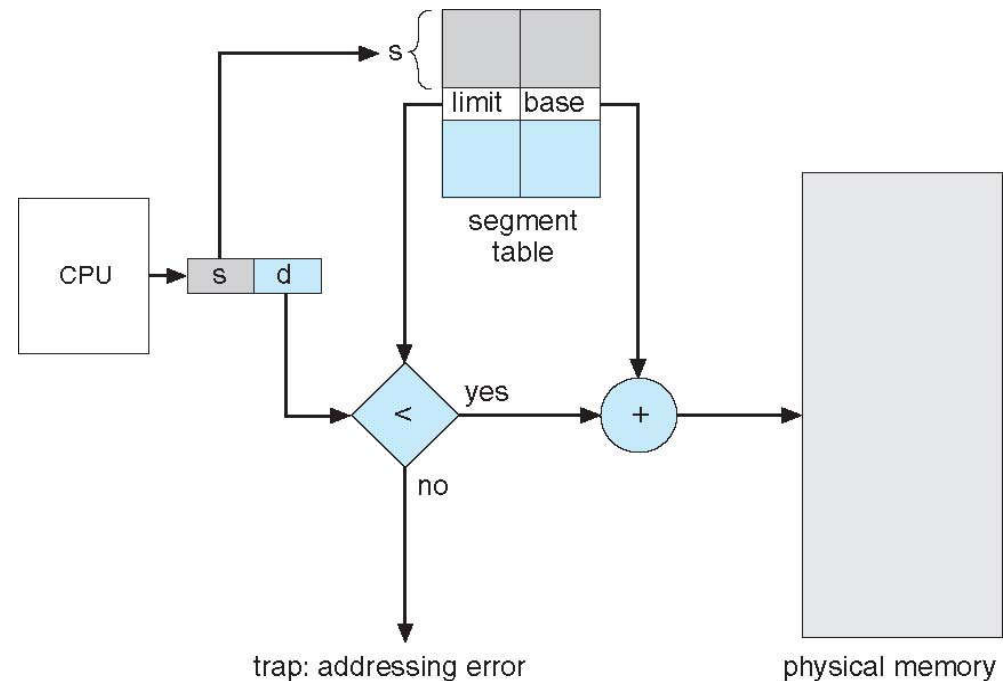
More Flexible Segmentation

- Logical View: multiple separate segments
 - Typical: Code, Data, Stack
 - Others: memory sharing, etc
- Each segment is given region of contiguous memory
 - Has a base and limit
 - Can reside anywhere in physical memory

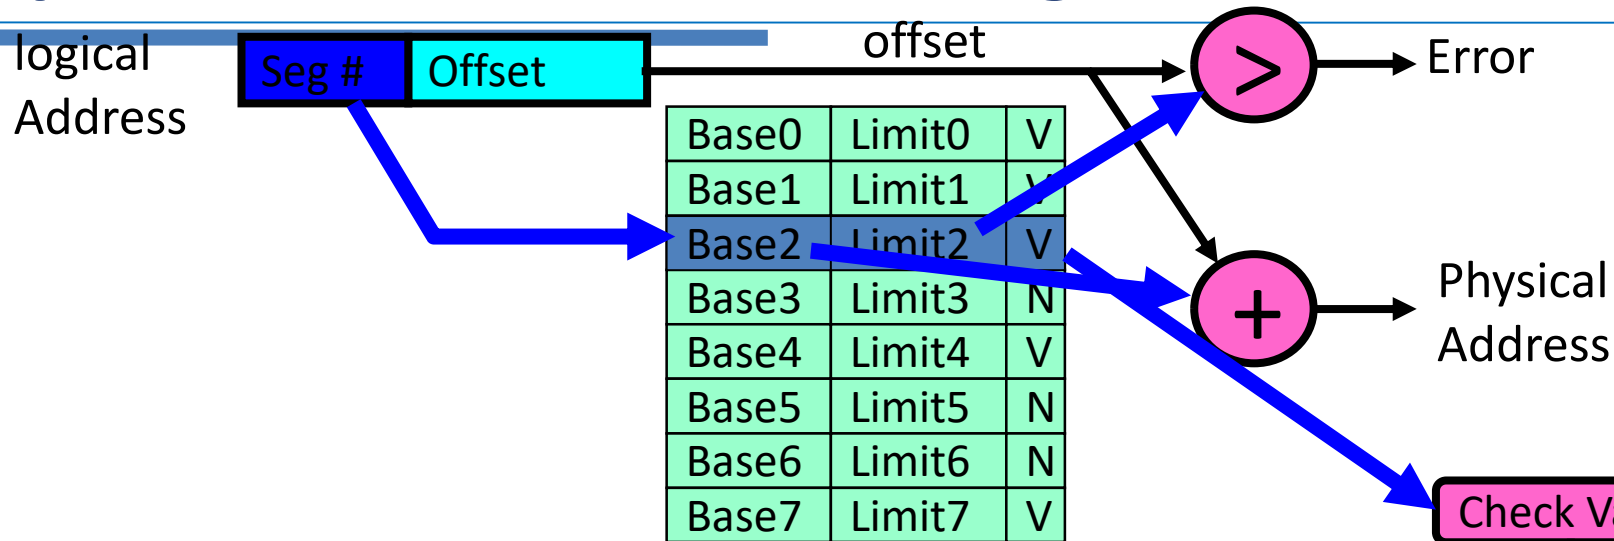


Segmentation implementation

- Logical address consists of a two tuple:
<segment-number, offset>
- **Segment table**
 - Maps **two-dimensional** physical addresses;
 - each table entry has:
 - **base** – contains the starting physical address where the segments reside in memory
 - **limit** – specifies the length of the segment



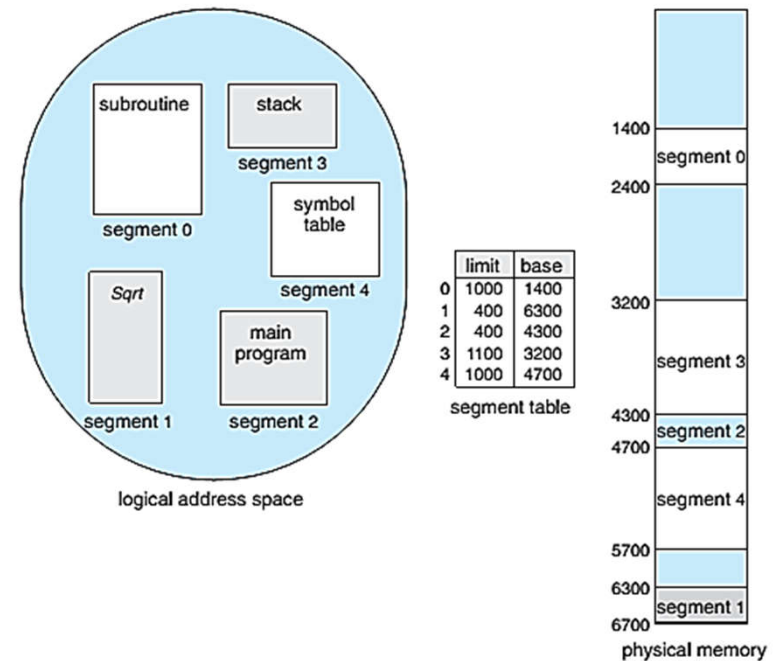
Implementation of Multi-Segment Model



- Segment map resides in processor
 - Segment number mapped into base/limit pair
 - Base added to offset to generate physical address
 - Error check catches offset out of range
- As many chunks of physical memory as entries
 - Segment addressed by portion of logical address
 - However, could be included in instruction instead:
 - x86 Example: `mov [es:bx],ax.`
- What is “V/N” (valid / not valid)?
 - Can mark segments as invalid; requires check as well

Example of segmentation

- A reference to byte 53 of segment 2:
 $4300 + 53 = 4353$
- A reference to byte 852 of segment 3:
 $3200 + 852 = 4052$
- A reference to byte 1222 of segment 0:
Trap to OS!




Example: Four Segments (16 bit addresses)




Virtual Address Format

Seg ID #	Base	Limit
0 (code)	0x4000	0x0800
1 (data)	0x4800	0x1400
2 (shared)	0xF000	0x1000
3 (stack)	0x0000	0x3000



Virtual
Address Space



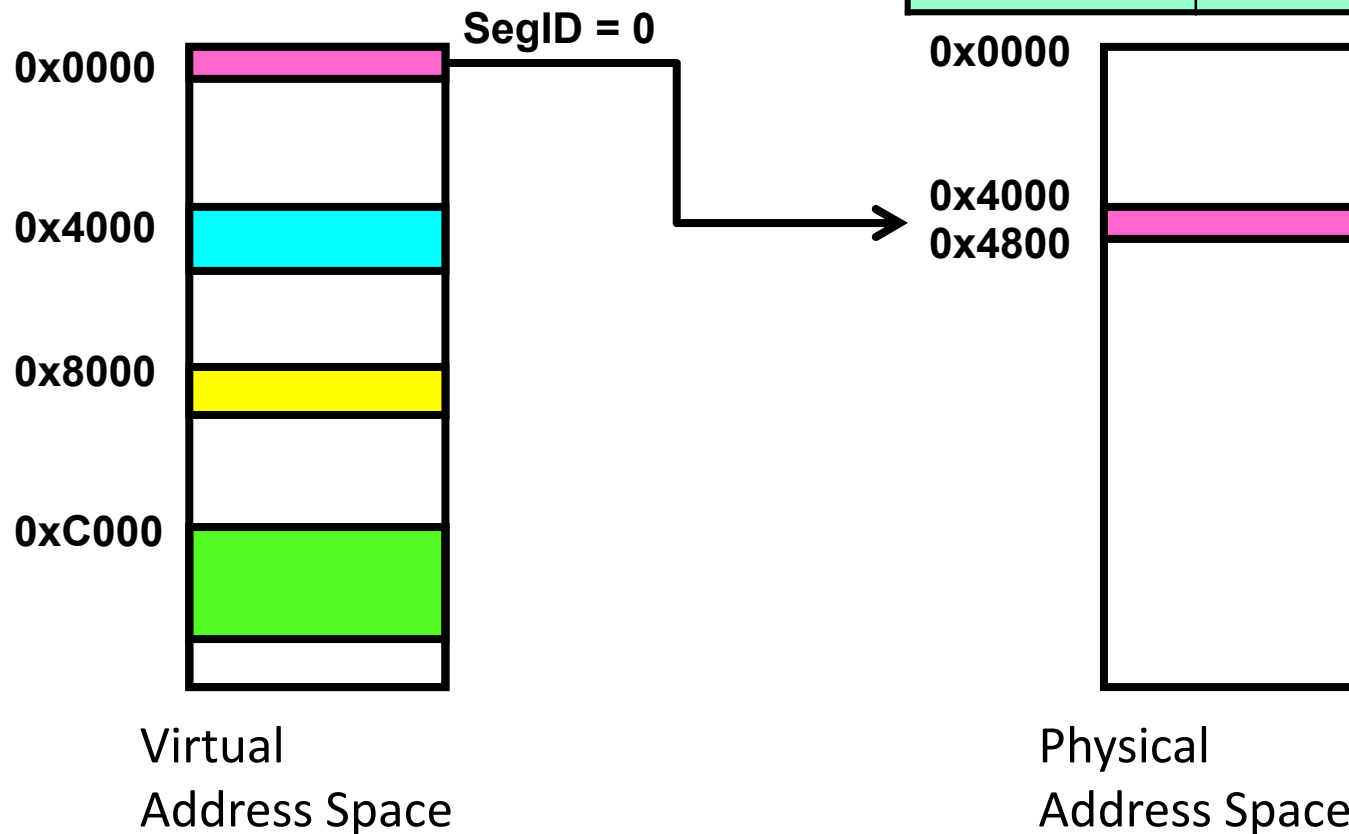
Physical
Address Space

Example: Four Segments (16 bit addresses)



Virtual Address Format

Seg ID #	Base	Limit
0 (code)	0x4000	0x0800
1 (data)	0x4800	0x1400
2 (shared)	0xF000	0x1000
3 (stack)	0x0000	0x3000

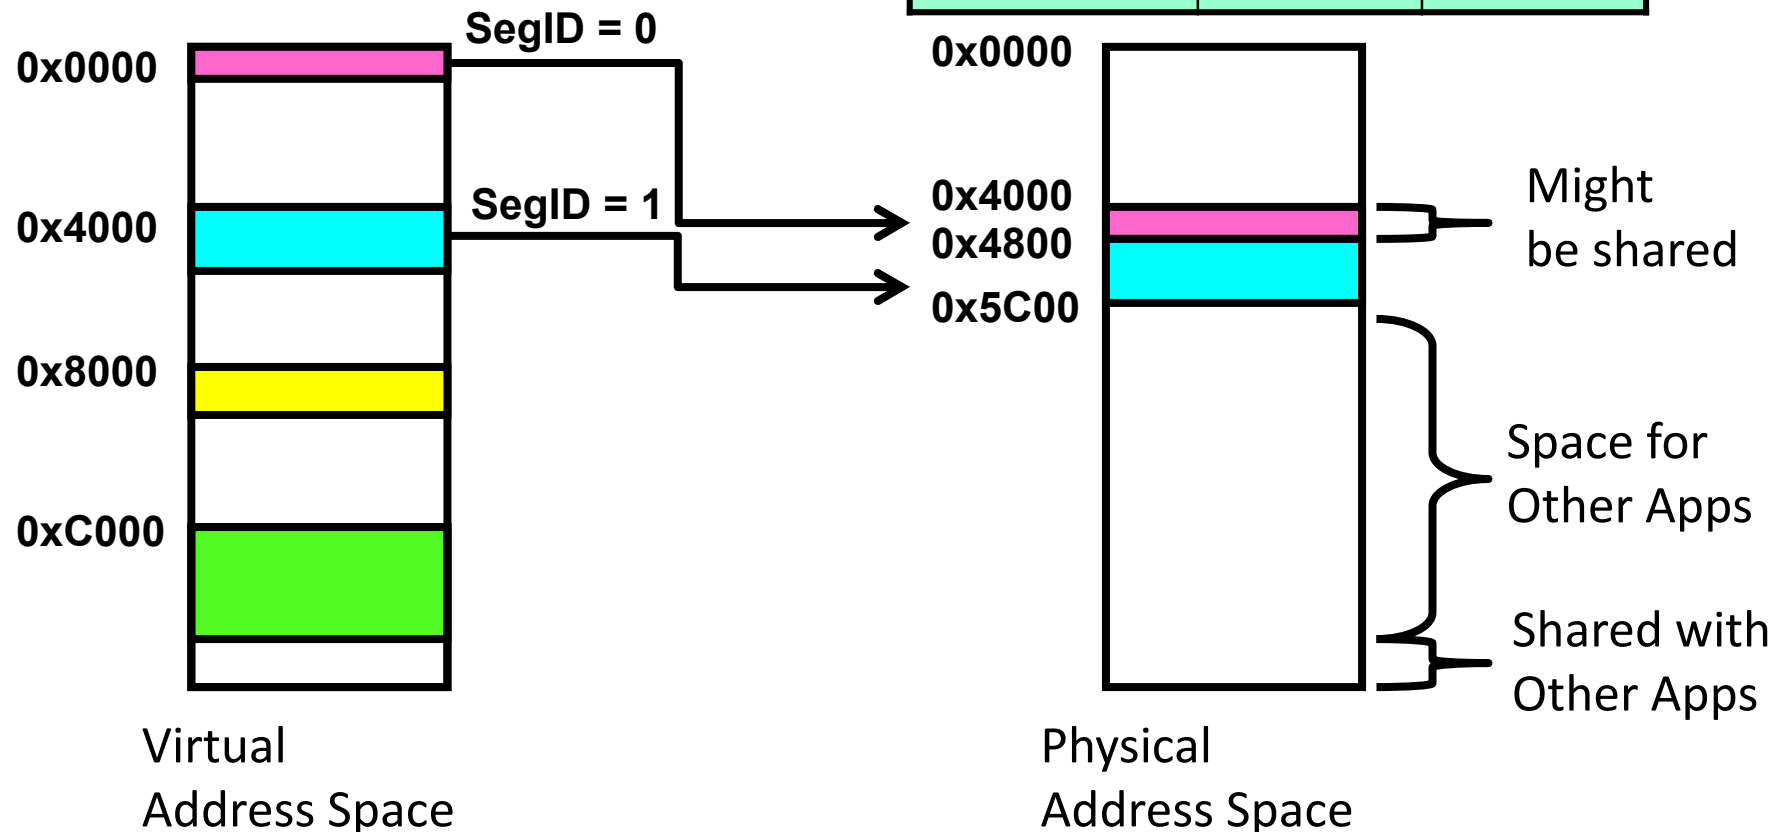


Example: Four Segments (16 bit addresses)



Virtual Address Format

Seg ID #	Base	Limit
0 (code)	0x4000	0x0800
1 (data)	0x4800	0x1400
2 (shared)	0xF000	0x1000
3 (stack)	0x0000	0x3000





Example: Four Segments (16 bit addresses)



Virtual Address Format

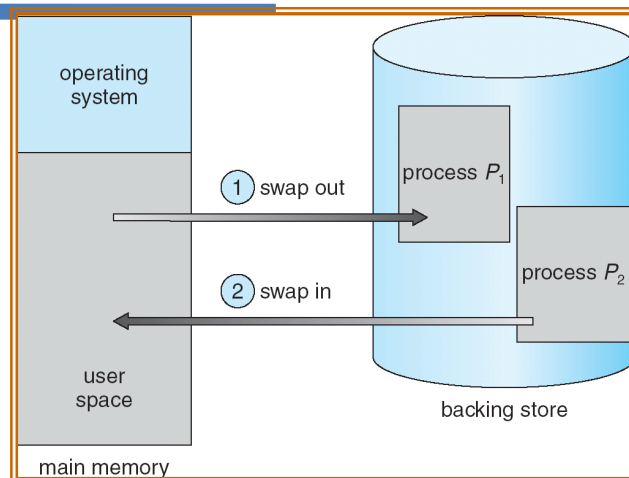
Seg ID #	Base	Limit
0 (code)	0x4000	0x0800
1 (data)	0x4800	0x1400
2 (shared)	0xF000	0x1000
3 (stack)	0x0000	0x3000



Observations about Segmentation

- Virtual address space has holes
 - Segmentation efficient for sparse address spaces
 - A correct program should never address gaps (except as mentioned in moment)
 - If it does, trap to kernel and dump core
- When it is OK to address outside valid range?
 - This is how the stack and heap are allowed to grow
 - For instance, stack takes fault, system automatically increases size of stack
- Need protection mode in segment table
 - For example, code segment would be read-only
 - Data and stack would be read-write (stores allowed)
 - Shared segment could be read-only or read-write
- What must be saved/restored on context switch?
 - Segment table stored in CPU, not in memory (small)

What if not all segments fit into memory?



- Extreme form of Context Switch: Swapping
 - In order to make room for next process, some or all of the previous process is moved to disk
 - Likely need to send out complete segments
 - This greatly increases the cost of context-switching
- What might be a desirable alternative?
 - Some way to keep only active portions of a process in memory at any one time
 - Need finer granularity control over physical memory

Problems with Segmentation

- Must fit variable-sized chunks into physical memory
- May move processes multiple times to fit everything
- Limited options for swapping to disk
- **Fragmentation**: wasted space
 - **External**: free gaps between allocated chunks
 - **Internal**: don't need all memory within allocated chunks

3) Paging


- Noncontiguous memory allocations:
 - Segmentation
 - Paging
- Paging avoids external fragmentation, and need of compaction, whereas segmentation does not.
- Process is allocated physical memory whenever physical memory is available
 - Avoid external fragmentation
 - Avoid problem of varying sized memory chunks
- Divide physical memory into fixed-sized blocks called frames.
 - Size is power of 2, between 512 bytes to 16 Mbytes
- Divide logical memory into blocks of same size called pages.

Paging

- Keep tracks of all free frames
- To run a program with N pages, need to find N different free frames and load program.
- Setup a page table to translate logical to physical address
- Still have internal fragmentation

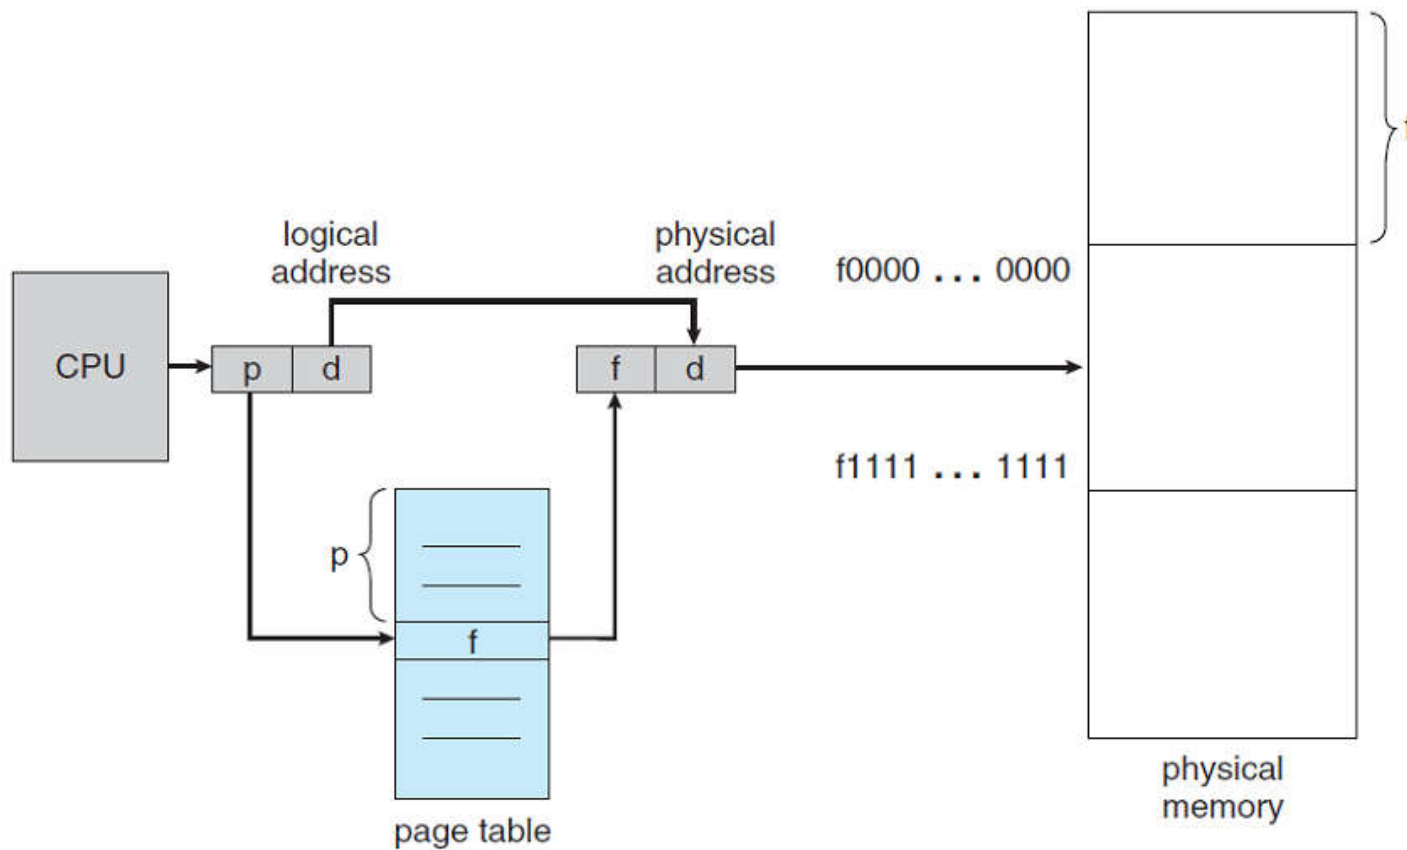
Address translation scheme

- CPU address (Logical address) is divided into **two parts**:
 - **Page number** (p): used as an index into a page table that contains **base register of each page** in physical memory
 - **Page offset** (d): combined with base address to define the physical memory address that is sent to the memory unit.

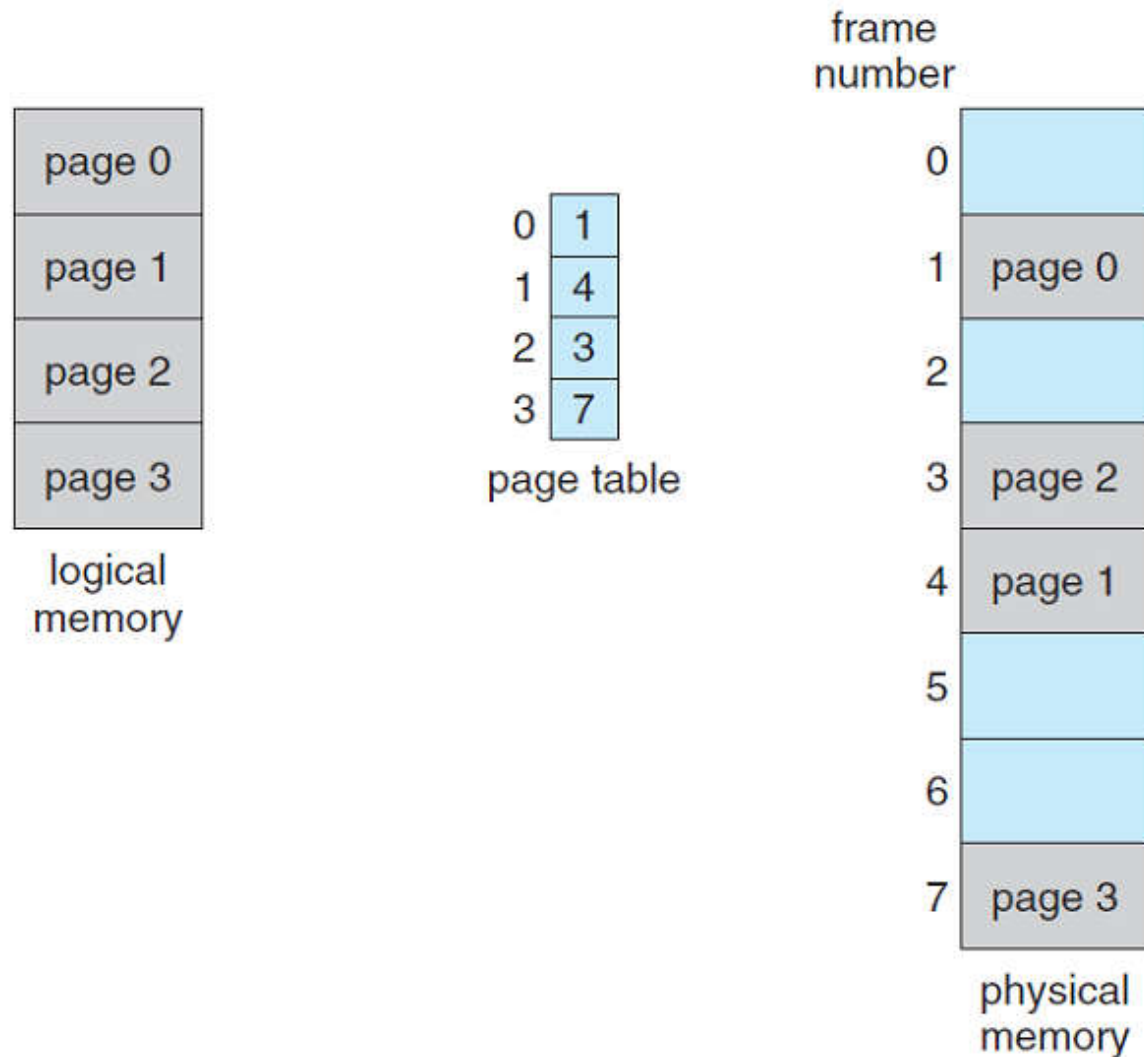


- For given logical address space 2^m and page size 2^n

Paging hardware



Paging model of logical and physical memory



Paging example

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

logical memory

0	5
1	6
2	1
3	2

page table

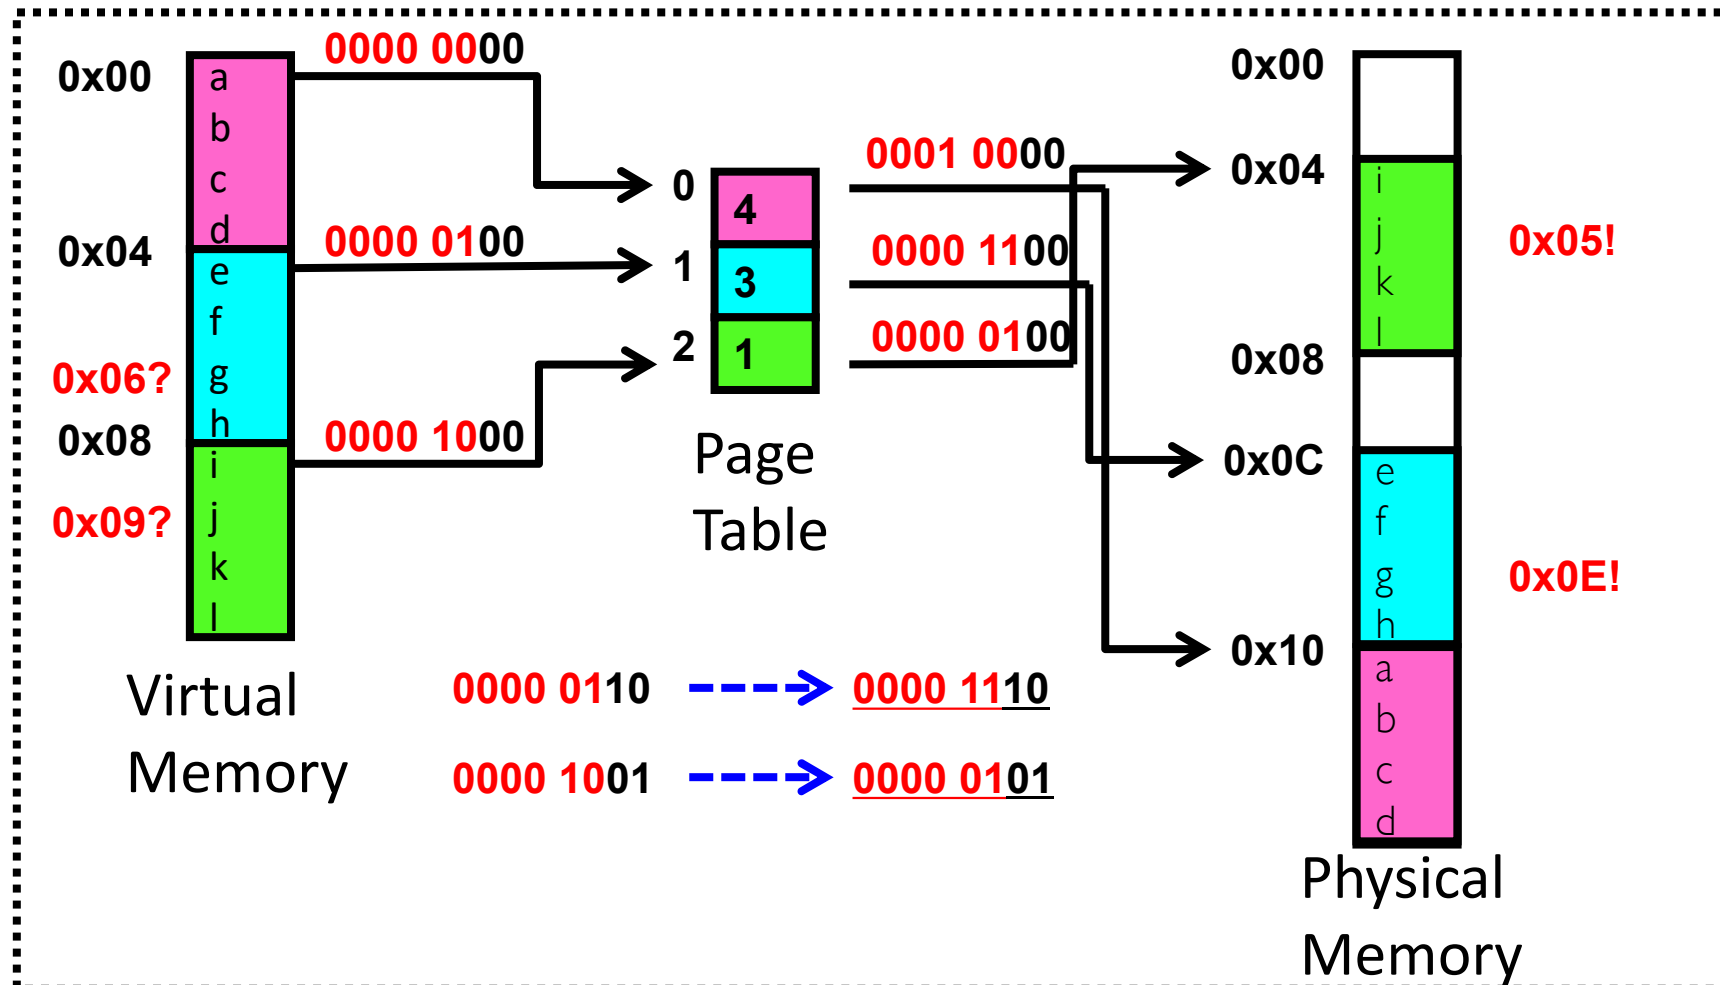
- Logical address: $n=2$ and $m=4$
- Physical memory: 32 byte memory and 4-byte pages.

0	
4	i j k l
8	m n o p
12	
16	
20	a b c d
24	e f g h
28	

physical memory

Simple Page Table Example



Example (4 byte pages)



Free frames


free-frame list

14
13
18
20
15




before allocation

free-frame list
15



0	14
1	13
2	18
3	20

new-process page table



after allocation

Paging example – Internal fragmentation

- Page size = 2048 bytes (2kB)
- Process size = 72766 bytes
- $72766/2048 = 35 \text{ pages} + 1086 \text{ bytes}$
- Internal fragmentation: $2048 - 1086 = 962 \text{ bytes}$
- Worst case fragmentation: 1 frame – 1 byte
- On average fragmentation: $\frac{1}{2}$ frame size

Small page size vs big page size

- On average fragmentation: $\frac{1}{2}$ page size, small page size are good.
- Small page size, more overhead is in the page-table, this overhead is reduced when page size increases.
- Disk I/O is more efficient when the amount of data being transferred is larger (e.g. big pages)
- Page typically are between 4 kB and 8 kB in size.



Page table implementation



Page table

- In the simplest case, the page table is implemented as a set of **dedicated registers**
- **Page table** is kept in **memory**
- **Page-table base register (PTBR)** points to the **page table**
- **Page-table length register (PTLR)** indicates **size of the page table**
- In this scheme every data/instruction access requires **two memory accesses**:
 - **One** for the page table, **another** for data/instruction

Translation look-aside buffer (TLB)


- The **two memory access problem** can be solved by the use of a special fast-lookup hardware cache called **associative memory** or **translation look-aside buffer (TLB)**.
- **Associative memory**: parallel search

Page #	Frame #

- Address translation (p, d) :
 - If p is in associative memory, get $frame\#$ out
 - Otherwise, get $frame\#$ from page table in memory

Paging hardware with TLB

- The standard solution to this problem is to use a special, small, fast lookup hardware cache called a **translation look-aside buffer (TLB)**.



Effective access time

- **Hit ratio**: percentage of times that a page number is found in the TLB.

Effective Access Time (EAT)

- α : memory access latency
- h : hit ratio
- $EAT = h \times \alpha + (1 - h) \times 2\alpha$

$$h = 80\%, \alpha = 100ns \Rightarrow EAT = 0.80 \times 100 + 0.20 \times 200 = 120ns$$

$$h = 99\%, \alpha = 100ns \Rightarrow EAT = 0.99 \times 100 + 0.01 \times 200 = 101ns$$

More about TLB

- Some TLBs store **address-space identifier (ASID)** in each TLB entry
 - **Uniquely identifies each process** to provide address-space protection for that process
 - Otherwise, need to **flush** at every context switch
- TLB is typically **small** (64 to 1024 entries)
- On a **TLB miss**, value is **loaded** into the TLB for faster access next time.
 - **Replacement policies** must be considered.

Memory protection

- Memory protection is implemented by **protection bit** with each frame to indicate if read-only or read-write access is allowed.
- **Valid-invalid bit** attached to **each entry** in page table:
 - **Valid** indicates that the page is in the process logical address space (legal page)
 - **Invalid** indicates that the page is not in the process logical address space (illegal page)
 - Or use **page-table length register (PTLR)**
- Any violation result in **trap to the kernel**

Valid/invalid bit in a page table

14-bit address space(0 to 16383)
a page size: 2 KB

00000	page 0
	page 1
	page 2
	page 3
	page 4
10,468	page 5
12,287	

frame number		valid-invalid bit
0	2	v
1	3	v
2	4	v
3	7	v
4	8	v
5	9	v
6	0	i
7	0	i

page table

0	
1	
2	page 0
3	page 1
4	page 2
5	
6	
7	page 3
8	page 4
9	page 5
	⋮
	page <i>n</i>

Shared pages

- **Shared code**
 - One copy of **read-only (reentrant) code** shared among processes (e.g., text editors)
 - Similar to multiple threads sharing the same process space.
- **Private code and data**
 - Each process keeps a **separate copy of the code and data**
 - The page for the **private code** and data can appear **anywhere** in the logical address space

What about Sharing?

Virtual Address
(Process A):

Virtual Page #	Offset
----------------	--------

PageTablePtrA

page #0	V,R
page #1	V,R
page #2	V,R,W
page #3	V,R,W
page #4	N
page #5	V,R,W

PageTablePtrB

page #0	V,R
page #1	N
page #2	V,R,W
page #3	N
page #4	V,R
page #5	V,R,W

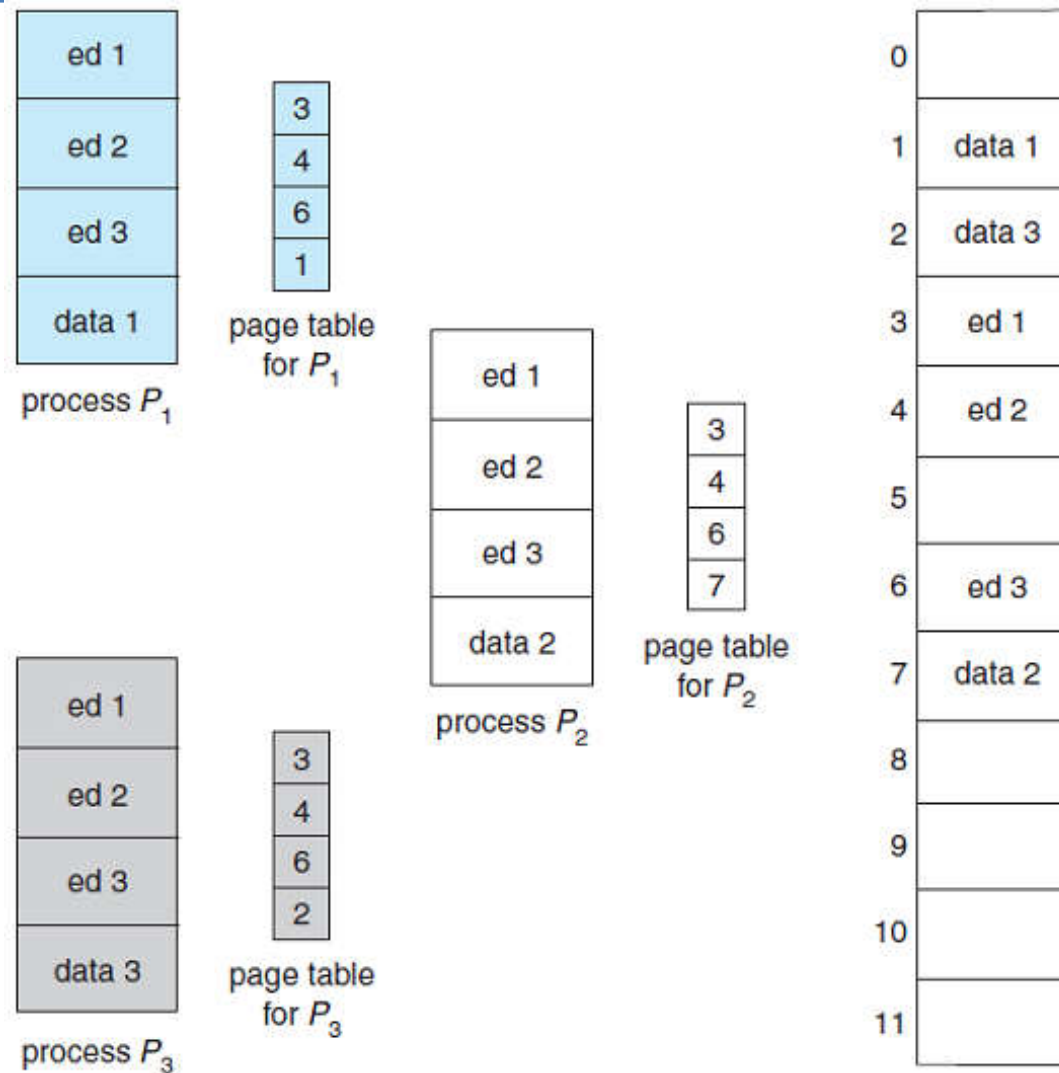
Virtual Address
(Process B):

Virtual Page #	Offset
----------------	--------

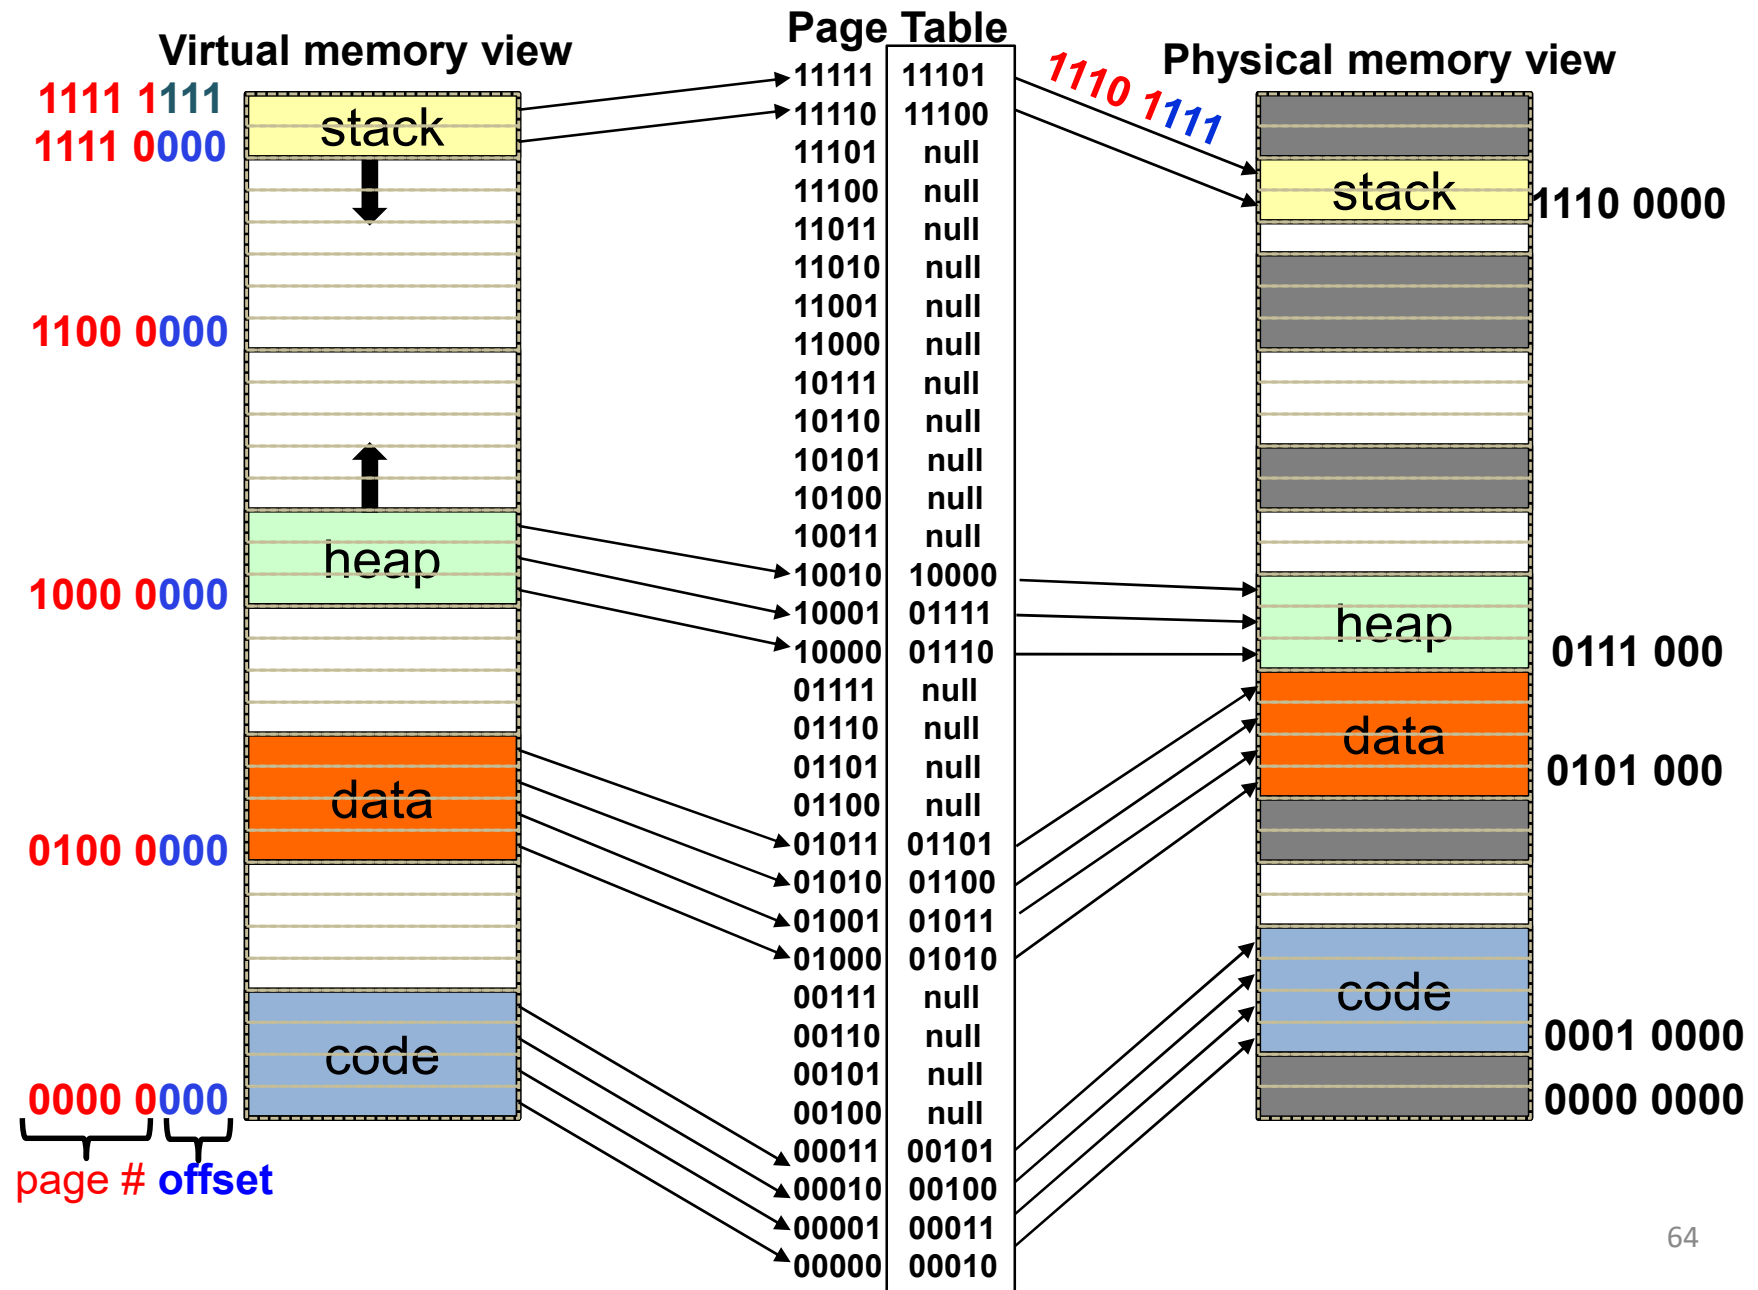
Shared
Page

This physical page
appears in address
space of both processes

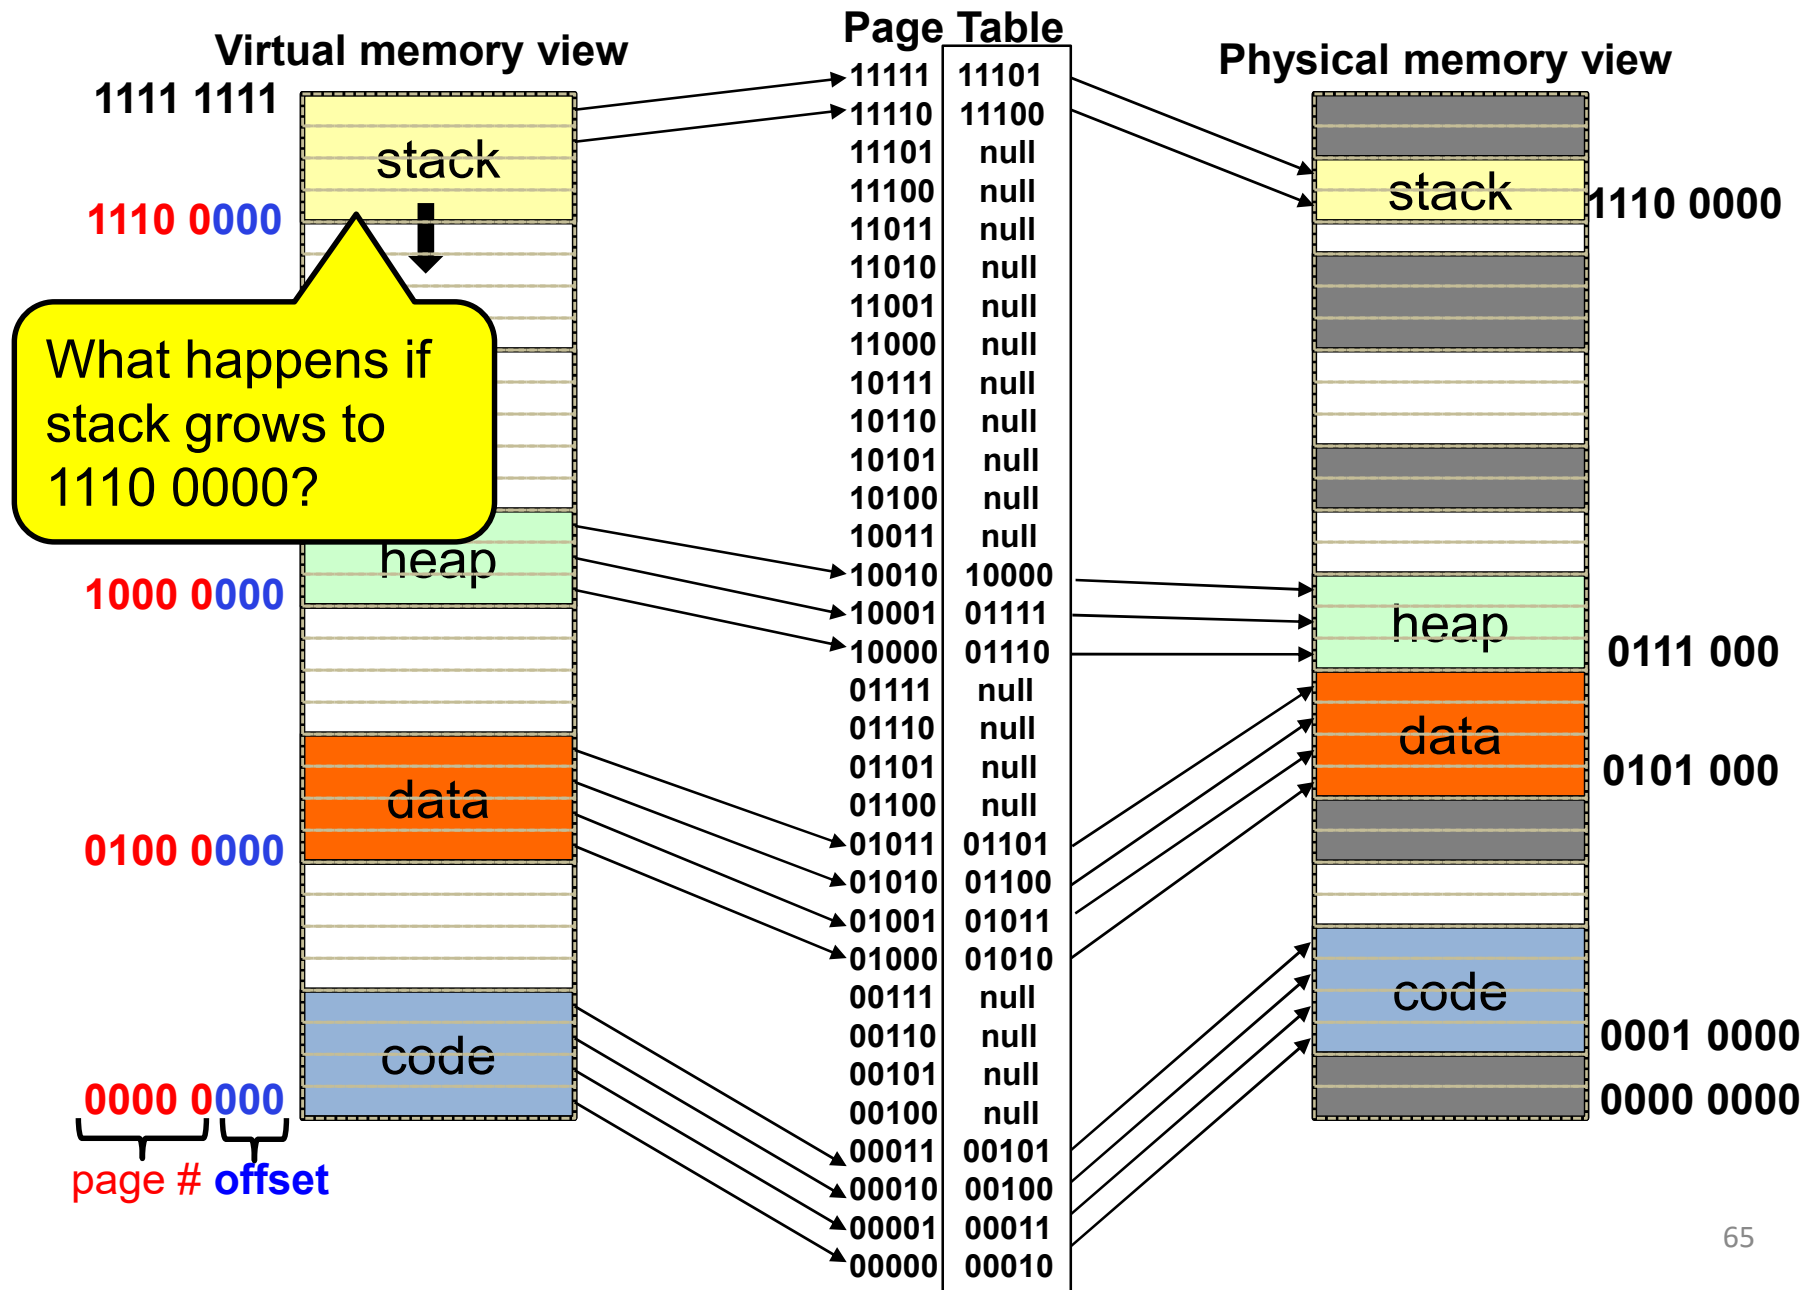
Shared pages - Example



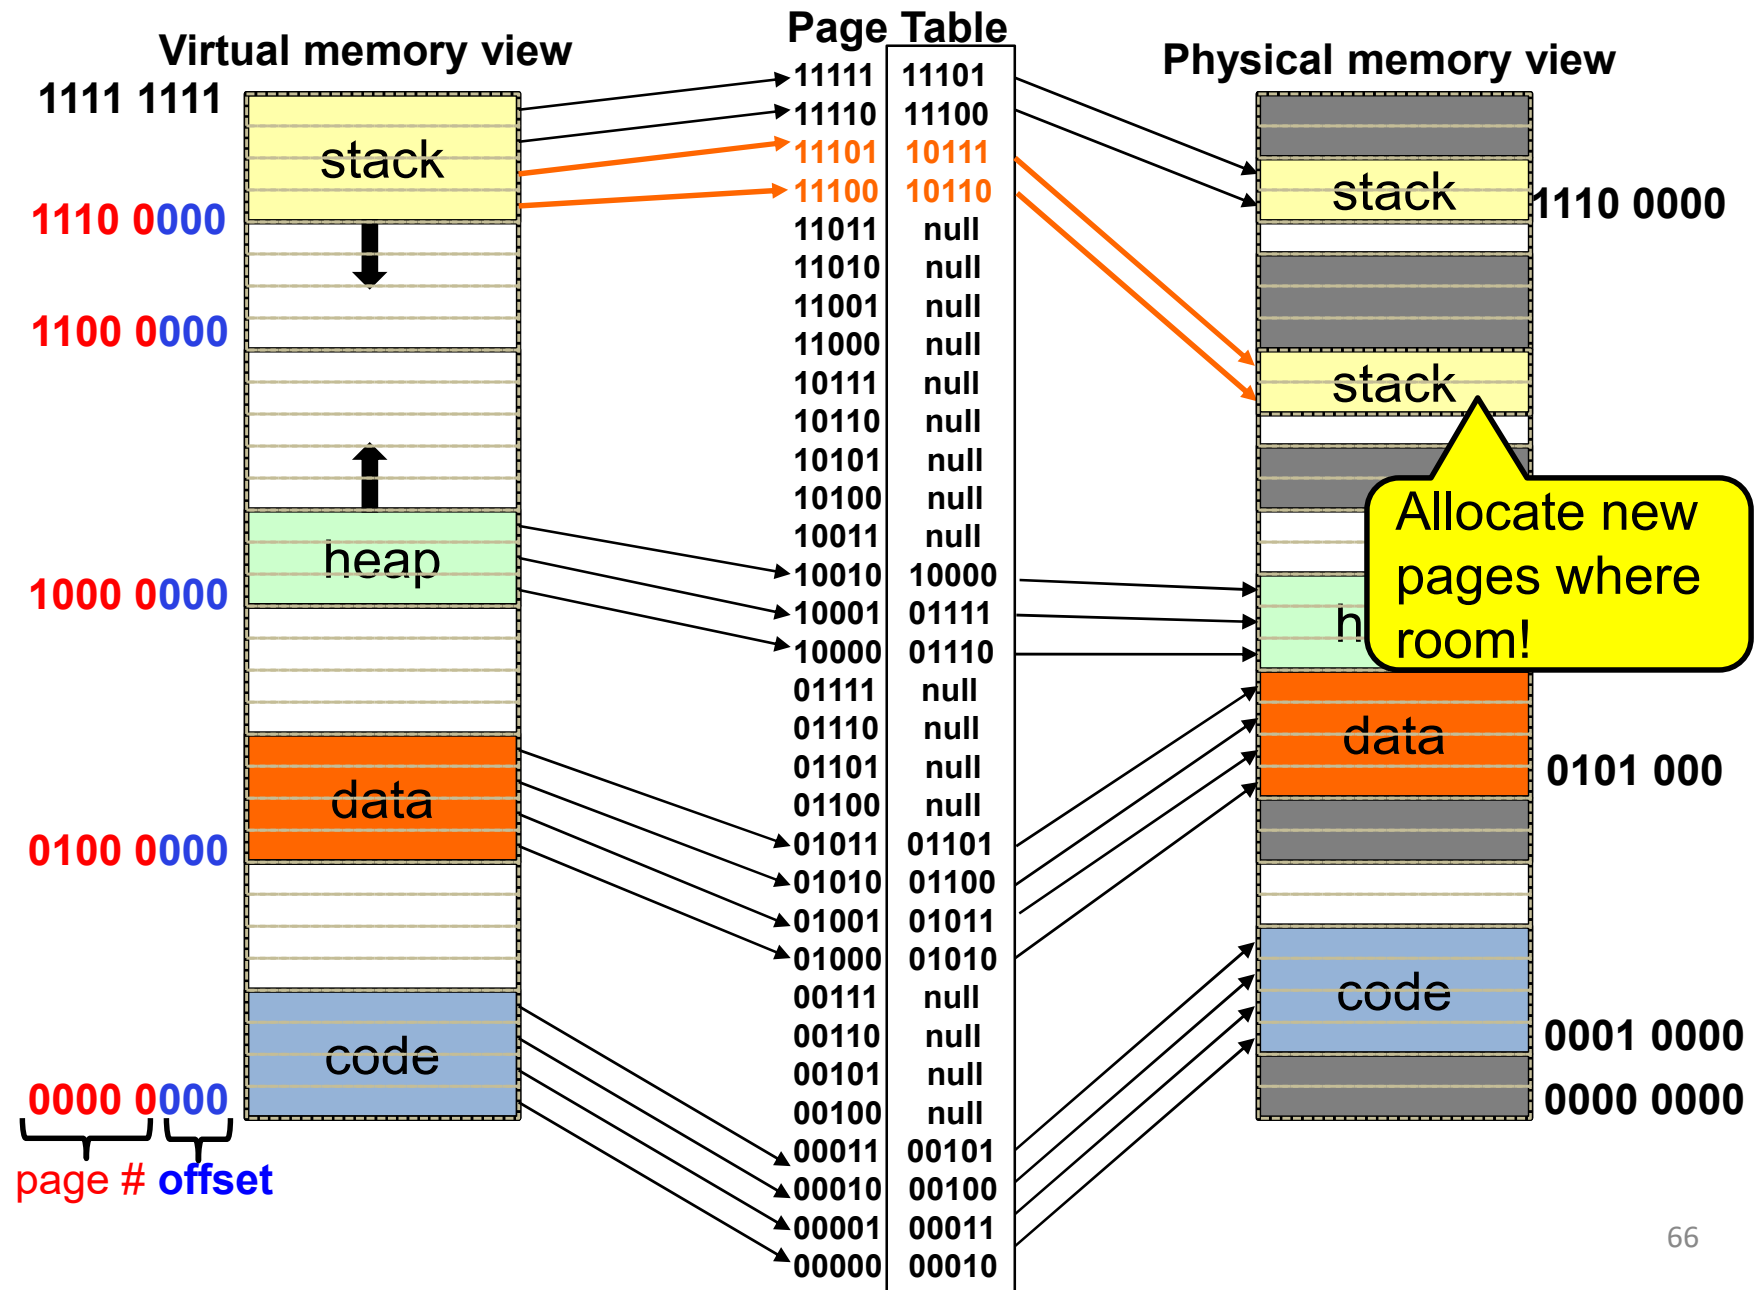
Summary: Paging



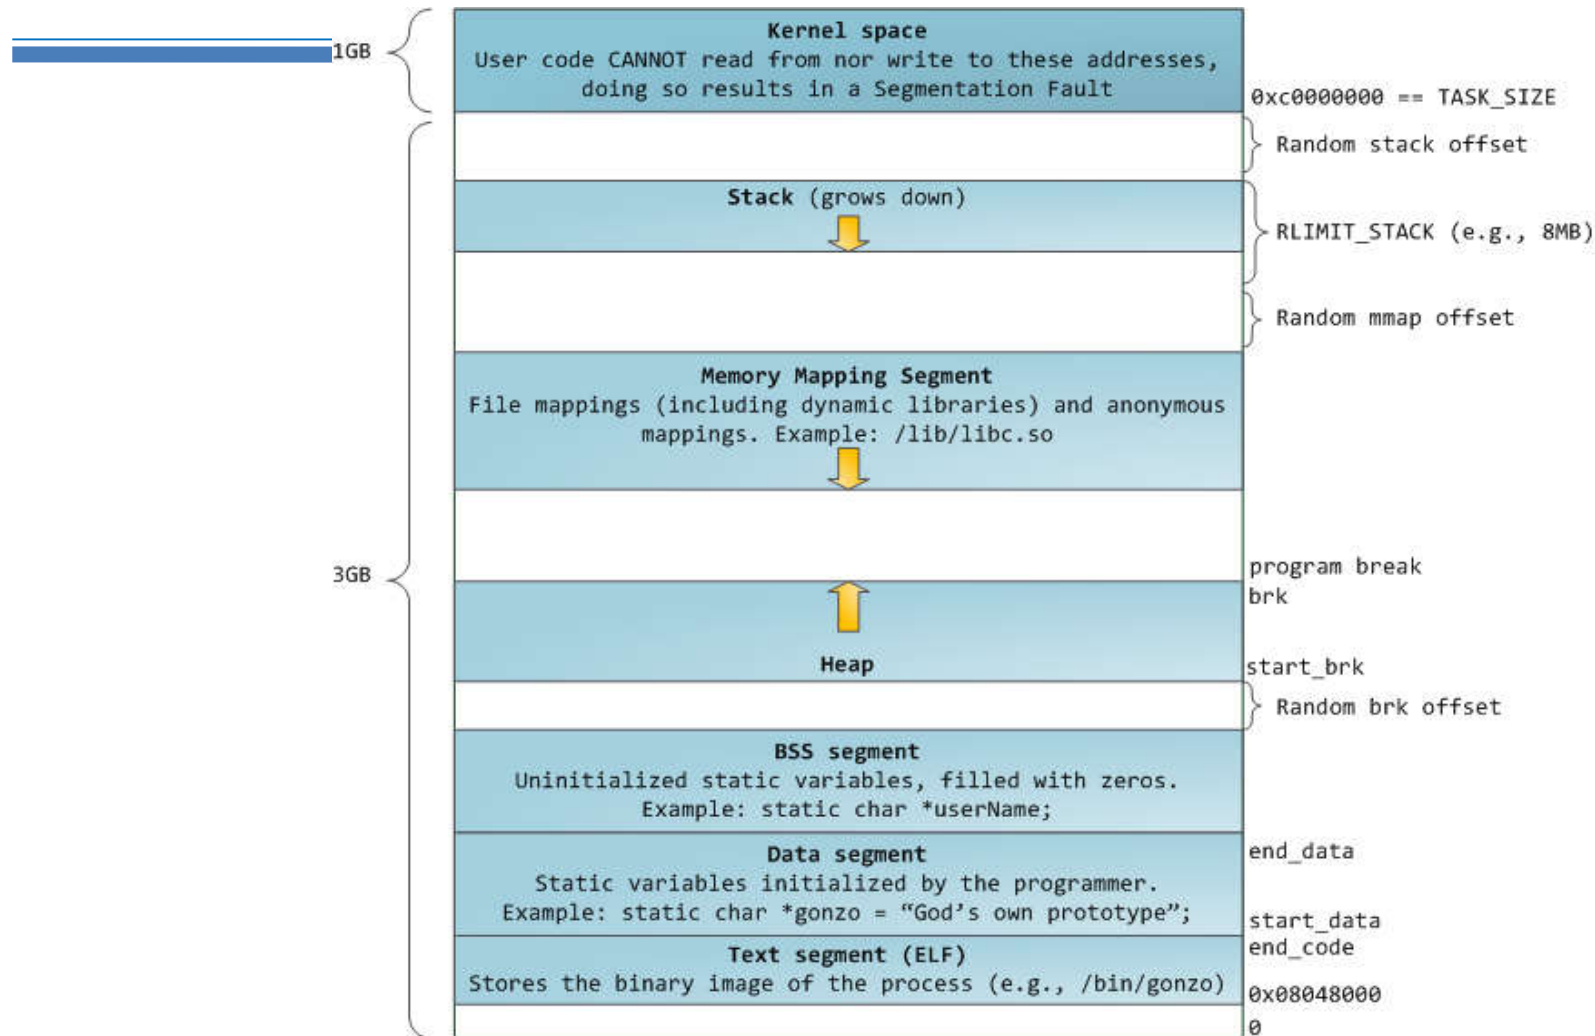
Summary: Paging



Summary: Paging



Example: Memory Layout for Linux 32-bit



<http://static.duartes.org/img/blogPosts/linuxFlexibleAddressSpaceLayout.png>

Problem of big page tables

- Memory structure for paging can get huge using straight-forward methods.
- Consider a 32-bit logical address space on a modern computers:
 - Page size of 4 kB = 2^{12}
 - Page table would have 1 million entries ($2^{32} / 2^{12}$)
 - If each entry is 4 B: 4 MB of physical address space memory for page table alone.
 - The amount of memory used, cost a lot
 - Don't want to allocate that contiguously in the memory

Page Table Discussion


- What needs to be switched on a context switch?
 - Page table pointer and limit
- Analysis
 - Pros
 - Simple memory allocation
 - Easy to share
 - Con: What if address space is sparse?
 - E.g., on UNIX, code starts at 0, stack starts at $(2^{31}-1)$
 - With 1K pages, need 2 million page table entries!
 - Con: What if table really big?
 - Not all pages used all the time \Rightarrow would be nice to have working set of page table in memory
- How about multi-level paging or combining paging and segmentation?

Solutions to maintain huge page tables

- A) Hierarchical paging
- B) Hashed page tables
- C) Inverted page tables


A) Hierarchical paging

- Use of **two-level page table**!
- We then **page** the **page table**!



Two-level paging - example

- A logical address, on 32-bit machine with 1 kB page size, is divided:
 - A page offset consisting 10 bits
 - A page number consisting of 22 bits
- Since the page table is paged, the page number is divided into:
 - A 12-bit page number
 - A 10 bit page offset
- Thus a logical address is:

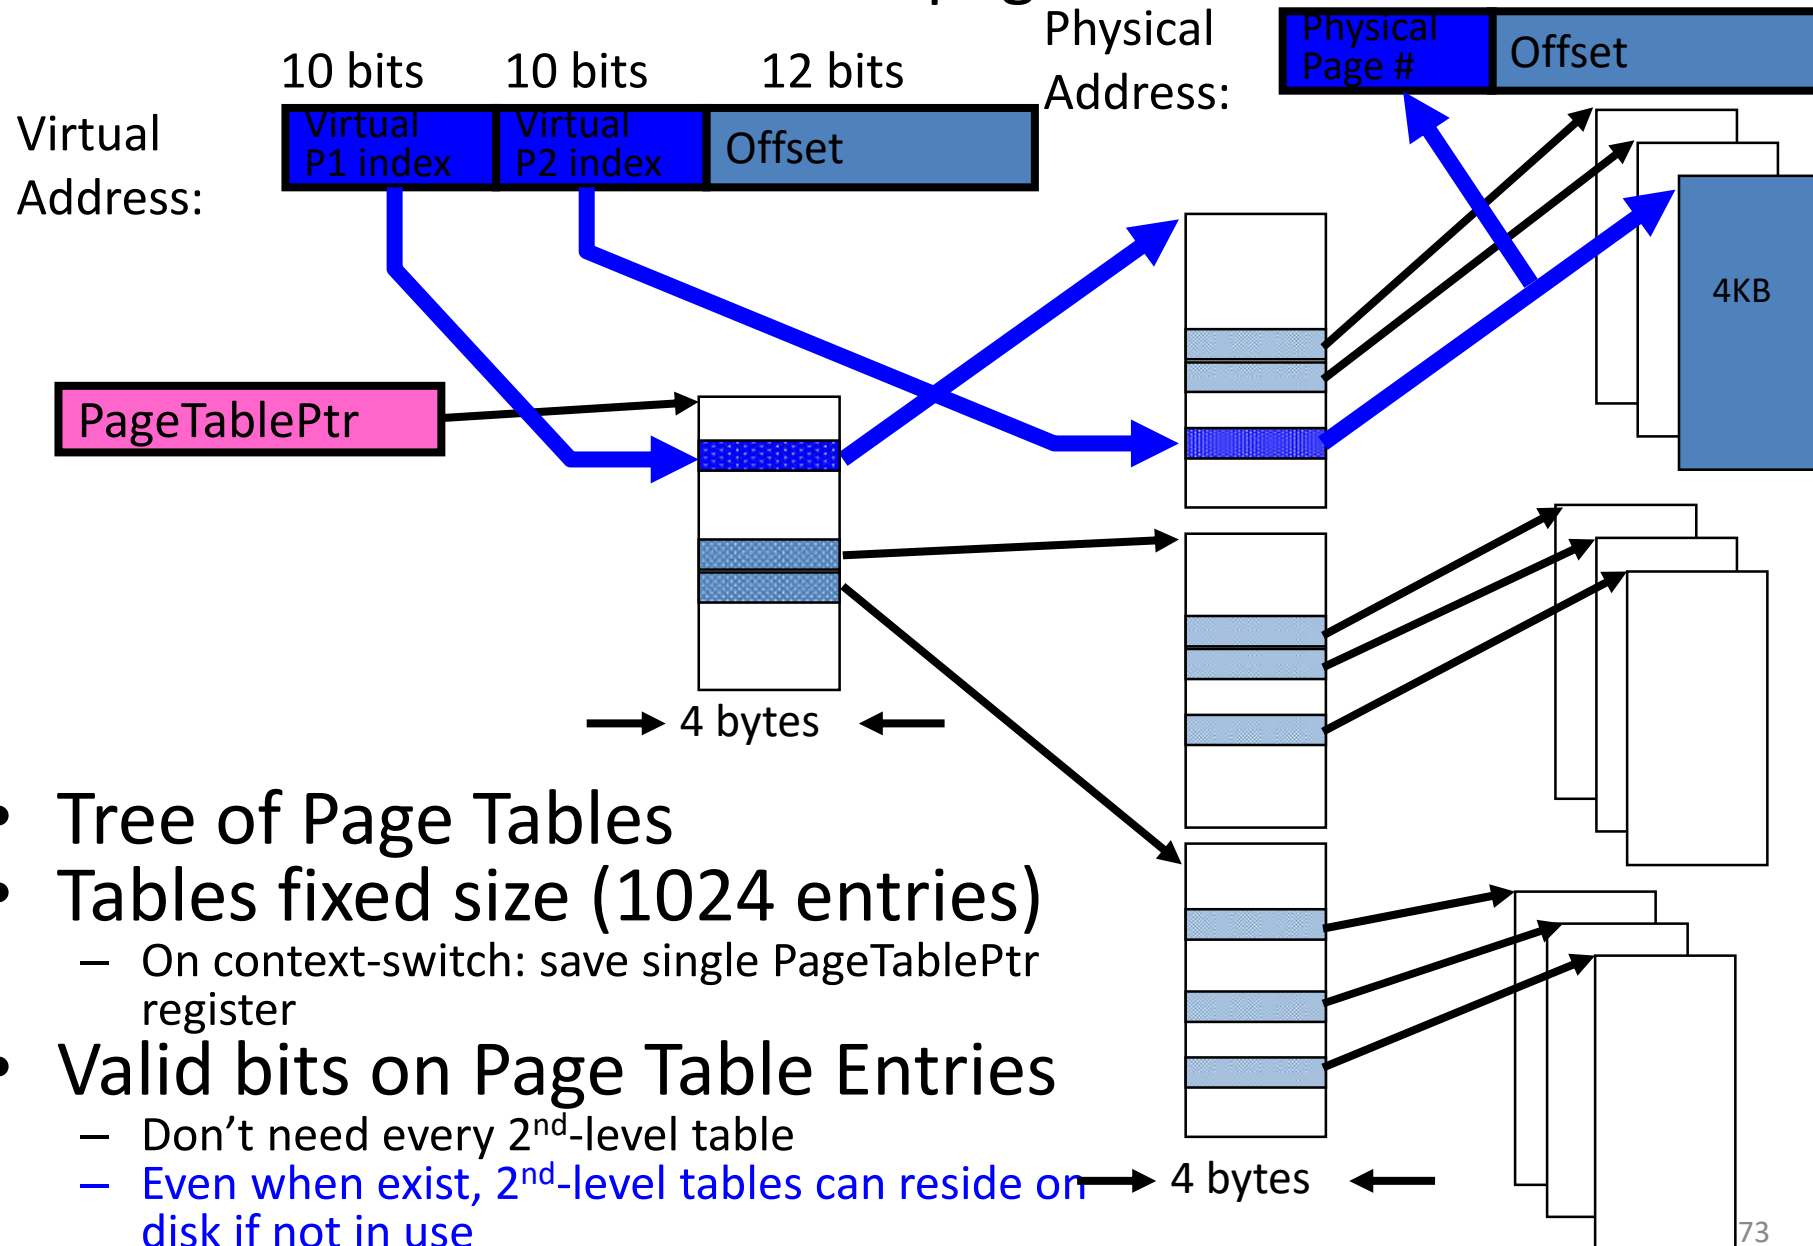


where p_1 is an index into outer page table, and p_2 is the displacement within the page of inner page table

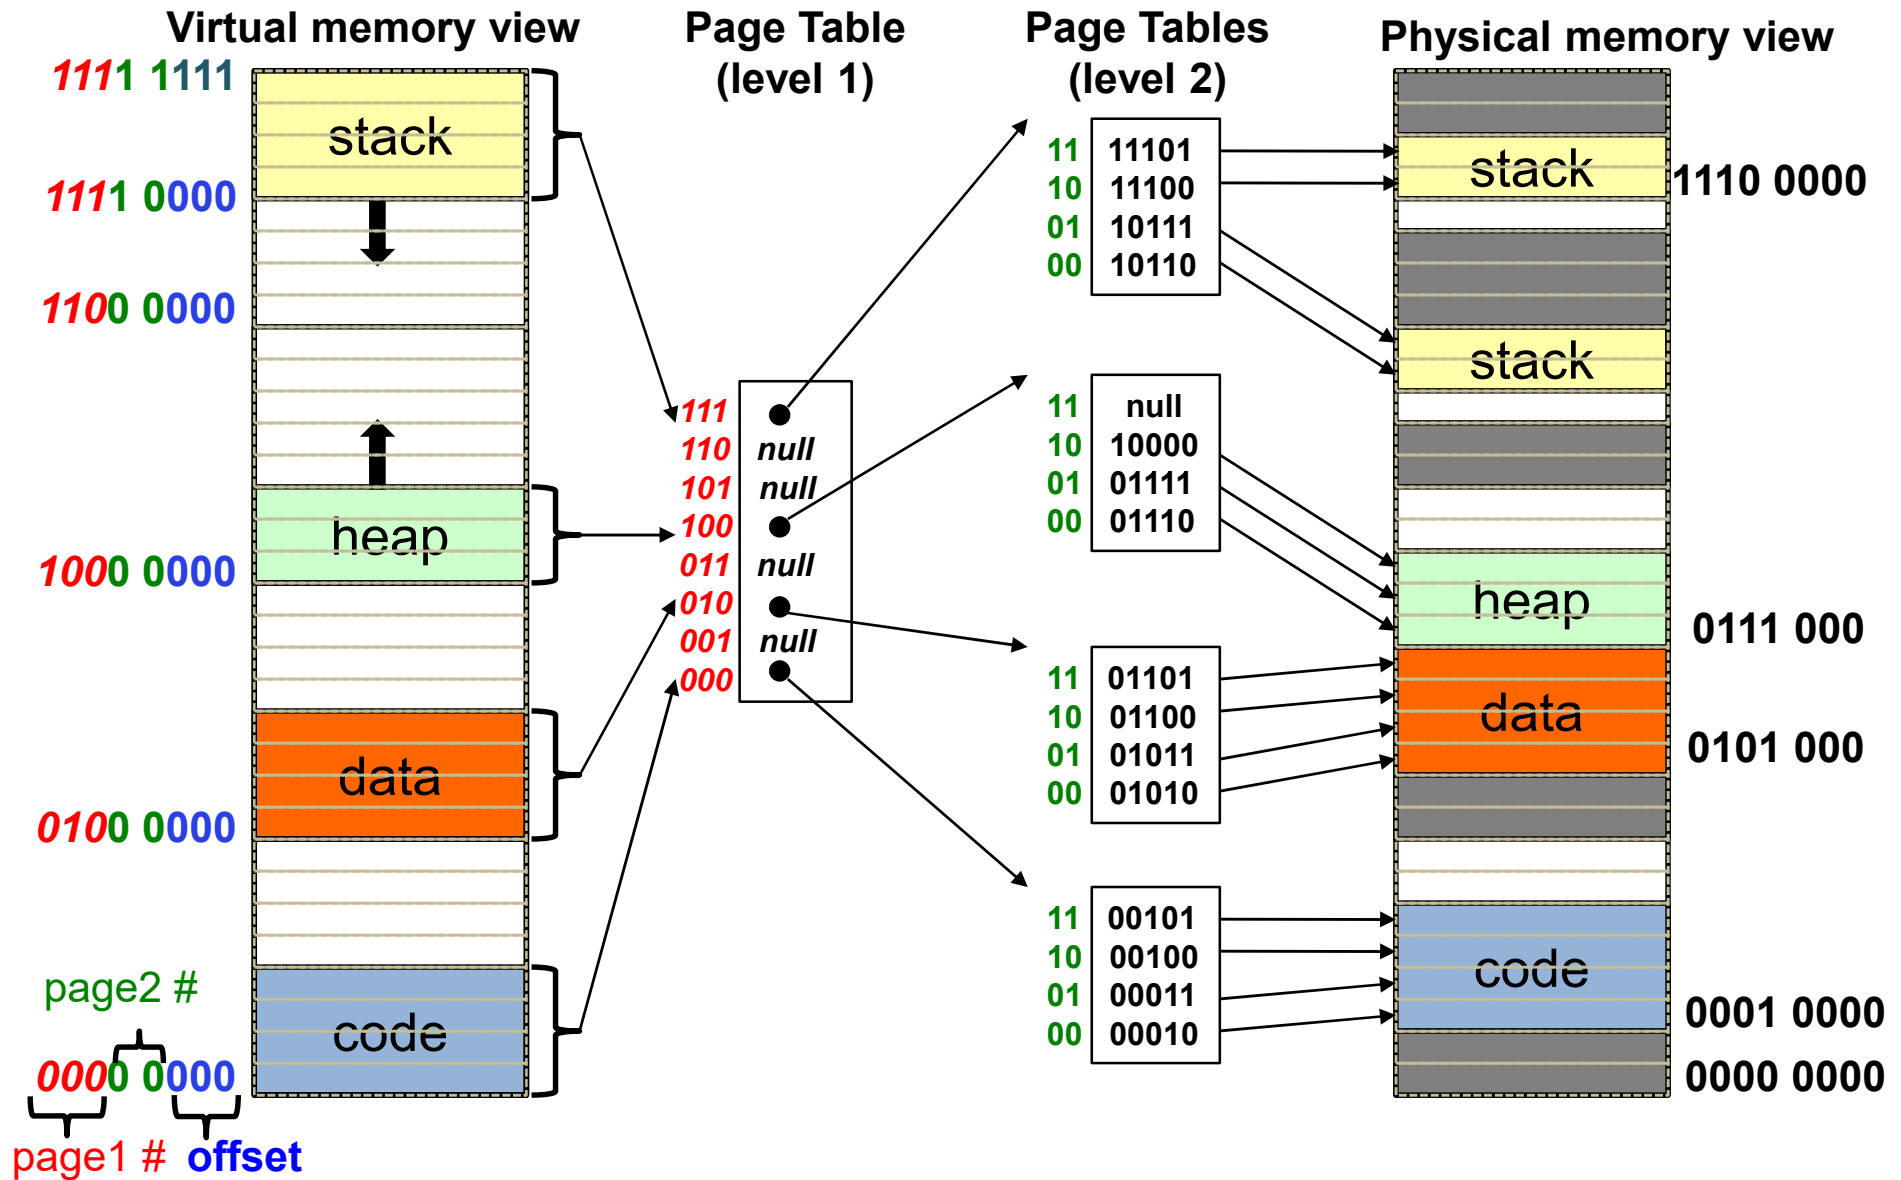
- Known as forward-mapped page table

Fix for sparse address space:

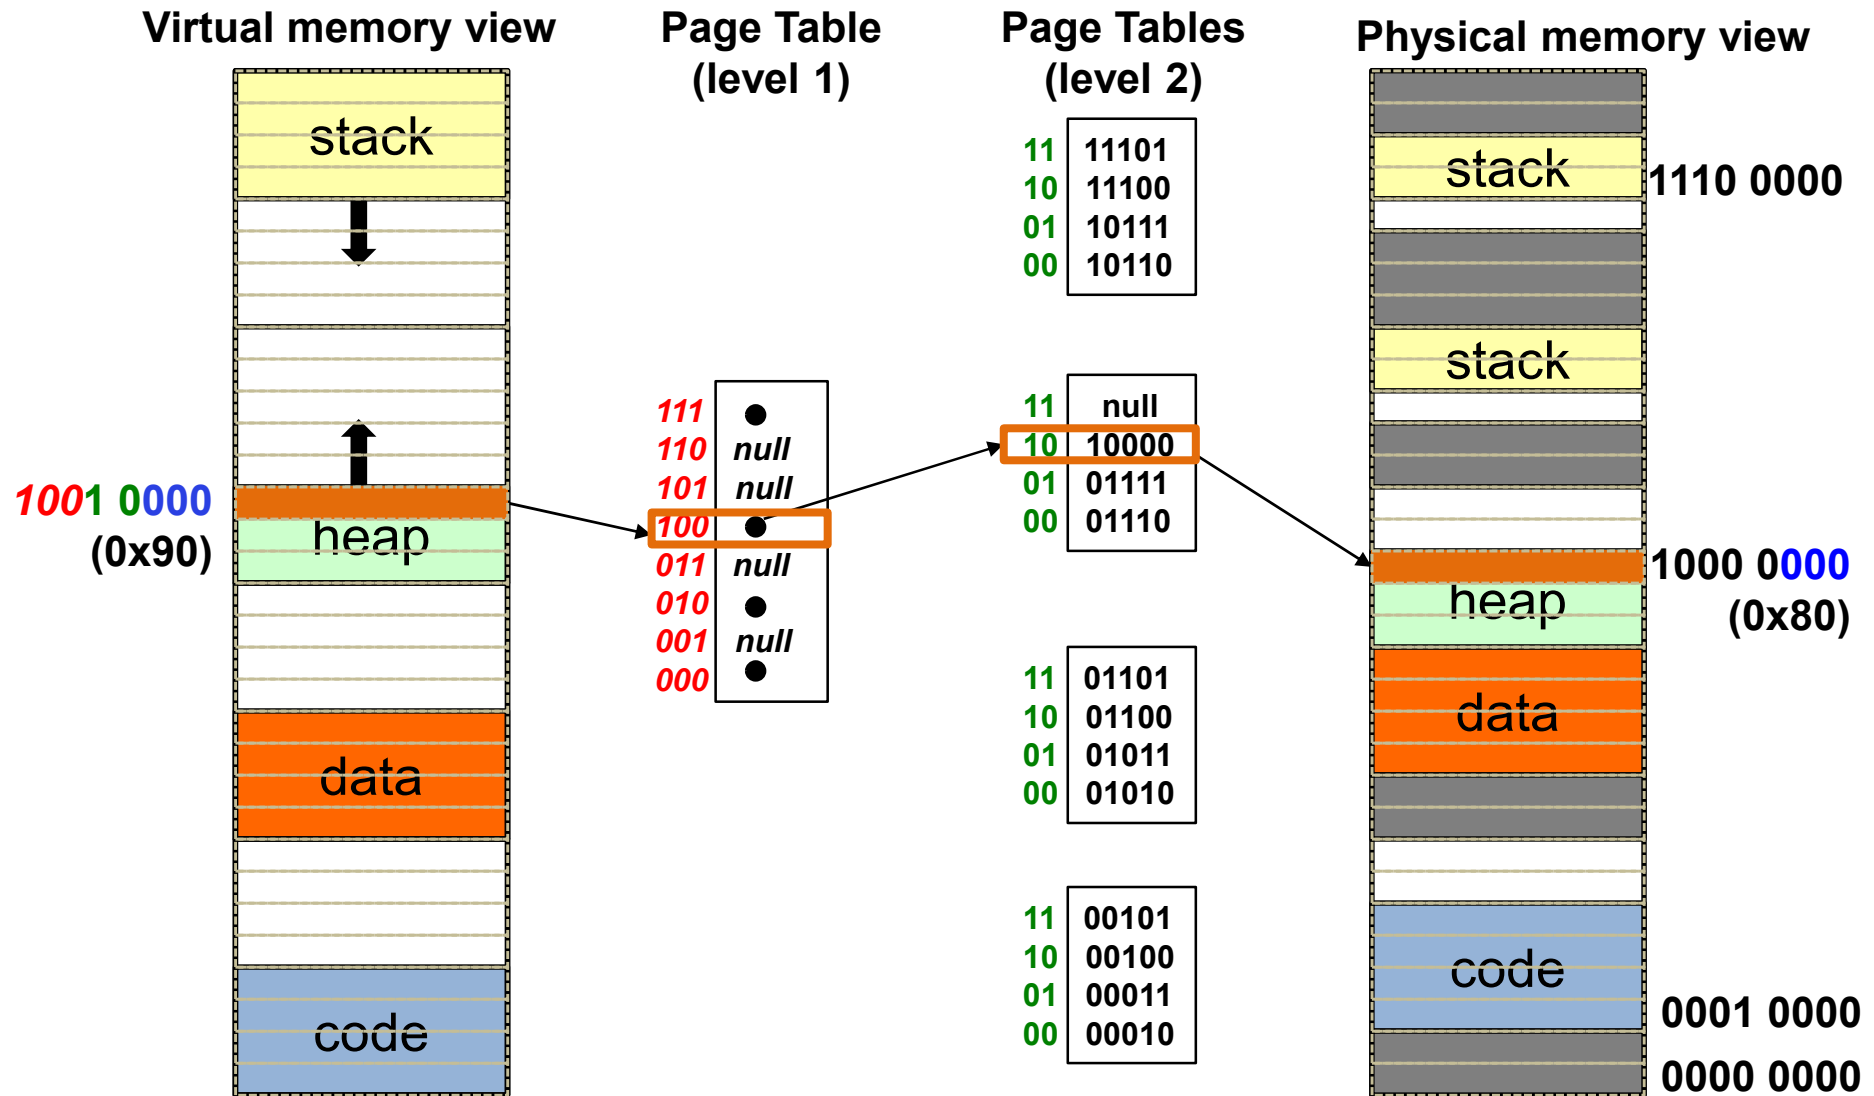
The two-level page table



Summary: Two-Level Paging




Summary: Two-Level Paging



64-bit logical address space

- Even two-level paging scheme is not sufficient
- If page size of 4 kB (2^{12})
 - Page table entries = 2^{52}
 - Assuming each page table entry size = 4 B
 - If two-level scheme, inner page table could be 2^{10} , 4 kB entries
 - Outer page table has 2^{42} entries or 2^{44} B
 - Address would look like:



Three-level paging scheme

- One solution is to add a 2nd outer page table
- But in the following example, the 2nd outer page table is still 2^{34} bytes in size!
- And possible 4 memory access to get to one physical memory location:

outer page	inner page	offset
p_1	p_2	d
42	10	12

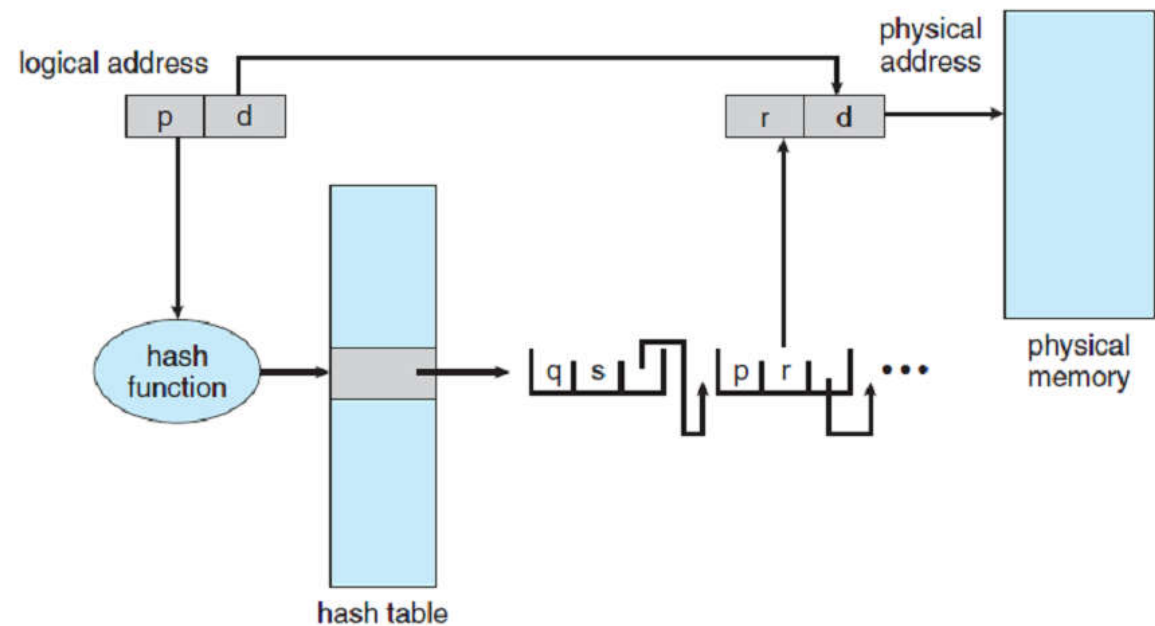
2nd outer page	outer page	inner page	offset
p_1	p_2	p_3	d
32	10	10	12

B) Hashed page tables

- Common in address space > 32 bits
- The logical page number is hashed into a page table
- This page table contains a chain of elements hashing to the same location

Hash page table scheme

- Each element contains
 - Logical page number
 - Physical frame number
 - Pointer to the next node
- Search is done serially in the linked list

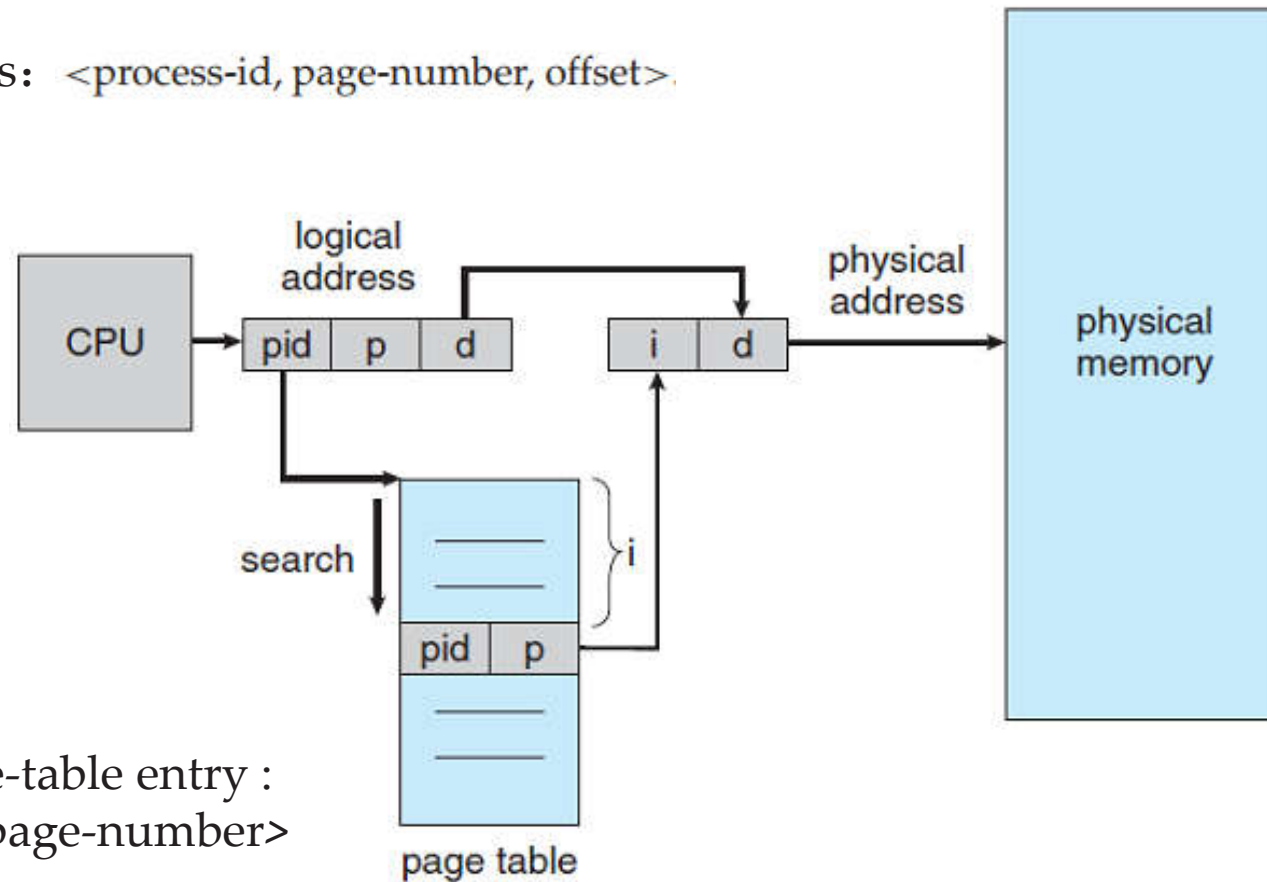


C) Inverted page table

- Rather than keeping all possible logical page numbers, track all physical pages (frame) numbers
- One entry for each physical page of memory
- Entry consists of
 - Virtual address of the page stored in that frame
 - + Process info

Inverted page table scheme

virtual address: $\langle \text{process-id, page-number, offset} \rangle$



inverted page-table entry :
 $\langle \text{process-id, page-number} \rangle$

Problem and solutions

- Good:
 - Decrease memory needed
- Bad:
 - Increase time needed to search the table
- Use hash table to limit the search to one, or at most a few, page-table entries.
- How to implement shared memory?
 - One mapping of a virtual address to the shared physical address




X86 architecture



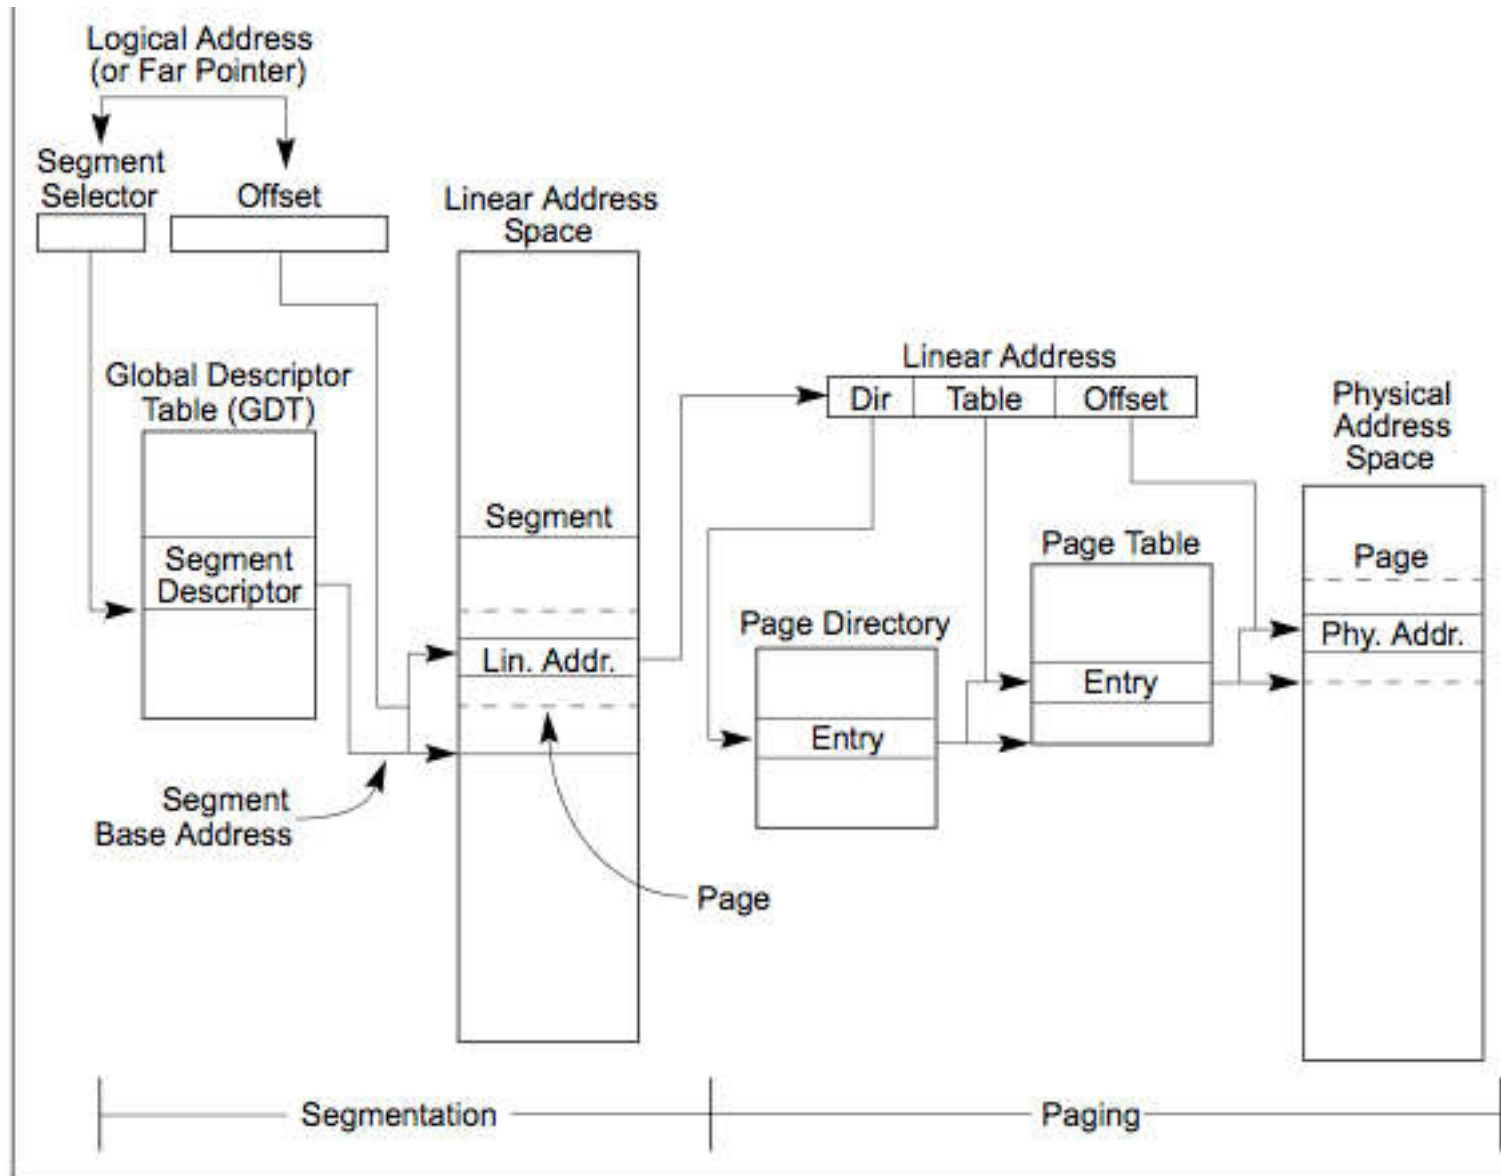
Example: Intel 32 Architectures

- The IA-32 architecture supported both **paging** and **segmentation**
 - The CPU generates logical addresses, which are given to the segmentation unit.
 - The segmentation unit produces a linear address for each logical address.
 - The linear address is then given to the paging unit, which in turn generates the physical address in main memory.



Making it real:

X86 Memory model with segmentation (16/32-bit)

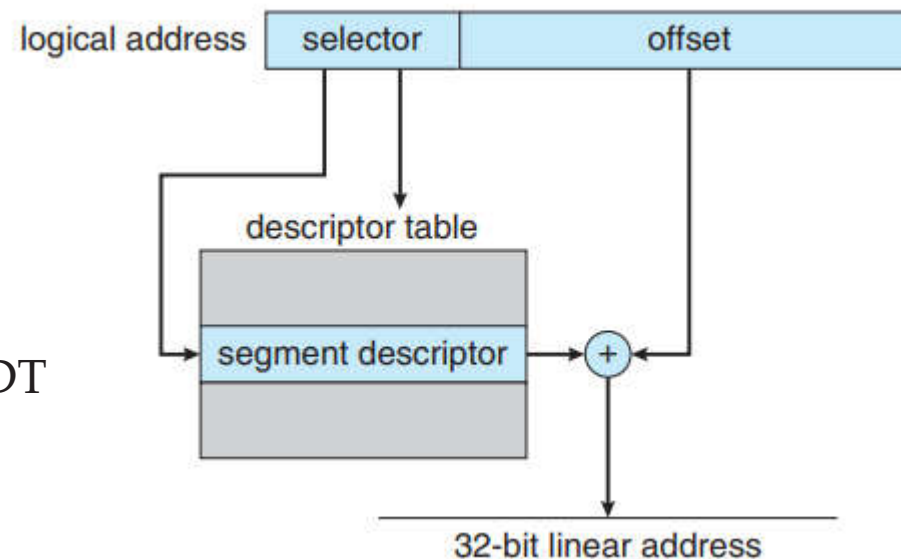


IA-32 Segmentation

The logical address is a pair (selector, offset), where the selector is a 16-bit number:




The segment register points to the appropriate entry in the LDT or GDT




X86 Segment Descriptors (32-bit Protected Mode)

- Segments are either implicit in the instruction (say for code segments) or actually part of the instruction
 - There are 6 registers: SS, CS, DS, ES, FS, GS
- What is in a segment register?
 - A *pointer* to the actual segment description:



G/L selects between GDT and LDT tables (global vs local descriptor tables)

- Two registers: GDTR and LDTR hold pointers to the global and local descriptor tables in memory
 - Includes length of table (for $< 2^{13}$) entries
- Descriptor format (64 bits):



G: Granularity of segment [Limit Size] (0: 16bit, 1: 4KiB unit)

DB: Default operand size (0: 16bit, 1: 32bit)

A: Freely available for use by software

P: Segment present


DPL: Descriptor Privilege Level

S: System Segment (0: System, 1: code or data)

Type: Code, Data, Segment


IA-32 Paging

- The IA-32 architecture allows a page size of either 4 KB or 4 MB



uses a two-level paging scheme

- The 10 high-order bits reference an entry in the outermost page table, which IA-32 terms the **page directory**
- The page directory entry points to an inner page table that is indexed by the contents of the innermost 10 bits
- the low-order bits 0–11 refer to the offset in the 4-KB page pointed to in the page table



Paging in the IA-32 architecture

Summary

- Segment Mapping
 - Segment registers within processor
 - Segment ID associated with each access
 - Often comes from portion of virtual address
 - Can come from bits in instruction instead (x86)
 - Each segment contains base and limit information
 - Offset (rest of address) adjusted by adding base
- Page Tables
 - Memory divided into fixed-sized chunks of memory
 - Virtual page number from virtual address mapped through page table to physical page number
 - Offset of virtual address same as physical address
 - Large page tables can be placed into virtual memory
- Multi-Level Tables
 - Virtual address mapped to series of tables
 - Permit sparse population of address space

Questions?

